

Bachelor's Thesis

in the Bachelor's Program 'Applied Mathematics and
Computer Science'

Backtesting and Live-Testing of Classic and AI-Powered Trading Strategies

by

Jason Becker

Enrollment Number: 3567285

Supervisor: Prof. Dr. Bodo Kraft
Co-Supervisor: Hendrik Karwanni, M. Sc.
Submitted on: 8th July, 2025

Declaration of Authorship

This bachelor thesis has been written by:

Jason Becker
Wachtelstraße 24
40789 Monheim am Rhein

I hereby declare that this bachelor thesis has been written independently and without unauthorized assistance. I confirm that I have used only the sources and aids indicated in the bibliography. Furthermore, I declare that I have marked all passages that are either verbatim or paraphrased versions of sources.

I am aware that any form of plagiarism or dishonesty in academic work constitutes a serious offense and may lead to disciplinary measures.

Monheim am Rhein, 8th July, 2025

.....
Jason Becker

Abstract

Contents

1	Introduction	1
1.1	Motivation	1
1.2	Aim of this Paper	1
2	Data Source and Broker Selection	2
2.1	Broker Selection	2
2.2	API Connection and Data Retrieval Process	3
3	Exploratory Data Analysis	4
3.1	Statistics	4
3.2	Dividing the Data	4
3.3	Using Log>Returns	6
3.4	Additional Features	7
3.4.1	Trend Following Indicators	7
3.4.2	Volatility Indicators	8
3.4.3	Momentum Indicators	9
3.4.4	Price Transformation Indicators	9
3.5	Scaling the Data	9
3.6	Pricipal Component Analysis	9
4	Market Regimes	12
4.1	Introduction to Market Regimes	12
4.2	How to Categorize the Market?	12
4.3	Recognizing Market Regimes	13
4.3.1	Recognizing Trend Behavior	13
4.3.2	Recognizing Volatility	15
4.3.3	Combining Trend and Volatility Classification	16
4.4	Dividing Market Regimes in Durations	17
5	Money- and Risk-Management	18
5.1	Calculating the Position Size	18
5.2	Validating an Entry Signal	19
5.2.1	Risk-Reward-Ratio	19
5.2.2	Maximum Account Risk	20
5.2.3	Minimum Take Profit	20
5.3	Dealing with Trading Fees	21

6	Deep Learning Models	22
6.1	Metrics	22
6.2	Optuna	22
6.3	Neural Networks	23
6.3.1	Base Model with Flatten-Architecture	23
6.3.2	Dropout Neural Network	24
6.3.3	Residual Neural Network	25
6.4	Long Short-Term Memory Models	26
6.4.1	Base LSTM	27
6.4.2	Dropout LSTM	27
6.4.3	Bidirectional LSTM (BiLSTM)	28
6.4.4	Encode-Decode Model with Repeat-Vector	29
6.5	Convolutional Neural Networks	30
6.5.1	Base CNN	31
6.5.2	Deep CNN	32
6.5.3	Attention CNN	33
6.5.4	Concatenation CNN	33
6.5.5	CNN-GRU Hybrid Model	34
6.6	Transformer Models	35
6.6.1	Transformer Architecture	36
6.7	Regression-Models	38
6.7.1	Loss-Function	38
6.8	Classification-Models	39
6.9	Model Evaluation	39
7	Trading-Strategies	40
7.1	AI-Trading-Strategies	40
7.2	Classic Trading-Strategies	40
7.3	Performance Comparison	40
8	Trading-Engine	41
8.1	Use-Cases of the Trading-Engine	41
8.2	Architecture	42
8.2.1	Plugin Architecture	43
8.2.2	Core Modules	44
8.2.3	Applications	45
8.3	Demo Broker	45
8.3.1	Order Execution and Position Management	46
8.3.2	Adapting Trailing Stop Positions	49

9	Backtesting Trading-Strategies	51
9.1	Executing Backtests	51
9.2	Evaluating Backtests	51
9.3	Fee-Impacts	51
10	Stress Testing	52
11	Live-Testing	53
11.1	Broker Setup	53
11.2	Connecting to ByBit	53
11.3	Live-Test Results	53
12	Conclusion	54
12.1	Key Findings	54
13	Aim of further Works	54
14	References	55

List of Figures

1	Price Fluctuation of ETH in USDC	6
2	Log Returns of ETH in USDC	7
3	Cumulative Explained Variance	11
4	Trend Classification	15
5	Volatility Classification	16
6	Regime Classification	16
7	Long Position Decision	38
8	Long Position Decision	39
9	Trading-Engine Components	42
10	Opening a Single Position	47
11	Closing a Position	48
12	Opening a Position	49
13	Trailing Stop Example	50

List of Tables

1	Broker Comparison	3
2	Data Split	4
3	Train Data	5
4	Validation Data	5
5	Test Data	5
6	Backtest Data	6

Code Listing

A	Base NN	24
B	Dropout NN	25
C	Residual NN	26
D	Base LSTM	27
E	Dropout LSTM	28
F	Bidirectional LSTM	28
G	Encode-Decode LSTM	29
H	Base CNN	31
I	Deep CNN	32
J	Attention CNN	33
K	Concatenation CNN	34
L	CNN + GRU	35
M	Transformer Model	36
N	SPI Definition	43
O	File-Repository Implementation	43
P	Database-Repository Implementation	43

List of Abbreviations

EMH	Efficient Market Hypothesis
BTC	Bitcoin

1 Introduction

1.1 Motivation

1.2 Aim of this Paper

+ Noch ein Kapitel welches Kurz erklärt, das ETHUSDC M1 genommen wird und warum

2 Data Source and Broker Selection

Cryptocurrency brokers (also called crypto brokers) play an important role in cryptocurrency trading. Among other things, they act as intermediaries between different market participants. Their key tasks include:

1. **Providing access:** Individuals can participate in the market through a broker and thereby trade various cryptocurrencies. This includes executing orders such as buying cryptocurrencies at the lowest available price or selling them at the highest available price.
2. **Security, and Compliance:** They also provide customers with a secure platform for executing transactions and adhere to the financial regulations established by authorities.
3. **Leveraging:** Brokers offer customers the opportunity to borrow money, and thus trade with more capital than they actually have in their account.

This has the advantage that trading with cryptocurrencies is much easier and safer, but one of the biggest disadvantages are the fees that are incurred when using [1].

2.1 Broker Selection

For this paper, one broker must be selected for data retrieval, and live testing. Since the process is fully automated in short time-frames, the broker must meet certain requirements.

The API must be able to stream market data, request historical data, the current account balance, closed trades, and currently open positions, placing orders, and positions, as well as canceling unfilled orders.

Apart from the API the broker must support leveraged long/short products like CFDs or margin trading. They also must provide data in high quality as well as a demo depot. The further they must be regulated in the European Union with the lowest possible fees.

Table 1 summarizes the required features for some potential brokers. All listed there meet the API functionality requirements.¹

¹Sources: [2], [3], [4], [5], [6], [7]

Broker	Tradable assets	Fees		Leverage
		Maker	Taker	
ByBit	Spot, Spot with leverage, Futures, Options	0.02%	0.055%	10:1
IG	CFDs, Knock-out-Options	Spread (approx. \$1.30)		2:1
Capital.com	CFDs	Spread (approx. \$1.75)		2:1

Table 1: Broker Comparison

Taking into account [Table 1](#), ByBit is the best broker because it has the lowest fees, high quality data, the highest possible leverage as well as a regularization in the EU.

2.2 API Connection and Data Retrieval Process

Before starting with the Machine Learning process, and the backtests, the first step is to download historical ETH/USDC M1 via the ByBit API. The request was executed on the `/v5/market/kline` API-Endpoint [8] with the category `linear`, symbol `ETHPERP`, and interval `1` at 17th June, 2025, 11:30 UTC+2. Since ByBit only returns 1000 candlesticks per request, the same request with different start-, and end-times was executed until the ByBit API does no longer return older candlestick data. This resulted in a candlestick data pool with data on a minute basis from 5th August, 2022, 10:00 UTC+2 to 17th June, 2025, 11:30 UTC+2. Chapter 3.1 will go into more detail about the data.

3 Exploratory Data Analysis

3.1 Statistics

3.2 Dividing the Data

Unlike classic machine learning processes, where data is splitted in three subsets, named train-, validation-, and test-set, here the data is splitted in four subsets. The fourth data-set is used for backtesting the real trading strategy, and is therefore not part of the machine learning process but plays an important role in developing the final trading strategy.

Set	From	To	No. of Datapoints	% of All Data
Complete	08/05/2022 10:00 UTC+2	06/17/2025 11:30 UTC+2	1,507,598	
Train	08/05/2022 10:00 UTC+2	04/30/2024 23:59 UTC+2	913,684	60.6%
Validation	05/01/2024 00:00 UTC+2	09/30/2024 23:59 UTC+2	220,320	14.6%
Test	10/01/2024 00:00 UTC+2	12/31/2024 23:59 UTC+2	132,482	8.8%
Backtest	01/01/2025 00:00 UTC+2	06/17/2025 11:30 UTC+2	241,111	15.9%

Table 2: Data Split

After the splitting, the four subsets have the following summaries:

	Open	High	Low	Close	Volume
	Open	High	Low	Close	Volume
count	913684.0	913684.0	913684.0	913684.0	913684.0
mean	1933.4	1933.95	1932.85	1933.4	7.58
std	619.6	619.92	619.28	619.6	105.09
min	1074.35	1077.45	1064.05	1074.35	0.0
25%	1578.44	1578.89	1577.99	1578.44	0.0
50%	1801.52	1801.81	1801.13	1801.52	0.05
75%	2083.4	2083.84	2083.05	2083.4	2.02
max	4098.5	4099.48	4096.35	4098.5	31350.65

Table 3: Train Data

	Open	High	Low	Close	Volume
	Open	High	Low	Close	Volume
count	220320.0	220320.0	220320.0	220320.0	220320.0
mean	3050.36	3051.3	3049.4	3050.36	2.89
std	473.4	473.45	473.34	473.4	21.45
min	2111.8	2160.4	2088.13	2111.8	0.0
25%	2617.31	2618.04	2616.71	2617.31	0.0
50%	3063.4	3064.31	3062.43	3063.4	0.12
75%	3462.22	3463.29	3461.21	3462.22	1.32
max	3974.68	3976.96	3969.94	3974.68	4972.18

Table 4: Validation Data

	Open	High	Low	Close	Volume
	Open	High	Low	Close	Volume
count	132482.0	132482.0	132482.0	132482.0	132482.0
mean	3093.93	3095.27	3092.58	3093.94	4.42
std	531.85	532.29	531.4	531.85	18.42
min	2309.01	2311.73	2307.73	2309.01	0.0
25%	2541.8	2542.65	2540.94	2541.8	0.08
50%	3143.69	3145.38	3142.02	3143.7	0.68
75%	3487.9	3489.46	3486.33	3487.9	2.92
max	4107.28	4112.68	4102.6	4107.28	1341.05

Table 5: Test Data

	Open	High	Low	Close	Volume
count	241111.0	241111.0	241111.0	241111.0	241111.0
mean	2432.57	2433.77	2431.35	2432.57	7.0
std	574.72	574.95	574.49	574.72	32.11
min	1386.6	1395.8	1382.99	1386.6	0.0
25%	1887.22	1888.28	1886.1	1887.22	0.18
50%	2510.29	2511.4	2509.1	2510.29	1.26
75%	2732.0	2733.21	2730.7	2732.0	4.94
max	3742.33	3745.13	3739.65	3742.33	2534.66

Table 6: Backtest Data

3.3 Using Log>Returns

In the summary statistics of the subsets (Table 3, Table 4, Table 5, Table 6) it is noticeable that the mean values change over time. This becomes also clear when visualizing the data (Figure 1).




Figure 1: Price Fluctuation of ETH in USDC

To avoid this data drift, the price is transformed to its logarithmic returns (also called log-returns). These are calculated as follows:

$$LogReturn_t = \ln\left(\frac{Price_t}{Price_{t-1}}\right) \quad (1)$$

After the transformation, the means, and standard deviations in the subsets are very similar, and the data does no longer drift over time.




Figure 2: Log Returns of ETH in USDC

The transformation is applied to the open, high, low, and close prices, and the original prices are replaced by the logarithmic returns, so that all subsequent actions are carried out with the prices on the logarithmic returns.

3.4 Additional Features

To provide the machine learning models more context about the price, additional features from different categories were added to the raw data.

3.4.1 Trend Following Indicators

In financial analysis trend following indicators play an essential role while modeling and predicting future price movements. This occurs because markets move in trends that are repeatedly interrupted by outliers. This results in a zigzag movement that nevertheless moves in one direction. Trend-following indicators can be used to filter out these outliers [9].

The exponential moving average (EMA) is one type of trend following indicator. It is a variation of the classic simple moving average (SMA), placing more emphasis on newer prices. It is often used by traders in length of 10-, 50-, and 200-period. One limitation is that many traders believe that

new data better reflects the current trend, where many others believe that overweighting recent prices creates a bias [10].

Because the EMA reacts faster to price changes than the SMA, and the aim of this paper are short-term predictions, the EMA could provide more relevant context for the model. The EMA was added in 5, 10, 20, 30, 50, and 200 period to the data.

Another Trend following indicator is the moving average convergence/divergence (MACD) which does not only help to identify price trends, but also helps to measure the trend momentum. It shows the relationship between two exponential moving averages. To calculate the MACD line, an EMA(12) is subtracted from an EMA(26). Additionally a signal line is calculated as an EMA(9) of the MACD line.

Although the MACD can signal possible reversals, it is also known for creating many false positives. This often happens if the market moves sideways [11].

3.4.2 Volatility Indicators

To measure volatility there are also other indicators, and techniques to measure the volatility, in addition to those described in [subsection 4.2](#).

One indicator is the average true range (ATR) which decomposes the entire range of an asset price for a period. It is calculated by determining the so-called true range (TR) for each candlestick - the maximum of: current high minus low; distance from the previous closing price up, and down. The ATR is then the moving average of these TR values, usually over 14 periods.

The ATR has two main limitations. The first is that an ATR value must always be set into comparison to previous ATR values, because one single value is not enough to tell if a trend is going to reverse. The second limitation is that the ATR does not tell anything about the direction of the price [12]. The ATR was added in 5, 7, 10, 14, and 18 period to the data.

Another volatility indicator are Bollinger Bands, which consist of three lines. The middle line is a SMA of the closing prices, the lower line is calculated by subtracting a certain number of standard deviations from the middle line, and the upper line is calculating by adding a certain number of standard deviations to the middle line. Usually the double of the standard deviation is added, and subtracted from the middle line.

The higher the volatility of the market is in the last closing prices, the wider the band gets. If the price of the market rises near the upper band, traders see the market as overbought. Similar if the market falls near the lower band, the market could be oversold. This allows to generate possible entry, and exit signals [13]. The three Lines were added to the data with a

15, 20, and 25 period SMA.

3.4.3 Momentum Indicators

Momentum measures the strength, and direction of a price movement over a certain period of time. Momentum indicators are useful because they give insights into the strength of trending prices. Therefore they can indicate possible reversals in the trend direction [14].

A common momentum indicator is the relative strength index (RSI). It measures the speed, and magnitude of an assets price by comparing the average gains, and losses of the asset, and can be used to detect overbought, and oversold conditions. The RSI ranges between zero, and 100. Usually an RSI over 70 indicates an overbought, and an RSI below 30 indicates an oversold market. Commonly the default RSI period to compare the average gains, and losses is 14 [15]. The RSI was added in periods 7, 14, and 20 to the data.

To depict relative trend strength, a sophisticated momentum indicator was constructed that compares the log returns of two different time frames. This feature allows the model to distinguish phases of accelerating price movements from stable or declining trends. The use of logarithmic returns simultaneously achieves scale independence, and improved comparability, which is particularly advantageous for modeling financial market-related time series. This indicator was added for time frames M2, M3, M6, M9, and M12.

3.4.4 Price Transformation Indicators

Apart from the mentioned indicators shifted logarithmic returns for the last six minutes have been added to provide additional context about the last price movements in a compact form. This could help the models to recognize trend reversals, volatility changes or short-term patterns.

Lastly the logarithmic returns of other time frames (M2, M3, M6, M9, and M12) have been added to the data providing another more stable trend context which helps to correctly classify short-term price movements. This creates a balanced feature set that takes into account both rapid reactions, and long-term patterns.

3.5 Scaling the Data

3.6 Principal Component Analysis

The principal component analysis (PCA) is a process for dimensional reduction, by linearly transforming high dimensional datasets to a small number

of uncorrelated principal components (directions of the new coordinate system). During the transformation, it can be specified how much variance in the data can be eliminated. After a transformation using PCA, it is ensured that at least the specified variance is retained [16].

Especially when processing numerous technical indicators or derived features in financial data, the high dimensionality can become problematic - a phenomenon known as the curse of dimensionality. This term describes the increasing challenges in modeling as the number of dimensions or features increases. Data points become increasingly sparsely distributed, computational costs increase, and many models lose their ability to generalize. Applying PCA allows redundant or correlated information to be condensed, making the model more robust, faster, and easier to interpret. At the same time, the risk of overfitting is reduced because the model focuses on the most important structures in the dataset [17].

Figure 3 shows the cumulative explained variance for the 56 added features in subsection 3.4 for each quantile market regime. It shows that the cumulative variance increases rapidly at the beginning. In this case, the reduction of the project to just 3 to 6 principal components already explains at least 80% of the variance. This means that the majority of the statistically relevant structures in the dataset are retained, even though the number of features has been massively reduced. Even if 20% of the variance is lost, the benefit outweighs this: The remaining principal components capture the statistical essence of the original feature space in a significantly more compact, and robust form, which is particularly well-suited for machine learning.




Figure 3: Cumulative Explained Variance

4 Market Regimes

Cryptocurrency markets are subject to constant change, which is reflected not only in price movements but also in the underlying structures, and dynamics. In quantitative analysis, and algorithmic trading, understanding these changes is essential for developing, and adapting robust trading strategies. A central concept in this context are market regimes.

4.1 Introduction to Market Regimes

Market regimes describe phases with distinct statistical, and economic characteristics, such as volatility, and trend behavior. They can be understood as different "states" of the market in which certain trading patterns dominate. Distinguishing between, for example, upward, sideways, and downward trends or high, and low volatility enables more targeted strategy selection, and adaption. Accordingly, the identification, and classification of market regimes plays an increasingly important role in modern trading analysis. For example, a strategy that performs well in a stable uptrend may fail in a sideways movement or in periods of high volatility [18].

Understanding market regimes leverages traders, and analysts to design strategies that are adaptive, and therefore more robust. Through targeted adaptation of the parameters or the selection of different models, the performance can be increased, and the risk can be reduced.

4.2 How to Categorize the Market?

Categorizing different market characteristics is a common step in identifying market regimes. In literature, and practical applications, there exist different approaches to classify markets. A fundamental categorization is often made by the following dimensions:

1. **Trend behavior:** Markets can be categorized trend following (bullish/bearish) or trendless (sideways).
2. **Volatility:** The volatility of a market is often an indicator for insecurity or stability. High volatility can indicate periods of stress, while low volatility indicates calm markets.
3. **Liquidity:** In illiquid markets pricing processes can be different compared to liquid markets which has effect in strategies.

According to the analysis goal the categorization can be binary (e.g. bullish vs. bearish) or granular (e.g. a combination of trend behavior, and volatility). Also a combination of multiple indicators, named regime scores, is possible to capture more complex market structures.

In the current context the market is categorized by trend behavior, and volatility. This results in six categories:

1. Downtrend + Low Volatility
2. Downtrend + High Volatility
3. Sideways Trend + Low Volatility
4. Sideways Trend + High Volatility
5. Uptrend + Low Volatility
6. Uptrend + High Volatility

This takes into account the two most central aspects that lead traders to different trading decisions. However, the market is not divided into too many small segments, which can lead to overfitting.

4.3 Recognizing Market Regimes

As described in [subsection 4.2](#), the market will be divided into six categories which are the result of combinations of two individual categories. This makes it possible to categorize the two individual categories individually, and finally merge them.

4.3.1 Recognizing Trend Behavior

The first step is to categorize the market into uptrends, downtrends, and sideways trends. Commonly a combination of a short-term moving average (e.g. SMA(50)), and a long-term moving average (e.g. SMA(200)) is used to identify superior trends. The SMA(200) is considered the classic boundary between bullish, and bearish market phases. If the short-term moving average is above the long-term moving average the market is considered bullish. Vice versa, if the short-term moving average is below the long-term moving average the market is considered bearish [\[19\]](#).

For the purpose of the current context, a modified combination of SMA(50), and SMA(100) was chosen. This decision is based on two considerations:

1. **Faster reaction:** The SMA(100) is intended to achieve faster reaction to medium-term trend changes without heavily weighting short-term volatility.
2. **Inertia of trend definition:** A shorter trend window, compared to the SMA(200) reduces the inertia of trend definition, which can be particularly advantageous for more refined classification into uptrends, downtrends, and sideways trend.

Additionally, a minimum slope threshold over the last 15 minutes was integrated for the SMA(50) to avoid that minimal direction changes are mistakenly interpreted as a meaningful trend. This short time span ensures that current market movements are adequately incorporated into the trend classification without being dominated by short-term noise (e.g., individual volatility peaks), and therefore increases the robustness of the trend recognition, and addresses the weaknesses of moving averages in sideways phases. A slope above $+0.05$ signals a significant short-term uptrend, while a slope below -0.05 suggests a clear downtrend. Values in between are interpreted as ambiguous, and are included in the sideways classification accordingly.

The market can therefore be divided into three trend phases based on the following criteria:

1. **Uptrend:** The SMA(50) is above the SMA(100), and the slope of the SMA(50) in the last 15 minutes is greater than 0.05.
2. **Downtrend:** The SMA(50) is below the SMA(100), and the slope of the SMA(50) in the last 15 minutes is less than -0.05.
3. **Sideways trend:** The market currently does not meet condition 1 or 2.

Figure 4 shows an example of trend classification of ETH/USDC M1 for 700 minutes.




Figure 4: Trend Classification

4.3.2 Recognizing Volatility

The second step is to categorize the market into phases with high, and low volatility. The volatility is calculated as the standard deviation of the logarithmic returns over the last 30 minutes [20]. This locally calculated volatility depicts short-term fluctuation intensity, and enables a context-dependent assessment of current market behavior.

To classify this local volatility, a comparison is made with the median of all available volatilities in the training dataset which was used for fitting the volatility classification indicator. If the current volatility is greater than the median the market is classified as highly volatile. Otherwise, the market is classified as low volatility.

This threshold definition is deliberately based on a dynamic, data-dependent approach rather than using a fixed absolute threshold. This automatically adapts the volatility classification to each specific market, and can therefore theoretically be applied to other financial markets.

Figure 5 shows the volatility classification on the same base data used in Figure 4.




Figure 5: Volatility Classification

4.3.3 Combining Trend and Volatility Classification

After the two separate classifications in [subsubsection 4.3.1](#) and [subsubsection 4.3.2](#) the results can be combined. The results of the combination can be seen in [Figure 6](#).




Figure 6: Regime Classification

While classification of the market regimes some discrepancies may occur in individual cases between the algorithmically determined category, and the visually perceived category. Such misclassifications are particularly possible during transition phases or in the case of short-term outliers. For example the downtrend, highly volatile market regime after 10:00 PM. The algorithm detected a downtrend. Visually an uptrend is perceived. At this time, the market is in a transition, which causes this misclassification.

Crucially, the classification distinguishes consistently, and meaningfully, which is the case when viewed as a whole. The robustness, and usefulness of the approach does not arise from absolute freedom from errors, but rather from the systematic reduction of uncertainty compared to a purely visual or subjective assessment.

4.4 Dividing Market Regimes in Durations

Based on the combination of trend direction, and volatility level, the duration of each market regime is also taken into account to enable a more differentiated classification. The regime duration is divided into three quantiles (33%, 66%, 100%) based on its distribution, allowing for classification into short-term, medium-term, and long-term phases.

The combination of these three dimensions results in a total of 18 different market regimes. This fine-grained segmentation makes it possible to more precisely capture, and differentiate different market conditions. For example, a short-term, volatile uptrend from a long-term, calm downtrend.

Even though this makes the market appear more fragmented, and regimes can change more frequently, this finer subdivision is analytically useful. It allows for context-sensitive market behavior to be examined, and differences in the dynamics of effectiveness of trading strategies in specific regime types to be systematically analyzed.

Instead of smoothing out reality with overly broad categories, a more nuanced picture is deliberately drawn that takes into account both the direction, intensity, and stability of market behavior. The resulting increased complexity is not a disadvantage, but rather a prerequisite for more meaningful, context-based analyses.

5 Money- and Risk-Management

Successful trading is not only based on a good strategy, but also on a disciplinant management of capital, and risk. Even the best prediction is useless if losses grow uncontrolled or a majority of the capital is risked by a few wrong decisions. This is where money, and risk management come into play. They define clear rules regarding how much to invest per trade, what level of risk is acceptable, and how losses can be limited. The goal is to protect capital over the long term, minimizing drawdowns, and profit from positive expected values in a controlled manner. This chapter introduces fundamental concepts, metrics, and methods helping to make rational, and sustainable decisions.

5.1 Calculating the Position Size

An important part of risk management is calculating the position size. It is helpful for maximizing the potential returns, as well as minimizing the financial risk. Many traders are only willing to risk 1% or 2% of the available capital per trade, to prevent a series of losing trades from decimating the available capital too much.

The distance between the estimated entry price, and the estimated stop loss price represents the maximum distance a market can move in an unprofitable direction before a position is automatically closed.

This allows the calculation of the position size to be carried out in three steps:

1. **Determining the risk per trade:** First, it must be determined how much of the available capital should be risked. If 1% of \$10,000 is to be risked, the maximum risk is \$100. It is important to note that the fraction of 1% does not have to be fixed. Thus, it is possible to risk more if the entry signal is very clear. If the entry signal is less clear, it can also be risked less. Only a fixed upper limit should be defined.
2. **Calculating the risk per unit:** To calculate the risk per share the absolute distance between the estimated entry price, and the estimated stop loss price must be calculated. This represents the risk per unit.
3. **Calculate the position size:** Dividing the risked capital by the risk per unit represents the number of units to buy or sell.

In total, these three steps can be combined into one formula [\[21\]](#):

$$PositionSize = \frac{AvailableBalance * RiskPerTrade}{RiskPerUnit} \quad (2)$$

So if the available account balance is \$10,000.00, the risk per trade is 1%, and the distance from the estimated entry price to the stop loss is \$5. The position size is calculated by:

$$PositionSize = \frac{\$10,000 * 1\%}{\$5} = \frac{\$100}{\$5} = 20[Units] \quad (3)$$

It is important to note that numbers can result with many decimal places, and many brokers only allow positions with a certain number of decimal places. If this is the case, the position size must be subsequently rounded to the maximum number of decimal places supported.

5.2 Validating an Entry Signal

Not every entry signal generated by a trading strategy is necessarily profitable. To be profitable in the long term, it is important to ensure that entry signals that are too risky or unrealistic in advance are filtered out, thus preventing positions from being opened. This chapter presents some techniques that can be used to validate entry signals.

5.2.1 Risk-Reward-Ratio

The risk reward ratio (RRR) is a fundamental key figure in trading. It describes the ratio between the potential profit (reward), and the potential loss (risk) of a single trade. The RRR helps in deciding whether an entry signal is too risky or not. It ensures that not only the hit rate determines the success of a strategy, but also the ratio of profit to loss in each individual trade.

The RRR is calculated by:

$$RRR = \frac{PossibleProfit}{PossibleLoss} = \frac{|OpenPrice - TakeProfitPrice|}{|OpenPrice - StopLossPrice|} \quad (4)$$

For example, if a long trade is opened at \$100, with a take profit at \$110, and a stop loss at \$95, the result is:

$$RRR = \frac{|\$100 - \$110|}{|\$100 - \$95|} = \frac{\$10}{\$5} = 2 \quad (5)$$

This means that for every dollar risked, a potential profit of two dollar is targeted.

A RRR greater 1 is generally considered positive because the expected profit is higher than the potential loss. However, the RRR should not be viewed in isolation. The essential factor is the combination of RRR, and hit rate:

1. **High RRR, low hit rate:** e.g. $RRR = 3$ with only a 30% probability of winning \Rightarrow potentially profitable.
2. **Low RRR, high hit rate:** e.g. $RRR = 0.5$ with an 80% hit rate \Rightarrow also potentially profitable.

The following rule of thumb clarifies when a strategy has a positive expected value in the long term:

$$ExpectedValue = PossibleProfit * HitRatio - PossibleLoss * (1 - HitRatio) \quad (6)$$

Only when this expected value is above zero a trading strategy is statistically profitable.

The RRR is not only a mathematical metric, but a central component of risk management. It helps traders systematically plan how much they are willing to lose per trade, relative to the expected profit. By consistently applying a minimum RRR (e.g. ≥ 1.5), many inefficient setups can be eliminated in advance [22].

5.2.2 Maximum Account Risk

(Nicht mehr als 10% des Riskieren in Summe)

5.2.3 Minimum Take Profit

In every trading strategy transaction fees play an important role. Especially in short-term trading transaction fees can turn seemingly profitable trades negative if they are not adequately considered. A common mistake is setting the take profit level too narrowly, resulting in a profit that is smaller than the costs incurred. To trade profitably, and sustainably, it is therefore essential that the take profit at least covers the fees incurred, but ideally, significantly higher.

5.3 Dealing with Trading Fees

As shown in [Table 1](#) ByBit charges two different fees named maker-, and taker-fee. The maker fee is charged when a limit order is placed in the order book, thereby creating liquidity. In contrast, the taker fee is charged when a market order is executed. This removes liquidity from the market. It is better for a broker if a market is as liquid as possible. Therefore, maker orders incur lower fees than taker orders.

Especially in higher-frequency trading, it is better to charge as few fees as possible. Therefore, it is better for a trader to execute as many limit orders as possible. Two orders are required for a complete trade: one for entry, and one for exit. But typically, three orders are placed (entry, stop-loss, take-profit), with either the take-profit or stop-loss order being executed.

If a market order is to be executed as an entry order, it is possible to convert it into a limit order by setting the order price slightly below (for long positions) or slightly above the current price (for short positions). The same procedure can be followed for the take-profit order.

However, this procedure should not be used for a stop-loss order. If this is converted to a limit position, there is no longer any guarantee that the stop order will be executed at all, as there is no longer any guarantee that the price will rise above or below the order level. This means that a position may remain open significantly longer than intended. If the price then continues to move in an unprofitable direction, the worst-case scenario is that the position is automatically closed by the broker because there is no longer any capital in the account.

The same problem can theoretically occur with a take-profit order. The difference here, is that the set stop market order still defines the maximum risk. Thus, while this particular trade may not be profitable, it is impossible to lose all of the capital. It is similar with the opening order. If it is not executed, the entire trade is not going to be executed. This means a missed potential profit, but there is no risk of losing capital.

If all orders are executed as planned, and the price moves in a profitable direction, the conversion of the orders will result in a reduction of fees by a factor of 2.75.

6 Deep Learning Models

This chapter presents the key components for developing and evaluating deep learning models. This includes both the selection of suitable models and the criteria for assessing their performance.

For this work, Keras version 3.8.0 was used as a high-level API for modeling and training neural networks [23]. Instead of the standard TensorFlow backend, PyTorch version 2.5.1 was used as the backend. A key reason for this decision was its better support on Windows, compared to TensorFlow [24], particularly with regard to installation and compatibility with existing CUDA drivers. By using Keras in combination with the PyTorch backend, user-friendly modeling could be combined with stable and well-supported execution on Windows systems.

For both regression and classification models, 13 different architectures were used, which can be divided into neural networks, convolutional neural networks, long short-term memories, and transformers.

6.1 Metrics

6.2 Optuna

In modern machine learning methods, the selection of hyperparameters plays a central role in model performance. Hyperparameters such as learning rate, the number of layers in a neural network, or regularization strengths directly influence the behavior and generalization of the model. The search for optimal values for these parameters, the so-called hyperparameter optimization, is often a time-consuming and computationally intensive process [25].

Optuna is a modern framework for automated hyperparameter optimization designed for efficiency, flexibility, and ease of use. It was developed to enable easy integration into existing machine learning pipelines while providing powerful, goal-oriented optimization.

The goal of Optuna is to automatically find those hyperparameter combinations that satisfy a specific optimization criterion (e.g., minimum validation error rate or maximum accuracy). The search process should be as efficient as possible, requiring as few model training sessions as possible [26].

In this work, Optuna was used to automatically run multiple training models with different hyperparameters to find the best hyperparameters. In the relevant sections of the next chapters, the range of the hyperparameters for each model will be mentioned.

6.3 Neural Networks

Neural networks are among the central methods of machine learning and form the basis of many modern deep learning models. The simplest and most widely used form is the feedforward neural network (FNN), also known as a multilayer perceptron (MLP). Such networks are universally applicable and, with appropriate structuring, can be used for classification, regression, and pattern recognition.

A classic neural network consists of several layers of artificial neurons that forward information in a fixed order from input to output. The structure can be divided as follows [27]:

1. **Input Layer:** Accepts the raw input data, e.g., a vector form of a time series or preprocessed features.
2. **Hidden Layers:** One or more layers of artificial neurons, each of which calculates a weighted sum of the inputs and further processes it using an activation function. Typical activation functions are ReLU, sigmoid, or tanh.
3. **Output Layer:** Provides the final result, e.g., class membership (for classification) or continuous value (for regression).

In this work, three different neural network models have been tested. The tested models are different in their complexity, depth, and methodical approach to data processing.

6.3.1 Base Model with Flatten-Architecture

The first model follows a classic feedforward approach, with one to three dense layers, each with 32 to 128 neurons and a ReLU activation function. Either before or after the dense layers, the input data are flattened. This decision influences whether the model processes all time points as a vector early on or treats each time step separately.

The learning rate of the Adam optimizer is varied between 1^{-5} and 1^{-2} . The number of historical time points used as input ranges from 5 to 150. This architecture was used as a simple baseline model to obtain a quick training and a first baseline model.

The learning rate of the Adam optimizer ranges between 10^{-5} and 10^{-2} . Also a flatten layer is built in the model, which is either directly after the input layer or directly before the output layer. The input layer has a length of 5 to 150. All hyperparameters are managed by Optuna.


```

1 num_layers = trial.suggest_int('num_layers', 1, 3)
2 num_units = trial.suggest_int('num_units', 32, 128)
3 learning_rate = trial.suggest_float('learning_rate', 1e-5, 1e
  -2, log=True)
4 input_length = trial.suggest_int('input_length', 5, 150)
5 flatten_before = trial.suggest_categorical("flatten_before",
  [True, False])
6
7 model = Sequential()
8
9 model.add(InputLayer(shape=(input_length, input_dimension)))
10
11 if flatten_before:
12     model.add(Flatten())
13
14 model.add(Dense(num_units, activation='relu'))
15 for _ in range(num_layers - 1):
16     model.add(Dense(num_units, activation='relu'))
17
18 if not flatten_before:
19     model.add(Flatten())
20
21 model.add(Dense(...))
22
23 model.compile(optimizer=Adam(learning_rate=learning_rate),
  ...)

```

Listing A: Base NN

6.3.2 Dropout Neural Network

In the second model, the data is first converted into a one-dimensional vector using a flatten layer. This is followed by one to three dense layers using ReLu activation. Batch normalization supports stable and rapid convergence, while dropout between 5% and 35% of the data is used as a regularization method to prevent overfitting.

The parameters tuned via Optuna control the model depth, the number of nodes per layer, the training dynamics, and the number of historical time points used as input.

The flatten layer is directly after the input layer, followed by a dense layer. After the dense layer a batch normalization and a dropout layer are following. After that, between one and three hidden layers. After the hidden layers one additional dropout layer is added. Each dropout layer has a dropout from 0.05 to 0.35. All other parameters are identical to those in the simple neural network.

```

1 num_layers = trial.suggest_int('num_layers', 1, 3)
2 num_units_1 = trial.suggest_int('num_units_1', 32, 256)
3 num_units_2 = trial.suggest_int('num_units_1', 32, 256)
4 dropout = trial.suggest_float('dropout', 0.05, 0.35)
5 learning_rate = trial.suggest_float('learning_rate', 1e-5, 1e
    -2, log=True)
6 input_length = trial.suggest_int('input_length', 5, 150)
7
8 model = Sequential()
9
10 model.add(InputLayer(shape=(input_length, input_dimension)))
11
12 model.add(Flatten())
13
14 model.add(Dense(num_units_1, activation='relu'))
15 model.add(BatchNormalization())
16 model.add(Dropout(dropout))
17 for _ in range(num_layers - 1):
18     model.add(Dense(num_units_2, activation='relu'))
19 model.add(Dropout(dropout))
20 model.add(Dense(...))
21
22 model.compile(optimizer=Adam(learning_rate=learning_rate),
    ...)

```

Listing B: Dropout NN

6.3.3 Residual Neural Network

The third model integrates a residual concept, originally known from image processing but also increasingly used in time series contexts. The architecture consists of a linear shortcut branch connected to the main path via an additive connection. This construction facilitates the training of deeper networks by improving gradient flow and preventing information loss through multiple layers.

In addition to the previous parameters, two additional variables controlled by Optuna are added: `num_units_res_prep` and `num_units_res`. These determine the complexity of the residual branch. The distinction whether flattening is done before or after the hidden layers is also included here to flexibly adapt data processing to the underlying data structure.

This architecture is particularly suitable for nonlinear and complex relationships, such as those frequently encountered in ETH price forecasting. The residual path allows the model to pass on basic information directly, while deeper layers learn abstract features.

```

1 num_layers = trial.suggest_int('num_layers', 1, 3)
2 num_units = trial.suggest_int('num_units', 32, 128)
3 num_units_res_prep = trial.suggest_int('num_units_res_prep',
    32, 128)
4 num_units_res = trial.suggest_int('num_units_res', 32, 128)
5 learning_rate = trial.suggest_float('learning_rate', 1e-5, 1e
    -2, log=True)
6 input_length = trial.suggest_int('input_length', 5, 150)
7 flatten_before = trial.suggest_categorical("flatten_before",
    [True, False])
8
9 input_layer = Input(shape=(input_length, input_dimension))
10
11 x = Flatten()(input_layer)
12
13 res = Dense(num_units_res_prep, activation="relu")(x)
14 res = Dense(num_units_res, activation="linear")(res)
15
16 x = Dense(num_units_res, activation="relu")(x)
17
18 x = Add()([x, res])
19
20 x = Dense(num_units, activation="relu")(x)
21
22 output = Dense(...)(x)
23
24 model = Model(input_layer, output)
25
26 model.compile(optimizer=Adam(learning_rate=learning_rate),
    ...)

```

Listing C: Residual NN

6.4 Long Short-Term Memory Models

In addition to traditional dense neural networks, recurrent neural networks (RNNs) based on the Long Short-Term Memory (LSTM) architecture were also used. LSTMs are specifically designed to capture temporal dependencies in sequential data and retain long-term information across multiple time steps. This is particularly important for financial data such as ETH prices, as current prices are often influenced by past developments.

LSTM cells have an internal memory structure and three control mechanisms (input, forget, and output gates) that allow the network to decide which information should be stored, overwritten, or passed on. This effectively mitigates the so-called vanishing gradient problem of traditional RNNs.

6.4.1 Base LSTM

This model consists of one or two stacked LSTM layers followed by a dense output layer. If only one LSTM layer is added, it returns the result directly. With multiple layers, the temporal sequence information is passed to the second LSTM layer.

Similar to feedforward neural networks, the number of layers, the number of neurons per layer, the learning rate, and the length of the input window used are managed via Optuna.

This model represents the basic form of a sequential predictor and is well suited for simple time series patterns.

```
1 num_layers = trial.suggest_int('num_layers', 1, 2)
2 num_units_input = trial.suggest_int('num_units_input', 32,
3                                     128)
4 num_units = trial.suggest_int('num_units', 32, 128)
5 learning_rate = trial.suggest_float('learning_rate', 1e-5, 1e
6                                     -2, log=True)
7 input_length = trial.suggest_int('input_length', 30, 150)
8
9 model = Sequential()
10 model.add(InputLayer(shape=(input_length, input_dimension)))
11 if num_layers == 1:
12     model.add(LSTM(num_units_input, return_sequences=False))
13 else:
14     model.add(LSTM(num_units_input, return_sequences=True))
15
16 for i in range(num_layers - 1):
17     if i == num_layers - 2:
18         model.add(LSTM(num_units, return_sequences=False))
19     else:
20         model.add(LSTM(num_units, return_sequences=True))
21
22 model.add(Dense(...))
23
24 model.compile(optimizer=Adam(learning_rate=learning_rate),
25               ...)
```

Listing D: Base LSTM

6.4.2 Dropout LSTM

This variant extends the standard model with dropout layers inserted between the LSTM units, discarding between 5% and 50% of the data.

The additional parameter is also determined by Optuna and allows fine-tuned control over the regularization strength. This model is particularly

useful when working with noisy or volatile price trends, such as cryptocurrencies.

```
1 num_layers = trial.suggest_int('num_layers', 1, 2)
2 num_units_input = trial.suggest_int('num_units_input', 32,
   128)
3 num_units = trial.suggest_int('num_units', 32, 128)
4 learning_rate = trial.suggest_float('learning_rate', 1e-5, 1e
   -2, log=True)
5 input_length = trial.suggest_int('input_length', 30, 150)
6 dropout = trial.suggest_float('dropout', 0.05, 0.5)
7
8 model = Sequential()
9 model.add(InputLayer(shape=(input_length, input_dimension)))
10 if num_layers == 1:
11     model.add(LSTM(num_units_input, return_sequences=False))
12 else:
13     model.add(LSTM(num_units_input, return_sequences=True))
14
15 model.add(Dropout(dropout))
16 for i in range(num_layers - 1):
17     if i == num_layers - 2:
18         model.add(LSTM(num_units, return_sequences=False))
19     else:
20         model.add(LSTM(num_units, return_sequences=True))
21     model.add(Dropout(dropout))
22
23 model.add(Dense(...))
24
25 model.compile(optimizer=Adam(learning_rate=learning_rate),
   ...)
```

Listing E: Dropout LSTM

6.4.3 Bidirectional LSTM (BiLSTM)

The third model uses bidirectional LSTM layers. These process the input sequence forward and backward simultaneously, allowing both earlier and later information to be incorporated into the internal state.

This architecture can be particularly valuable in the context of ETH predictions based on historical data, as it allows for better detection of symmetric patterns or turning points in price trends.

The hyperparameter structure is similar to the previous models, but the modeling is performed using double LSTM paths.

```
1 num_layers = trial.suggest_int('num_layers', 1, 2)
2 num_units_input = trial.suggest_int('num_units_input', 32,
   128)
```

```

3 num_units = trial.suggest_int('num_units', 32, 128)
4 learning_rate = trial.suggest_float('learning_rate', 1e-5, 1e
  -2, log=True)
5 input_length = trial.suggest_int('input_length', 30, 150)
6
7 model = Sequential()
8 model.add(InputLayer(shape=(input_length, input_dimension)))
9
10 if num_layers == 1:
11     model.add(Bidirectional(LSTM(num_units_input,
12     return_sequences=False)))
13 else:
14     model.add(Bidirectional(LSTM(num_units_input,
15     return_sequences=True)))
16
17 for i in range(num_layers - 1):
18     if i == num_layers - 2:
19         model.add(Bidirectional(LSTM(num_units,
20         return_sequences=False)))
21     else:
22         model.add(Bidirectional(LSTM(num_units,
23         return_sequences=True)))
24
25 model.add(Dense(...))
26
27 model.compile(optimizer=Adam(learning_rate=learning_rate),
28     ...)

```

Listing F: Bidirectional LSTM

6.4.4 Encode-Decode Model with Repeat-Vector

The fourth model is based on an encoder-decoder approach, as used in machine translation. The input sequence is first converted into a fixed context vector by an LSTM layer, which is then duplicated multiple times using RepeatVector. A decoder-LSTM interprets this vector sequentially to generate a temporally structured output.

This architecture is particularly suitable for multi-step forecasting, as it can map not only the next value but an entire sequence of future values. The final preprocessing is performed using TimeDistributed(Dense(...)) to calculate a separate output for each time step.

```

1 num_units_input = trial.suggest_int('num_units_input', 32,
  128)
2 num_units = trial.suggest_int('num_units', 32, 128)
3 learning_rate = trial.suggest_float('learning_rate', 1e-5, 1e
  -2, log=True)

```

```

4 input_length = trial.suggest_int('input_length', 30, 150)
5
6 model = Sequential()
7 model.add(InputLayer(shape=(input_length, input_dimension)))
8
9 model.add(LSTM(num_units_input, return_sequences=False))
10
11 model.add(RepeatVector(30))
12
13 model.add(LSTM(num_units, activation="relu", return_sequences
    =True))
14 model.add(TimeDistributed(Dense(...)))
15
16 model.compile(optimizer=Adam(learning_rate=learning_rate),
    ...)

```

Listing G: Encode-Decode LSTM

6.5 Convolutional Neural Networks

Convolutional neural networks (CNNs) represent a special architecture within deep learning that was developed specifically to process data with spatial structures, such as images. Compared to traditional feedforward networks, CNNs are significantly more efficient in handling high-dimensional inputs and have established themselves as standard methods in numerous application areas such as image recognition, object detection, medical image analysis, and time series analysis.

The key difference between CNNs and classical neural networks lies in the use of so-called convolutional layers, which exploit the local structures in the input data. Instead of connecting every neuron to all input values (as with a fully connected layer), local filters (kernels) are used that respond only to a small portion of the input space. These filters are automatically learned during the training process.

Time series typically exhibit temporal dependencies, repetitions, seasonal patterns, or short-term fluctuations. Classic methods for processing this data, such as LSTM networks, explicitly model such dependencies using recursive states. CNNs, on the other hand, use an alternative approach, called local convolutions, which also serve time series to detect relevant patterns or changes over short or medium periods of time.

The advantages of using CNNs for time series include:

1. **Parallel Processing:** CNNs do not require recursive loops, e.g., compared to LSTMs.

2. **Local Pattern Detection:** CNNs can detect local patterns, such as peaks, fluctuations, or trend changes.
3. **Computational Complexity:** CNNs require lower computational power compared to recursive architectures.

6.5.1 Base CNN

The simplest variant consists of two stacked one-dimensional convolutional layers, followed by maxpooling, and two dense layers for output preparation. The position of the flattening layer affects the processing order between the dense and the sequence structure.

The hyperparameters tuned by Optuna are the number of units in the convolutional layers, the length of the filters (Kernel size), the size of the maxpooling region, the position of the flattening layer, and the number of neurons in the dense layers.

```
1 num_units_cnn = trial.suggest_int('num_units_cnn', 32, 128)
2 num_units = trial.suggest_int('num_units', 32, 128)
3 learning_rate = trial.suggest_float('learning_rate', 1e-5, 1e
  -2, log=True)
4 input_length = trial.suggest_int('input_length', 5, 150)
5 flatten_before = trial.suggest_categorical("flatten_before",
  [True, False])
6 kernel_size = trial.suggest_int("kernel_size", 2, 5)
7 pool_size = trial.suggest_int("pool_size", 1, 5)
8
9 model = Sequential()
10 model.add(InputLayer(shape=(input_length, input_dimension)))
11 model.add(Conv1D(num_units_cnn, kernel_size=kernel_size,
  activation="relu"))
12 model.add(MaxPooling1D(pool_size=pool_size))
13 if flatten_before:
14     model.add(Flatten())
15 model.add(Dense(num_units, activation="relu"))
16 if not flatten_before:
17     model.add(Flatten())
18 model.add(Dense(...))
19
20 model.compile(optimizer=Adam(learning_rate=learning_rate),
  ...)
```

Listing H: Base CNN

6.5.2 Deep CNN

The model shown here represents an extension of a simple CNN model (Base CNN). While the simple model uses only a single 1D convolutional layer followed by a pooling and dense layer, this Deep CNN model uses two consecutive Conv1D layers. This additional depth allows the network to learn more complex and hierarchical feature representations from the input data. This is particularly useful for time series with multi-level dependencies or subtle patterns. The model is complemented by a MaxPooling layer, which reduces the temporal resolution and counteracts overfitting. The optional placement of the Flatten layer allows for the evaluation of different configurations of feature densification, which in turn increases model flexibility. Overall, this more deeply structured CNN enables more powerful modeling than the flattened variant.

```
1 num_units_cnn = trial.suggest_int('num_units_cnn', 32, 128)
2 num_units = trial.suggest_int('num_units', 32, 128)
3 learning_rate = trial.suggest_float('learning_rate', 1e-5, 1e
4   -2, log=True)
5 input_length = trial.suggest_int('input_length', 5, 150)
6 flatten_before = trial.suggest_categorical("flatten_before",
7   [True, False])
8
9 model = Sequential()
10 model.add(InputLayer(shape=(input_length, input_dimension)))
11 model.add(Conv1D(num_units_cnn, kernel_size=kernel_size,
12   activation="relu", padding="same"))
13 model.add(Conv1D(num_units_cnn, kernel_size=kernel_size,
14   activation="relu", padding="same"))
15 model.add(MaxPooling1D(pool_size=pool_size))
16 if flatten_before:
17     model.add(Flatten())
18 model.add(Dense(num_units, activation="relu"))
19 if not flatten_before:
20     model.add(Flatten())
21 model.add(Dense(...))
22
23 model.compile(optimizer=Adam(learning_rate=learning_rate),
24   ...)
```

Listing I: Deep CNN

6.5.3 Attention CNN

An alternative architecture uses multiple parallel convolutional layers with different kernel sizes. This allows the model to detect both short-term and medium-term patterns simultaneously. The resulting feature maps are combined using an attention layer to highlight important sequence segments.

This is followed by global average pooling and a dense layer for output generation.

```
1 num_units_cnn = trial.suggest_int('num_units_cnn', 32, 128)
2 num_units = trial.suggest_int('num_units', 32, 128)
3 learning_rate = trial.suggest_float('learning_rate', 1e-5, 1e
  -2, log=True)
4 input_length = trial.suggest_int('input_length', 5, 150)
5 kernel_size_1 = trial.suggest_int("kernel_size_1", 2, 9)
6 kernel_size_2 = trial.suggest_int("kernel_size_2", 2, 9)
7 kernel_size_3 = trial.suggest_int("kernel_size_3", 2, 9)
8
9 input_layer = Input(shape=(input_length, input_dimension))
10
11 conv_1 = Conv1D(num_units_cnn, kernel_size=kernel_size_1,
  padding="same", activation="relu")(input_layer)
12 conv_2 = Conv1D(num_units_cnn, kernel_size=kernel_size_2,
  padding="same", activation="relu")(input_layer)
13 conv_3 = Conv1D(num_units_cnn, kernel_size=kernel_size_3,
  padding="same", activation="relu")(input_layer)
14
15 concat = Attention()([conv_1, conv_2, conv_3])
16
17 gap = GlobalAveragePooling1D()(concat)
18 dense1 = Dense(num_units, activation="relu")(gap)
19
20 output_layer = Dense(...)(dense1)
21
22 model = Model(input_layer, output_layer)
23
24 model.compile(optimizer=Adam(learning_rate=learning_rate),
  ...)
```

Listing J: Attention CNN

6.5.4 Concatenation CNN

A slightly modified version replaces the attention module with a simple concatenation of the feature maps. This allows the model to pass the extracted features directly without performing weighting.

This variant is computationally less expensive than the attention-based one and is well-suited when all filter outputs should be treated equally.

```

1 num_units_cnn = trial.suggest_int('num_units_cnn', 32, 128)
2 num_units = trial.suggest_int('num_units', 32, 128)
3 learning_rate = trial.suggest_float('learning_rate', 1e-5, 1e
  -2, log=True)
4 input_length = trial.suggest_int('input_length', 5, 150)
5 kernel_size_1 = trial.suggest_int("kernel_size_1", 2, 9)
6 kernel_size_2 = trial.suggest_int("kernel_size_2", 2, 9)
7 kernel_size_3 = trial.suggest_int("kernel_size_3", 2, 9)
8 pool_size = trial.suggest_int("pool_size", 1, 5)
9
10 input_layer = Input(shape=(input_length, input_dimension))
11 conv_1 = Conv1D(num_units_cnn, kernel_size=kernel_size_1,
  padding="same", activation="relu")(input_layer)
12 conv_2 = Conv1D(num_units_cnn, kernel_size=kernel_size_2,
  padding="same", activation="relu")(input_layer)
13 conv_3 = Conv1D(num_units_cnn, kernel_size=kernel_size_3,
  padding="same", activation="relu")(input_layer)
14
15 concat = Concatenate()([conv_1, conv_2, conv_3])
16
17 gap = GlobalAveragePooling1D()(concat)
18 dense1 = Dense(num_units, activation="relu")(gap)
19
20 output_layer = Dense(...)(dense1)
21
22 model = Model(input_layer, output_layer)
23
24 model.compile(optimizer=Adam(learning_rate=learning_rate),
  ...)

```

Listing K: Concatenation CNN

6.5.5 CNN-GRU Hybrid Model

The Gated Recurrent Unit (GRU) model is a simplified variant of the LSTM and was developed to reduce computational effort and model complexity. Compared to LSTM, GRU has only two gates, an update gate and a reset gate, instead of three (input, forget, and output gates in LSTM). This makes GRU faster to train, requires less memory, and delivers comparable results in many tasks. Both models are capable of learning long-term dependencies in sequence data, but GRU is often the preferred choice when a good balance between efficiency and performance is required.

To combine the advantages of both architectures, a hybrid CNN-GRU model was also implemented. A Conv1D layer first extracts local features

from the sequence, which are then sequentially processed using a GRU module.

This model is particularly well-suited for linking local patterns (using CNN) with temporal dependencies (using GRU). The final dense layers perform the transformation to the target variable as usual.

```

1 num_units_cnn = trial.suggest_int('num_units_cnn', 32, 128)
2 num_units = trial.suggest_int('num_units', 32, 128)
3 num_units_gru = trial.suggest_int('num_units_gru', 32, 128)
4 learning_rate = trial.suggest_float('learning_rate', 1e-5, 1e
    -2, log=True)
5 input_length = trial.suggest_int('input_length', 5, 150)
6 kernel_size = trial.suggest_int("kernel_size", 2, 5)
7 pool_size = trial.suggest_int("pool_size", 1, 5)
8
9 model = Sequential()
10 model.add(InputLayer(shape=(input_length, input_dimension)))
11 model.add(Conv1D(num_units_cnn, kernel_size=kernel_size,
    activation="relu"))
12 model.add(MaxPooling1D(pool_size=pool_size))
13 model.add(GRU(num_units_gru, return_sequences=False))
14 model.add(Dense(num_units, activation="relu"))
15 model.add(Dense(...))
16
17 model.compile(optimizer=Adam(learning_rate=learning_rate),
    ...)

```

Listing L: CNN + GRU

6.6 Transformer Models

In addition to recurrent architectures such as LSTMs, a model based on the Transformer architecture has also been implemented. Transformers are considered a powerful alternative to RNNs, especially in tasks that require the detection of long dependencies within the input sequence. Originally introduced in natural language processing, Transformers have also proven extremely effective in time series analysis.

In contrast to RNNs, transformers process all positions of a sequence simultaneously, i.e., in parallel. Instead of sequential propagation, the transformer uses a mechanism called self-attention to determine which parts of a sequence are particularly relevant to the current position. This allows the transformer to model long-term dependencies more efficiently, without information becoming diluted as the sequence length increases.

A transformer encoder typically consists of the following components:

1. **Multi-head self-attention:** Multiple parallel attention heads allow the model to consider different aspects of the input sequence simultaneously.
2. **Residual connections & layer normalization:** The original inputs are added to the output of the attention and feedforward steps, stabilized by normalization.
3. **Feedforward network:** A small, position-independent MLP layer for applying nonlinear transformations to the position embeddings.
4. **Dropout layers:** These are used for regularization to avoid overfitting.

6.6.1 Transformer Architecture

The implemented model is based on a stacked transformer encoder consisting of two encoder blocks. The architecture is as follows:

1. **Input:** A sequential input with variable length and dimensionality corresponding to the number of features.
2. **Dual transformer encoder:** Each block uses multi-head attention with variable number of heads, a feedforward layer, and normalization and dropout layers.
3. **Global average pooling:** Reduces the sequence to a fixed vector by calculating the mean across all time points.
4. **Dense layers:** A fully connected layer with ReLU activation and dropout for further regularization.
5. **Output:** A linear dense layer that provides the prediction of the target variable.

The most important hyperparameters managed by Optuna are the number of parallel attention heads, the dimension of the key, query, and value vectors in the attention, the number of neurons in the feedforward network, the dropout rate, as well as the learning rate.

The method `transformer_encoder(...)` implements a basic encoder block:

```

1 def transformer_encoder(inputs, head_size, num_heads, ff_dim,
2   dropout=0.):
    x = MultiHeadAttention(key_dim=head_size, num_heads=
      num_heads, dropout=dropout)(inputs, inputs)

```

```

3     x = Dropout(dropout)(x)
4     x = LayerNormalization(epsilon=1e-6)(x)
5     res = x + inputs
6
7     x = Conv1D(filters=ff_dim, kernel_size=1, activation='
relu')(res)
8     x = Dropout(dropout)(x)
9     x = Conv1D(filters=inputs.shape[-1], kernel_size=1)(x)
10    x = LayerNormalization(epsilon=1e-6)(x)
11    return x + res
12
13
14    head_size = trial.suggest_int('head_size', 32, 128)
15    num_heads = trial.suggest_int('num_heads', 2, 6)
16    learning_rate = trial.suggest_float('learning_rate', 1e-5, 1e
-2, log=True)
17    input_length = trial.suggest_int('input_length', 30, 150)
18    ff_dim = trial.suggest_int("ff_dim", 64, 256)
19    dropout = trial.suggest_float("dropout", 0.05, 0.3)
20
21    inputs = Input(shape=(input_length, input_dimension))
22
23    x = self.transformer_encoder(inputs, head_size=head_size,
num_heads=num_heads, ff_dim=ff_dim, dropout=dropout)
24    x = self.transformer_encoder(x, head_size=head_size,
num_heads=num_heads, ff_dim=ff_dim, dropout=dropout)
25
26    x = GlobalAveragePooling1D()(x)
27    x = Dense(128, activation="relu")(x)
28    x = Dropout(0.2)(x)
29    outputs = Dense(...)(x)
30
31    model = Model(inputs, outputs)
32    model.compile(optimizer=Adam(learning_rate=learning_rate),
...)

```

Listing M: Transformer Model

The usage of a one-dimensional convolutional layer instead of a feedforward neural network is a common practice in time-series-based Transformer implementations and enables local processing along the time axis.

In contrast to LSTM models, which work sequentially and propagate their internal states step by step, a transformer model can analyze all time points in parallel. This offers advantages in the modeling of long-term dependencies (e.g., trends or cyclical patterns). However, transformers are also more computationally intensive, especially for long sequences and large headcounts, since self-attention scales quadratically with the sequence length.

6.7 Regression-Models

6.7.1 Loss-Function

In order to incorporate the trading context during model training and evaluation, traditional loss functions such as the root mean squared error have been avoided. Instead, a proprietary loss function was implemented that calculates realized profit and loss. The value of the loss function decreases the more profit is generated.

As mentioned above the output of the regression models are either a sequence of the next expected logarithmic returns by minute. The first step of the loss function is to decide whether the opened position should be a long or short position. This is decided by cumulating the predictions. The absolute maximum is then determined for the cumulative values greater than and less than zero. If the absolute maximum of the values greater than zero is greater than the absolute maximum of the values less than zero, a long position is opened. Otherwise, a short position is opened. [Figure 7](#) illustrates the decision using a long position as an example.




Figure 7: Long Position Decision

After the direction decision, the stop loss and take profit levels must be determined. For long and short positions, there are two possibilities at which specific level the stop loss and take profit are set, based on the global high and low of the cumulative predicted logarithmic returns. If a long position was previously decided and the global high comes after the global low, the stop

loss is set at the global low and the take profit at the global high. However, if the global high comes before the global low, the stop loss is set at the lowest low before the global high. For short positions, the decision is exactly the opposite.

Figure 8 shows the stop loss and take profit levels for all four cases. The red areas mark the predicted unprofitable zones and the green areas mark the predicted profitable zones.




Figure 8: Long Position Decision

After the stop loss and take profit levels are determined, the actual logarithmic returns are used to calculate the profit or loss for each position. If neither the stop loss or the take profit is reached, the profit calculation defaults to the profit if the stop loss is reached to ensure that the model predicts the correct take profit more often.

6.8 Classification-Models

6.9 Model Evaluation

7 Trading-Strategies

7.1 AI-Trading-Strategies

7.2 Classic Trading-Strategies

7.3 Performance Comparison

8 Trading-Engine

For backtesting, and live execution of trading strategies, a Java framework was built. This chapter describes the use cases as well as the most important components of the trading engine.

8.1 Use-Cases of the Trading-Engine

The developed trading engine represents a flexible, and modular software platform that covers various use cases in the field of algorithmic trading. The focus is on combining usability, adaptability, and performance. The following describes the most important use cases, and features:

1. **Backtesting trading strategies:** The trading engine enables the simulated execution of trading strategies on historical market data. This allows strategies to be tested under realistic conditions, and evaluate their performance, robustness, and risk characteristics. By replicating historical market conditions, incorrect decisions can be identified early, and the trading strategy can be adjusted without real capital.
2. **Live Execution in Real-Time Operation:** In addition to backtesting, the engine also supports real-time execution of strategies in the market. By an interface to brokers, the engine can receive market data in real time, make trading decisions, and execute orders directly. This enables the engine to be used as the basis for automated trading in production environments.
3. **Using money, and risk management strategies:** An important use case is the implementation of different money, and risk management approaches. Users can integrate various strategies for position sizing, stop-loss setting, or profit-taking, and examine their impact on a strategies performance in backtests or in live operation. This supports the development of more stable, and profitable trading approaches.
4. **Connecting to any broker:** The trading engine is designed to connect to various brokers via interchangeable interfaces. This allows the engine to be combined with different trading systems, and platforms. This flexibility enables its use in a wide variety of markets, and infrastructures.
5. **Development, and integration of custom trading strategies:** A key feature of the engine is the ability for users to implement their own

trading strategies. A clear separation between core functionality, and strategy logic allows for flexible integration, and testing of individual algorithms.

Most components of the trading engine are interchangeable, and can be implemented by the user. Only the core of the framework, the basic process, and control logic, are fixed. This allows users to tailor the engine to their individual needs, develop their own modules for data feeds, order management, strategy, or risk control, and thus ensure a high degree of flexibility in the use, and further development of the platform.

8.2 Architecture

The trading engine consists of many modules which can be clustered in three main categories:

1. Core Modules (Trading Engine)
2. Applications
3. Adapters

Figure 9 shows the components of the trading engine.




Figure 9: Trading-Engine Components

8.2.1 Plugin Architecture

To achieve the flexibility mentioned in [subsection 8.1](#), a plugin architecture was used. Within the core, interfaces have been defined that define which external functionalities are required (Service provider interface, or SPI). To enable easy interchangeability by users, the implementations of the interfaces in the external modules (Adapters) are loaded by the Java `java.util.ServiceLoader`.

The `ServiceLoader` loads classes at runtime that implement a specific interface or abstract class. The classes to be loaded are specified via a configuration file in the `resources/META-INF/services` directory, which the `ServiceLoader` searches for in every JAR file in the classpath. If a corresponding configuration file is found, the `ServiceLoader` can create an instance of the respective class using the default constructor [28].

As an example, an service provider interface for read access can be defined.

```
1 // Contained in module spi
2 package com.example.spi.read;
3
4 public interface ReadRepository{
5     public DomainObject findAll();
6 }
```

Listing N: SPI Definition

This interface can be implemented in multiple modules. One implementation can be used to read data from files:

```
1 // Contained in module file-adapter
2 package com.example.adapter.file.read;
3
4 public class FileReadRepository implements ReadRepository {
5
6     @Override
7     public DomainObject findAll(){
8         // Logic to read a file
9     }
10 }
```

Listing O: File-Repository Implementation

Here, a configuration file named `com.example.spi.read.ReadRepository` is located in the directory `resources/META-INF/services`. This file contains the fully qualified class name of the implementation class (`com.example.adapter.file.read.FileReadRepository`).

Another implementation can be used to read data from databases:

```
1 // Contained in module db-adapter
```

```

2 package com.example.adapter.database.read;
3
4 public class DatabaseReadRepository implements
    ReadRepository {
5
6     @Override
7     public DomainObject findAll(){
8         // Logic to read a database
9     }
10 }

```

Listing P: Database-Repository Implementation

Similar to the `FileReadRepository` a file named `com.example.spi.read.ReadRepository` is located in the directory `resources/META-INF/services`. However, this file contains `com.example.adapter.database.read.DatabaseReadRepository`.

The `ServiceLoader` searched each JAR in the classpath for files under `resources/META-INF/services` with the name `com.example.spi.read.ReadRepository` that contains the implementation of this interface, and creates them.

Since an external application has the ability to define the classpath, it can add or remove specific JARs. This allows it to determine which implementations to use. For example, it can decide to add only the `file-adapter` module or, if necessary, the `db-adapter` module if loading from a file, and a database simultaneously.

8.2.2 Core Modules

The core of the framework consists of several loosely coupled components that communicate with each other through well-defined APIs. The most central unit is the **Strategy-Engine** which contains the main logic for execution of trading strategies, and controls the cooperation of the other subsystems.

The following components form the core:

1. **Data-Provider:** Obtains market data by a configurable `Datasource-API`. The specific data source (local or remote) is integrated via adapters.
2. **Trading-Strategy:** The actual trading logic is integrated via the `Strategies-API`. This API is separate from the framework, and can be extended as required.
3. **Broker-Requestor & Trade-Executor:** Both communicate with brokers via a generic `Broker-Connector-API`, and enable both order

placement, and broker status queries, such as the current account balance or all open positions.

4. **Entry-Signal-Adaptor & Entry-Signal-Validator:** These two components process, and validate generated entry signals. They are linked to their own **Management-API**, and enable additional risk checks, such as capital requirements or position limits.

Before the **Strategy-Engine** there are two other engines, named **Remote-Backtest-Engine**, and **Backtest-Engine**, which configure the **Strategy-Engine**. The difference between the engines is that the **Remote-Backtest-Engine** defines the **Strategy-Engine** asynchronously. This means that a remote data source can stream market data via an API. When new data arrives, the **Strategy-Engine** is notified by the data source, and the **Strategy-Engine** starts executing the trading strategy. For a local backtest, the **Strategy-Engine** is configured so that data synchronization is handled by the backtest engine, not by the data source.

Abstract apps are modules that accept, and parse user configurations. For example, the task of the **Abstract-Local-Backtest-App** is to ask the user via the console which strategy should be used for the backtest. With the **Abstract-Remote-Execution-App**, the configuration is not done via the console, but primarily via environment variables. The parsed configurations are then passed to the respective engines.

8.2.3 Applications

The application layer defines specific execution environments, by defining the classpath with adapters.

1. **Local-Backtest-App:** Executes trading strategies offline by using the **Local-CSV-Datasource**, and the **Demo-Broker-Connector**.
2. **Remote-Execution-App:** Executes strategies in real time, and communicates with live broker interfaces.

8.3 Demo Broker

For local backtests the connection to a real broker must be mocked. Therefore the **Demo-Broker-Connector** is used to replace a real broker. He takes over the main tasks including:

1. **Order execution and position management:** Executing market, limit, take-profit, and stop-loss orders, including execution fee calculation.
2. **Adapt trailing stop positions:** Monitor trailing stop positions, and adapt the stop-level according to the most recent price.
3. **Account balance management:** Manage the available account balance as well as the margin balance.
4. **Store executed trades:** Storing the executed trades is essential for risk management, and later analysis.

8.3.1 Order Execution and Position Management

In context of trading order execution and position management are closely connected, because positions can be translated in three orders where only two of them are actually executed. As described in [subsection 5.3](#) it is necessary that the trading engine supports market and limit orders. In real life environments the most orders are not executed at the expected price, due to slippage [29]. To simulate slippage, the trading engine can either use historical data of the underlying cryptocurrency, which is available in the lowest available timeframe (usually seconds or tick data), or simulate slippage randomly. Two problems exist with simulations using historical data. The first is that data in the seconds or tick range over a long period of time is difficult for private individuals to obtain. Furthermore, a type of latency must be simulated, which simulates the processing time and network traffic in live operation. With random simulations, the problem arises that the backtest results are no longer deterministic, thus a subsequent comparison of different strategies is subject to error. For these reasons, the slippage factor is not taken into account during order execution. This means that market orders are always executed at the most current available price and limit orders are always executed at the specified order level.

[Figure 10](#) shows the process of opening a single position. The margin in euro is calculated using the following formula. *EuroConversion* is the current price for converting the current counter currency in euro. For example, for ETH/USDC the *EuroConversion* is the current price of EUR/USDC:

$$Margin = \frac{PositionQuantity}{Leverage * EuroConversion} = \frac{PositionSize * OpenPrice}{Leverage * EuroConversion} \quad (7)$$




Figure 10: Opening a Single Position

Figure 11 shows the process of closing a single position. The profit in euro is calculated using the following formula:

$$Profit = \frac{(ClosePrice - OpenPrice) * PositionSize}{EuroConversion} * \begin{cases} -1, & \text{if Sell-Position} \\ 1, & \text{otherwise} \end{cases}$$




Figure 11: Closing a Position

Figure 12 shows the process of opening a position taking into account

that other open positions can already exist.

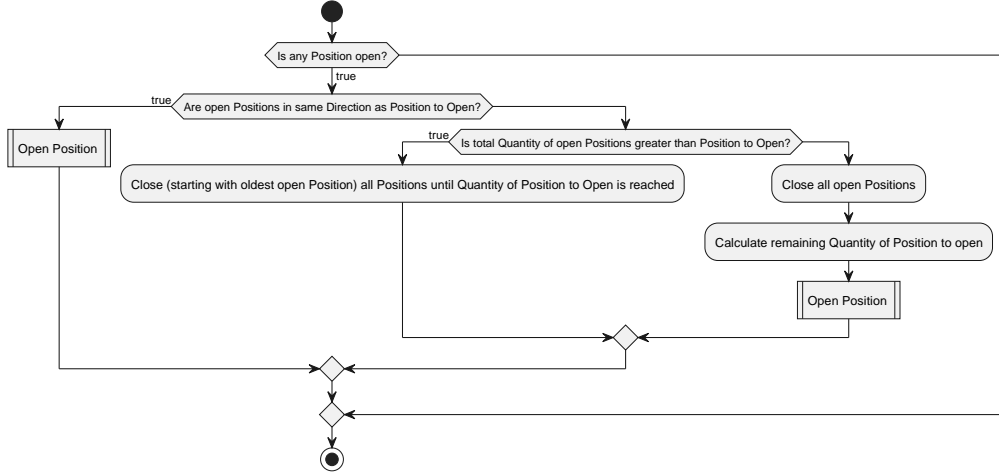


Figure 12: Opening a Position

8.3.2 Adapting Trailing Stop Positions

Trailing stop positions are positions with a dynamic stop order. The stop level moves automatically with the price as soon as it moves in the desired direction. For a buy position the stop level moves if the price rises. If the price falls, the stop remains unchanged [30].

Figure 13 shows a synthetic closing price (black) and a trailing-stop order with 60 points distance (red) for a buy position. In the green filled areas the price moves in the desired direction, so the trailing-stop level follows the price. In the red filled areas the prices does not move in the desired direction, so the trailing-stop level does not move. An exception is the red filled area after $t = 3$, where the price moves in the desired direction, but the trailing stop does not. This is because there is not yet a 60-point gap between the price and the stop level. Therefore, the stop only moves again once the 60-point gap is restored.




Figure 13: Trailing Stop Example

9 Backtesting Trading-Strategies

9.1 Executing Backtests

9.2 Evaluating Backtests

9.3 Fee-Impacts

10 Stress Testing

11 Live-Testing

11.1 Broker Setup

11.2 Connecting to ByBit

11.3 Live-Test Results

12 Conclusion

12.1 Key Findings

13 Aim of further Works

14 References

- [1] Investing.com - Die besten Krypto Broker – Juni 2025. 2025. URL: <https://de.investing.com/brokers/cryptocurrency-brokers/#:~:text=Zu%20den%20wichtigsten%20Aufgaben%20eines,an%20Kryptow%C3%A4hrungen%20als%20einzelne%20B%C3%B6rsen>. (visited on 06/18/2025).
- [2] ByBit.com. 2025. URL: <https://www.bybit.com/en/> (visited on 06/18/2025).
- [3] ByBit API Doc. 2025. URL: <https://www.bybit.com/en/%20https://bybit-exchange.github.io/docs/v5/intro> (visited on 06/18/2025).
- [4] IG.com. 2025. URL: https://www.ig.com/de/kryptowaehrungshandel?_gl=1*15gtia3*_up*MQ..*_gs*MQ..&gclid=CjwKCAjwpMTCBhA-EiwA_-MsmST2-xxZ7Kxhnw1PDrrdvL7kdRSPmknOMOR2x2ShxY2-5F5G5_mIFBoCUikQAvD_BwE&gclsrc=aw.ds&gbraid=0AAAAADzDEsR-vDC7urfY10DmU7i3x7bQU (visited on 06/18/2025).
- [5] IG API Doc. 2025. URL: <https://labs.ig.com/getting-started.html> (visited on 06/18/2025).
- [6] Capital.com. 2025. URL: <https://capital.com/de-de> (visited on 06/18/2025).
- [7] Capital API Doc. 2025. URL: <https://open-api.capital.com/#section/General-information> (visited on 06/18/2025).
- [8] ByBit API Doc - Get Kline. 2025. URL: <https://bybit-exchange.github.io/docs/v5/market/kline> (visited on 06/18/2025).
- [9] Investiopia.com - Trend Trading: The 4 Most Common Indicators. 2025. URL: <https://www.investopedia.com/articles/active-trading/041814/four-most-commonlyused-indicators-trend-trading.asp> (visited on 06/19/2025).
- [10] Investiopia.com - What is EMA? How to Use Exponential Moving Average With Formula. 2024. URL: <https://www.investopedia.com/terms/e/ema.asp> (visited on 06/19/2025).
- [11] Investiopia.com - What Is MACD? 2024. URL: <https://www.investopedia.com/terms/m/macd.asp> (visited on 06/19/2025).

- [12] Investiopia.com - Average True Range (ATR) Formula, What It Means, and How to Use It. 2025. URL: <https://www.investopedia.com/terms/a/atr.asp> (visited on 06/19/2025).
- [13] Investiopia.com - Bollinger Bands: What They Are, and What They Tell Investors. 2024. URL: <https://www.investopedia.com/terms/b/bollingerbands.asp> (visited on 06/19/2025).
- [14] Investiopia.com - What Is a Momentum Indicator? Definition and Common Indicators. 2025. URL: <https://www.investopedia.com/investing/momentum-and-relative-strength-index/> (visited on 06/19/2025).
- [15] Investiopia.com - Relative Strength Index (RSI) Indicator Explained With Formula. 2024. URL: <https://www.investopedia.com/terms/r/rsi.asp> (visited on 06/19/2025).
- [16] Wikipedia.org - Principal component analysis. 2024. URL: https://en.wikipedia.org/w/index.php?title=Principal_component_analysis&oldid=1296008542 (visited on 06/20/2025).
- [17] Wikipedia.org - Curse of dimensionality. 2024. URL: https://en.wikipedia.org/w/index.php?title=Curse_of_dimensionality&oldid=1296421397 (visited on 06/20/2025).
- [18] Macrosynergy - Classifying market regimes. 2025. URL: <https://macrosynergy.com/research/classifying-market-regimes/> (visited on 06/24/2025).
- [19] IG.com - Moving averages: a guide to trend trading. 2025. URL: <https://www.ig.com/en/trading-strategies/moving-averages--a-guide-to-trend-trading-241031> (visited on 06/25/2025).
- [20] wikipedia.org - Volatility (finance). 2025. URL: [https://en.wikipedia.org/w/index.php?title=Volatility_\(finance\)&oldid=1291878143](https://en.wikipedia.org/w/index.php?title=Volatility_(finance)&oldid=1291878143) (visited on 06/25/2025).
- [21] BritannicaMoney - Calculating position size in trading: The key to risk management. 2024. URL: <https://www.britannica.com/money/calculating-position-size> (visited on 06/23/2025).
- [22] bitpanda - Risk reward ratio: what is it & why is it important? 2024. URL: <https://www.bitpanda.com/academy/en/lessons/risk-reward-ratio-what-is-it-and-why-is-it-important/> (visited on 06/23/2025).
- [23] keras.io. 2025. URL: <https://keras.io/> (visited on 07/06/2025).

- [24] tensorflow.org - Windows Install. 2025. URL: <https://www.tensorflow.org/install/pip#windows-native> (visited on 07/06/2025).
- [25] Jasman Pardede and Khairul Rijal. “The Effect of Hyperparameters on Faster R-CNN in Face Recognition Systems”. In: *Jurnal RESTI (Rekayasa Sistem dan Teknologi Informasi)* 9 (May 2025), pp. 436–448. DOI: [10.29207/resti.v9i3.6405](https://doi.org/10.29207/resti.v9i3.6405).
- [26] Yosua Harmoni, Kartika Maulida Hindrayani, and Dwi Prasetya. “Optimizing Categorical Boosting Model with Optuna for Anti-Tuberculosis Drugs Classification”. In: *Indonesian Journal of Electronics, Electromedical Engineering, and Medical Informatics* 7 (May 2025), pp. 401–414. DOI: [10.35882/ijeemi.v7i2.92](https://doi.org/10.35882/ijeemi.v7i2.92).
- [27] Saleem Abdullah, Ihsan Ullah, and Fazal Ghani. “Heterogeneous wireless network selection using feed forward double hierarchy linguistic neural network”. In: *Artificial Intelligence Review* 57 (July 2024). DOI: [10.1007/s10462-024-10826-y](https://doi.org/10.1007/s10462-024-10826-y).
- [28] oracle.com - ServiceLoader (Java SE 17 & JDK 17). 2025. URL: <https://docs.oracle.com/javase/8/docs/api/java/util/ServiceLoader.html> (visited on 06/26/2025).
- [29] ig.com - Slippage definition. 2025. URL: <https://www.ig.com/uk/glossary-trading-terms/slippage-definition> (visited on 06/26/2025).
- [30] ig.com - Trailing stop orders definition. 2025. URL: <https://www.ig.com/uk/glossary-trading-terms/trailing-stops-definition> (visited on 06/26/2025).