

Bachelor's Thesis

in the Bachelor's Program 'Applied Mathematics and Computer
Science'

Backtesting and Live-Testing of Classic and AI-Powered Trading Strategies

by

Jason Becker

Enrollment Number: 3567285

Supervisor: Prof. Dr. Bodo Kraft
Co-Supervisor: Hendrik Karwanni, M. Sc.
Submitted on: 7th August, 2025

Declaration of Authorship

This bachelor thesis has been written by:

Jason Becker
Wachtelstraße 24
40789 Monheim am Rhein

I hereby declare that this bachelor thesis has been written independently and without unauthorized assistance. I confirm that I have used only the sources and aids indicated in the bibliography. Furthermore, I declare that I have marked all passages that are either verbatim or paraphrased versions of sources.

I am aware that any form of plagiarism or dishonesty in academic work constitutes a serious offense and may lead to disciplinary measures.

Monheim am Rhein, 7th August, 2025

.....
Jason Becker

Abstract

This bachelor's thesis examines the effectiveness of classical and AI-based algorithmic trading strategies using the ETH/USDC cryptocurrency pair on a per-minute basis (M1). The goal is to improve informed trading decisions through data-driven methods. The special feature of this currency pair lies in the high liquidity and volatility of ETH and the stability of USDC, which brings with it both advantages and disadvantages.

Therefore, a complete workflow is presented, from data acquisition via a Broker-API to feature engineering, the identification of market regimes, and the development and training of various deep learning models (FNN, LSTM, CNN, hybrid models), all the way to the implementation and evaluation. In addition to classic technical indicators such as moving averages and Bollinger Bands, AI-based regression and classification strategies were also developed and compared.

After retrieving the market data via the broker API, systematic feature engineering followed, using not only traditional price- and volume-based indicators but also technical oscillators such as the RSI and MACD. Principal component analysis (PCA) was also applied to reduce the dimensionality of the data and eliminate redundant information.

The Optuna framework, based on Bayesian optimization and enabling efficient hyperparameter search, was used to optimize the model parameters. A dedicated loss function was implemented for the regression models, which takes the practical application of the models in a trading context into account during training.

A particular focus of this work is the technical implementation of a modular trading engine. The architecture of this engine allows dynamic import and exchange of trading strategies, efficient processing of historical data, and both backtesting and live trading with real broker integration. The implementation includes essential components such as an order execution module or data source modules that enable live-streaming of data via a broker or local loading of data, for example via CSV files.

To ensure realistic evaluation, the developed strategies were first optimized through backtests on historical data. After identifying the best parameters, the best strategies were evaluated through out-of-sample tests, and one strategy was then deployed in a controlled live test with a real broker under real market conditions. During both development and live testing, various risk and money management concepts were used to ensure capital preservation and controlled loss limitation.

The results show that both classical and AI-based approaches have their strengths. However, it turned out that the various strategies were severely affected by overfitting. This was demonstrated by the fact that the strategies achieved profitable results on the data used for parameter selection, but failed on previously unused test data. This highlights the challenge of developing robust, generalizable strategies. Although none of the tested strategies were profitable in the long term, this work provides valuable insights into the connection between AI and algorithmic trading, the importance of risk control, and the implementation of a fully automated trading system with direct connectivity to real brokers.

Contents

1	Introduction	1
1.1	Motivation	1
1.2	Aim of this Paper	1
1.3	Structure of this Paper	2
1.4	Ethereum	3
1.5	USD Coin	4
2	Data Source and Broker Selection	5
2.1	Broker Selection	5
2.2	API Connection and Data Retrieval Process	6
3	Market Regimes	7
3.1	Introduction to Market Regimes	7
3.2	How to Categorize the Market?	7
3.3	Recognizing Market Regimes	8
3.3.1	Recognizing Trend Behavior	8
3.3.2	Recognizing Volatility	10
3.3.3	Combining Trend and Volatility Classification	10
3.4	Dividing Market Regimes in Durations	11
4	Money- and Risk-Management	13
4.1	Calculating the Position Size	13
4.2	Validating an Entry Signal	14
4.2.1	Risk-Reward Ratio	14
4.2.2	Maximum Account Risk	15
4.2.3	Minimum Take-Profit	15
4.3	Dealing with Trading Fees	15
5	Performance Comparison	17
5.1	Equity Curve	17
5.2	Maximum Drawdown	17
5.3	Maximum Gain and Maximum Loss	18
5.4	Standard Deviation of Profits	19
5.5	Win Rate	20
6	Exploratory Data Analysis	21
6.1	Dividing the Data	21
6.2	Using Log>Returns	23
6.3	Additional Features	25
6.3.1	Trend Following Indicators	25
6.3.2	Volatility Indicators	26
6.3.3	Momentum Indicators	27
6.3.4	Price Transformation Indicators	27

6.4	Further Transformations	28
6.4.1	Scaling the Data	28
6.4.2	Principal Component Analysis	28
7	Deep Learning Models	30
7.1	Optuna	30
7.2	Neural Networks	31
7.2.1	Base Model with Flatten-Architecture	32
7.2.2	Dropout Neural Network	32
7.2.3	Residual Neural Network	33
7.3	Long Short-Term Memory Models	34
7.3.1	Base LSTM	34
7.3.2	Dropout LSTM	35
7.3.3	Bidirectional LSTM (BiLSTM)	36
7.3.4	Encode-Decode Model with Repeat-Vector	37
7.4	Convolutional Neural Networks	37
7.4.1	Base CNN	38
7.4.2	Deep CNN	39
7.4.3	Attention CNN	40
7.4.4	Concatenation CNN	40
7.4.5	CNN-GRU Hybrid Model	41
7.5	Model Training	42
7.6	Regression Models	43
7.6.1	Loss-Function	43
7.6.2	Metrics	45
7.6.3	Model Evaluation	46
7.7	Classification-Models	47
7.7.1	Loss-Function	47
7.7.2	Model Evaluation	47
8	Trading Strategies	49
8.1	AI Trading Strategies	49
8.1.1	Regression AI Strategy	49
8.1.2	Classification AI Strategy	50
8.2	Classic Trading Strategies	51
8.2.1	Dual Simple Moving Average Strategy	51
8.2.2	Triple Exponential Moving Average Strategy	53
8.2.3	Bollinger Bands Strategy	55
9	Trading-Engine	57
9.1	Use-Cases of the Trading-Engine	57
9.2	Architecture	58
9.2.1	Plugin Architecture	58
9.2.2	Core Modules	60

9.2.3	Applications	61
9.3	Demo Broker	61
9.3.1	Order Execution and Position Management	61
9.3.2	Adapting Trailing Stop Positions	64
10	Backtesting Trading Strategies	66
10.1	Trading Strategy Parameter Selection	66
10.2	Executing Backtests	66
10.3	Evaluating Backtests	67
10.4	Fee-Impacts	68
11	Live-Testing	70
11.1	Broker Setup	70
11.2	Live-Test Results	70
12	Dealing with Strategy-Overfitting	73
13	Conclusion	76
14	Aim of further Works	77
15	References	79
16	Appendix	86
16.1	Strategy Results on Validation Data	86
16.2	Out-of-Sample Strategy Results	88
16.3	Best Strategy Parameters	89
16.3.1	Classification AI Strategy Backtest Results	89
16.3.2	Dual Simple Moving Average Strategy Backtest Results	95
16.3.3	Triple Exponential Moving Average Strategy Backtest Results	99
16.3.4	Bollinger Bands Strategy Backtest Results	103
16.4	Other Tested Strategies	106
16.4.1	Other Tested Strategies Descriptions	106
16.4.2	Other Tested Strategies Results	107

List of Figures

1	Trend Classification	9
2	Volatility Classification	10
3	Regime Classification	11
4	Equity Curve Comparison	17
5	Drawdown Comparison	18
6	Price Fluctuation of ETH in USDC	24
7	Log Returns of ETH in USDC	25
8	Cumulative Explained Variance	29
9	Long Position Decision	44
10	Long Position Decision	45
11	Dual Simple Moving Average Strategy Example	52
12	Triple Exponential Moving Average Strategy Example	54
13	Bollinger Bands Strategy Example	55
14	Trading-Engine Components	58
15	Opening a Single Position	62
16	Closing a Position	63
17	Opening a Position	64
18	Trailing Stop Example	65
19	Live-Test Results	71
20	Classification AI Strategy Results	86
21	Dual Simple Moving Average Strategy Results	86
22	Triple Exponential Moving Average Strategy Results	87
23	Bollinger Bands Strategy Results	87
24	Out-of-Sample Classification AI Strategy Results	88
25	Dual Simple Out-of-Sample Moving Average Strategy Results	88
26	Triple Exponential Out-of-Sample Moving Average Strategy Results	89
27	Out-of-Sample Bollinger Bands Strategy Results	89

List of Tables

1	Broker Comparison	6
2	Regime Durations	12
3	Data Split	22
4	Train Data	22
5	Validation Data	23
6	Test Data	23
7	Backtest Data	23
8	Statistics after Logarithmic Returns Transformation	24
9	Validation and Test Metrics for Classification Models	48
10	AI Regression Model Strategy Parameters	50
11	AI Classification Model Strategy Parameters	50
12	Dual Simple Moving Average Strategy Parameters	53
13	Triple Exponential Moving Average Strategy Parameters	54
14	Bollinger Band Strategy Parameters	56
15	Number of Parameters per Strategy	66
16	Combined Strategy Results	67
17	Combined Out-of-Sample Strategy Results	67
18	Total Trading Fees	68
19	Live-Test Statistics	71
20	Classification AI Strategy Results I	91
21	Classification AI Strategy Results II	92
22	Classification AI Strategy Results III	93
23	Dual Simple Moving Average Strategy Results I	95
24	Dual Simple Moving Average Strategy Results II	96
25	Dual Simple Moving Average Strategy Results III	97
26	Triple Exponential Moving Average Strategy Results I	99
27	Triple Exponential Moving Average Strategy Results II	100
28	Triple Exponential Moving Average Strategy Results III	101
29	Bollinger Bands Strategy Results I	103
30	Bollinger Bands Strategy Results II	104
31	Bollinger Bands Strategy Results III	105
32	Description of Other Tested Strategies	106
33	Results of Other Tested Strategies I	107
34	Results of Other Tested Strategies II	107

Code Listing

A	Base FNN	32
B	Dropout NN	33
C	Residual NN	33
D	Base LSTM	34
E	Dropout LSTM	35
F	Bidirectional LSTM	36
G	Encode-Decode LSTM	37
H	Base CNN	38
I	Deep CNN	39
J	Attention CNN	40
K	Concatenation CNN	40
L	CNN + GRU	41
M	SPI Definition	59
N	File-Repository Implementation	59
O	Database-Repository Implementation	59

List of Abbreviations

ETH	Ethereum
BTC	Bitcoin
USDC	USD Coin
DeFi	Decentralized Finance
dApps	Decentralized Applications
M1	Minute-Level
OHLCV	Open, High, Low, Close, Volume
PnL	Profit and Loss
TA	Technical Analysis
CFD	Contract For Difference
EMA	Exponential Moving Average
SMA	Simple Moving Average
MACD	Moving Average Convergence/Divergence
TR	True Range
ATR	Average True Range
RSI	Relative Strength Index
RRR	Risk-Reward Ratio
DD	Drawdown
MDD	Maximum Drawdown
PCA	Principal Component Analysis
AI	Artificial Intelligence
ML	Machine Learning
MLP	Multilayer Perceptron
FNN	Feedforward Neural Network
RNN	Recurrent Neural Network
GRU	Gated Recurrent Unit
LSTM	Long Short-Term Memory
CNN	Convolutional Neural Network
ReLU	Rectified Linear Unit
1D	One Dimensional
Conv1D	Convolutional 1D
SPI	Service Provider Interface
JAR	Java Archive
API	Application Programming Interface
Min	Minimum
Max	Maximum
No	Number
UTC	Coordinated Universal Time
PM	Post Meridiem

1 Introduction

1.1 Motivation

In recent years, the cryptocurrency market has emerged as a highly dynamic and rapidly evolving financial ecosystem. Trading pairs such as ETH/USDC M1 on minute-level (M1) intervals provide vast amounts of high-frequency data, reflecting extreme volatility and market shifts. This creates challenges and opportunities for traders and researchers.

The availability of detailed tick and minute data, combined with direct API access from crypto brokers, creates new opportunities for developing data-driven trading systems. Advances in machine learning and deep learning offer promising tools to identify patterns in market movements. This motivates an exploration of AI-based trading systems specifically designed for the ETH/USDC M1 pair, aiming to leverage technical features and advanced models to improve performance in the volatile cryptocurrency environment.

On the other hand, classical trading strategies also could perform well in this environment, especially if the strategies are adapted to specific market regimes. Techniques such as moving averages, Bollinger Bands, or momentum indicators have the advantage of simplicity and transparency, allowing traders to understand and trust their decision-making process.

The ETH/USDC trading pair at the minute timescale was deliberately chosen because it offers several key advantages for analyzing and modeling algorithmic trading strategies. Ethereum is one of the largest and most liquid cryptocurrencies in the world, while USDC, as a stablecoin, is subject to lower price fluctuations. This combination allows for a clear separation between the volatile component (ETH) and a stable reference currency (USDC), simplifying analysis. The high liquidity of this pair also ensures a reliable data basis with low slippage and tight spreads. In particular, the M1 timescale provides sufficiently granular information to identify short-term market movements and micro-patterns without completely relying on high-frequency tick data. This creates a practical scenario for the development and evaluation of data-driven trading algorithms.

1.2 Aim of this Paper

The primary goal of this paper is to develop and evaluate AI-powered and classical trading strategies focused on the ETH/USDC M1 cryptocurrency pair. This paper will:

1. Retrieve historical ETH/USDC M1 data via a broker API and process it for analysis.
2. Perform exploratory data analysis and feature engineering, including trend, volatility, and momentum indicators.
3. Identify and classify distinct market regimes to provide adaptive trading decisions.
4. Apply risk and money management techniques to realistically simulate trading performance.

5. Design, train, and benchmark multiple deep learning architectures such as LSTM, CNN, and hybrid models for forecasting the price and classify trading decisions.
6. Develop algorithmic trading strategies based on AI model predictions and compare them with classical technical trading approaches.
7. Build a modular trading engine capable of backtesting trading strategies and live execution using real broker connections.
8. Execution of the final best trading strategy in live operation on a demo account with a real broker.

By focusing on the ETH/USDC M1 pair, this paper aims to provide insights into the effectiveness of deep learning and classical strategies in cryptocurrency trading and contribute practical tools for automated, adaptive trading in this challenging asset class.

1.3 Structure of this Paper

[Section 2](#) first lays the foundation for technical integration with financial markets. This includes selecting a suitable broker and setting up an automated interface for data retrieval. Without this foundation, real-time access to market data and the subsequent live execution of strategies are not possible.

[Section 3](#) and [Section 4](#) introduce essential concepts such as market regimes and money and risk management. These theoretical foundations serve to analyze the markets in a structured manner while ensuring that trading decisions are made within the framework of solid risk management. Classification approaches for market phases as well as rules for position sizing and risk limitation are covered.

Building on this, [Section 6](#) provides a comprehensive exploratory data analysis (EDA) approach, describing key steps for processing and transforming historical price data. In addition to generating technical indicators, modern methods such as feature scaling and principal component analysis (PCA) are also discussed as a foundation for machine learning.

Subsequently, in [Section 7](#), various deep learning models are developed and trained to enable predictions of future market movements. Both regression and classification models are examined and evaluated for their suitability for use in trading.

[Section 8](#) is devoted to the development of concrete trading strategies. These are based either on classic technical indicators or on signals derived from the previously developed ML models. Each strategy is methodically described and prepared for later evaluation.

To execute and evaluate these strategies, [Section 9](#) introduces a modular trading engine that enables both backtesting and live trading. The structure follows a plug-in architecture to flexibly integrate different strategies.

[Section 10](#) verifies the performance of the developed strategies using backtests. Various metrics such as profit curves, drawdowns, and win rates are used. The impact of fees is also considered to ensure a realistic evaluation.

[Section 11](#) describes the transition to practice by conducting a live test. This section is intended to demonstrate the extent to which the results observed in the backtest can be confirmed in the real market environment. The results obtained are then critically discussed. [Section 12](#) then deals again with the possible overfitting of trading strategies.

[Section 13](#) wraps up the thesis with a summary of the main findings. [Section 14](#) provides suggestions for further work, for example, regarding improvements to the models, strategies, or trading engine extensions.

1.4 Ethereum

Ethereum is a decentralized, open-source blockchain platform first introduced in 2015 by Vitalik Buterin and other co-founders. Unlike Bitcoin, which was primarily designed as a digital currency, Ethereum is a Turing-complete platform capable of running smart contracts and decentralized applications (dApps). Ethereum's native cryptocurrency is called Ether (ETH) [1].

The Ethereum blockchain is a distributed, immutable mechanism that chronologically stores all transactions. Each block contains transactions linked together using cryptographic methods. The integrity and security of this chain are ensured by the Ethereum Virtual Machine (EVM), which replicates the execution of smart contracts identically on every full node [1].

Unlike centralized systems where a single entity exercises control, blockchain technology enables trustless consensus among network participants. This makes Ethereum particularly relevant for applications in decentralized finance (DeFi), digital identities, and asset tokenization [1].

The use of Ethereum introduces both technical opportunities and limitations that directly impact financial applications, DeFi protocols, and automated trading strategies. One of the key strengths lies in its decentralization and transparency. All transactions and contract states on the network are publicly verifiable and resistant to censorship or manipulation. Furthermore, Ethereum enables high programmability through smart contracts, allowing developers to encode complex financial logic, automate processes, and minimize reliance on intermediaries [2].

However, these strengths are accompanied by certain limitations. Ethereum still faces notable scalability issues. Despite recent upgrades, its transaction throughput remains limited compared to centralized systems, which can lead to congestion and elevated fees during periods of high usage. This is closely tied to the disadvantage of highly volatile transaction fees, which makes small or frequent transactions economically inefficient. Regulatory uncertainty also poses a challenge, as the legal classification of Ethereum remains unresolved in many jurisdictions, leading to problems, especially for institutional participants [2].

In summary, Ethereum provides a powerful and flexible foundation for building decentralized financial systems. Its ability to execute programmable logic on-chain, without the need for central authorities, makes it an ideal infrastructure for digital assets. However, limitations in scalability, cost, and regulatory clarity must be carefully considered when designing systems that depend on Ethereum.

1.5 USD Coin

In addition to traditional cryptocurrencies such as Bitcoin or Ethereum, a special form of digital assets, named stablecoins has emerged. One of the best-known stablecoin is the USD Coin (USDC), whose value is pegged to the US Dollar (USD).

USDC is a digital stablecoin that exists on various blockchains such as Ethereum, Solana, and Polygon. It was launched in 2018 by the company Circle in collaboration with the cryptocurrency platform Coinbase. The goal of USDC is to digitally represent the value of one US dollar. This means that one USDC token should always represent the equivalent of one USD [3].

Unlike volatile cryptocurrencies, USDC is pegged to a fiat currency. To ensure this peg, for every USDC issued, a real US dollar (or an equivalent liquid reserve instrument) is deposited with regulated financial institutions. The custody and transparency of the reserves are regularly audited by independent auditors and publicly documented [4].

USDC is fully programmable and can be transferred globally in near real time without the need for traditional financial intermediaries such as banks. This makes it particularly interesting for applications in decentralized finance (DeFi) systems, crypto trading, or international remittances [5].

Although USDC aims to be pegged 1:1 to the US dollar, they are not the same currency. The US dollar (USD) is the official fiat currency of the United States and is regulated by the US Federal Reserve. USD exists both in physical form (as cash) and digitally (e.g., in bank accounts) and is legal tender [4].

USDC, on the other hand, is a private product of a technology company. It is a tokenized representation of the US dollar that exists exclusively in the digital space. USDC is not legal tender and is not directly controlled by a central bank. Its peg to the US dollar is based on trust in the issuers and their ability to maintain sufficient reserves.

2 Data Source and Broker Selection

Cryptocurrency brokers (also called crypto brokers) play an important role in cryptocurrency trading. Among other things, they act as intermediaries between different market participants. Their key tasks include [6]:

1. **Providing access:** Individuals can participate in the market through a broker and thereby trade various cryptocurrencies. This includes executing orders such as buying cryptocurrencies at the lowest available price or selling them at the highest available price.
2. **Security and compliance:** They also provide customers with a secure platform for executing transactions and adhere to the financial regulations established by authorities.
3. **Leveraging:** Brokers offer customers the opportunity to borrow funds, allowing them to trade with more capital than they actually hold in their account.

This makes trading cryptocurrencies significantly more accessible and user-friendly, especially through simplified interfaces, integrated wallets, and automated order execution. Additionally, the security infrastructure provided by brokers, such as two-factor authentication, can reduce the risk of user error or loss of funds. However, one of the major disadvantages of using brokers is the fees they impose. They can include trading commissions, spreads, funding rates for leveraged positions, and withdrawal fees [7]. Especially in high-frequency or short-term trading, such costs can accumulate quickly and have a substantial impact on overall profitability.

2.1 Broker Selection

For this paper, one broker must be selected for data retrieval and live testing. Since the process is fully automated in short time-frames, the broker must meet certain requirements.

The API must be able to stream market data, request historical data, the current account balance, closed trades, and currently open positions, placing orders, and positions, as well as canceling unfilled orders.

Apart from the API, the broker must support leveraged long/short products like CFDs or margin trading, with the lowest possible fees. They also must provide data in high quality as well as a demo depot.

Table 1 summarizes the required features for some potential brokers. All listed there meet the API functionality requirements.¹

¹Sources: [8], [9], [10], [11], [12], [13]

Broker	Tradable assets	Fees		Leverage
		Maker	Taker	
ByBit	Spot, Spot with Leverage, Futures, Options	0.02%	0.055%	10:1
IG	CFDs, Knock-out-Options	Spread (approx. \$1.30)		2:1
Capital.com	CFDs	Spread (approx. \$1.75)		2:1

Table 1: Broker Comparison

Although ByBit’s fees are calculated as a percentage of the transaction volume, they are lower overall than those charged by IG or Capital.com. The average price of ETH is approximately 2280 USDC ². Based on the average ETH price, ByBit’s fees are approximately 1.25 USDC for a taker order and 0.45 USDC for a maker order per ETH bought or sold. This means it is possible to save up to 1.3 USDC compared to IG and Capital.

Taking into account [Table 1](#), ByBit is the best broker because it has the lowest fees, high quality data, as well as the highest possible leverage.

2.2 API Connection and Data Retrieval Process

Before starting with the Machine Learning process, and the backtests, the first step is to download historical ETH/USDC M1 data via the ByBit API. The request was executed on the `/v5/market/kline` API-Endpoint [14] with the category `linear`, symbol `ETHPERP`, and interval `1` at 17th June, 2025, 11:30 UTC. Since ByBit only returns 1000 candlesticks per request, the same request with different start-, and end-times was executed until the ByBit API does no longer return older candlestick data. This resulted in a candlestick data pool with data on a minute basis from 5th August, 2022, 10:00 UTC to 17th June, 2025, 11:30 UTC. The candlestick data includes the open, high, low, and close prices along with the traded volume (OHLCV data). [Section 6](#) will go into more detail about the data.

²The average price was calculated using all data collected in [Subsection 2.2](#)

3 Market Regimes

Cryptocurrency markets are subject to constant change, which is reflected not only in price movements but also in the underlying structures and dynamics. In quantitative analysis and algorithmic trading, understanding these changes is essential for developing and adapting robust trading strategies. A central concept in this context is the idea of market regimes.

3.1 Introduction to Market Regimes

Market regimes describe phases with distinct statistical and economic characteristics, such as volatility, and trend behavior. They can be understood as different *states* of the market in which certain trading patterns dominate. Distinguishing between, for example, upward, sideways, and downward trends or high and low volatility enables more targeted strategy selection and adaption. Accordingly, the identification and classification of market regimes plays an increasingly important role in modern trading analysis. For example, a strategy that performs well in a stable uptrend may fail in a sideways movement or in periods of high volatility [15].

Understanding market regimes leverages traders and analysts to design strategies that are adaptive, and therefore more robust. Through targeted adaptation of the parameters or the selection of different models, the performance can be increased, and the risk can be reduced.

3.2 How to Categorize the Market?

Categorizing different market characteristics is a common step in identifying market regimes. In literature and practical applications, different approaches to classifying markets exist. A fundamental categorization is often made by the following dimensions:

1. **Trend behavior:** Markets can be categorized trend following (bullish/ bearish) or trendless (sideways) [16].
2. **Volatility:** The volatility of a market is often an indicator for insecurity or stability [17]. High volatility can indicate periods of stress, while low volatility indicates calm markets.
3. **Liquidity:** In illiquid markets, pricing processes can differ significantly from those in liquid markets, which can affect the performance of trading strategies [18].

According to the analysis goal, the categorization can be binary (e.g., bullish vs. bearish) or granular (e.g., a combination of trend behavior, and volatility). While binary classification may be sufficient for high-level market orientation, more granular categorization enables a finer distinction of market conditions.

In the current context, the market is categorized by trend behavior, and volatility. This results in six categories:

1. Downtrend + Low Volatility
2. Downtrend + High Volatility
3. Sideways Trend + Low Volatility
4. Sideways Trend + High Volatility
5. Uptrend + Low Volatility
6. Uptrend + High Volatility

This takes into account the two most central aspects that lead traders to different trading decisions.

3.3 Recognizing Market Regimes

As described in [Subsection 3.2](#), the market will be divided into six categories which are the result of combinations of two individual categories. This makes it possible to categorize the two individual categories individually, and finally merge them.

3.3.1 Recognizing Trend Behavior

The first step is to categorize the market into uptrends, downtrends, and sideways trends. Commonly, a combination of a short-term simple moving average (e.g., SMA(50)), and a long-term simple moving average (e.g., SMA(200)) is used to identify superior trends. The SMA(200) is considered the classic boundary between bullish and bearish market phases. If the short-term simple moving average is above the long-term simple moving average, the market is considered bullish. Vice versa, if the short-term simple moving average is below the long-term simple moving average, the market is considered bearish [19].

For the purpose of the current context, a modified combination of SMA(50), and SMA(100) was chosen. This decision is based on two considerations:

1. **Faster reaction:** The SMA(100) is intended to achieve faster reaction to medium-term trend changes without heavily weighting short-term volatility.
2. **Inertia of trend definition:** A shorter trend window, compared to the SMA(200) reduces the inertia of trend definition, which can be particularly advantageous for more refined classification into uptrends, downtrends, and sideways trend.

Additionally, a minimum slope threshold over the last 15 minutes was integrated for the SMA(50) to avoid that minimal direction changes are mistakenly interpreted as a meaningful trend. This short time span ensures that current market movements are adequately incorporated into the trend classification without being dominated by short-term noise (e.g., individual volatility peaks), and therefore increases the robustness of the trend recognition, and addresses the weaknesses of simple moving averages in sideways

phases. A slope above $+0.05$ signals a significant short-term uptrend, while a slope below -0.05 suggests a clear downtrend. Values in between are interpreted as ambiguous, and are included in the sideways classification accordingly.

The threshold of ± 0.05 was determined using a targeted, visually supported trial-and-error approach. Different values were gradually tested within a range from ± 0.01 to ± 0.15 , and the graphical trend classification was assessed against historical charts. Values below ± 0.05 lead to hypersensitivity to short-term market noise, while values above ± 0.05 detected genuine trend reversals too late. Only with a threshold of ± 0.05 meaningful up- and downtrends could be detected consistently. Despite the lack of a formal optimization framework, this heuristic-guided approach enables a robust separation between noise and relevant market movements.

The market can therefore be divided into three trend phases based on the following criteria:

1. **Uptrend:** The SMA(50) is above the SMA(100), and the slope of the SMA(50) in the last 15 minutes is greater than 0.05.
2. **Downtrend:** The SMA(50) is below the SMA(100), and the slope of the SMA(50) in the last 15 minutes is less than -0.05.
3. **Sideways trend:** The market currently does not meet condition 1 or 2.

Figure 1 shows an example of trend classification of ETH/USDC M1 for 700 minutes.




Figure 1: Trend Classification

3.3.2 Recognizing Volatility

The second step is to categorize the market into phases with high and low volatility. The volatility is calculated as the standard deviation of the logarithmic returns over the last 30 minutes [20]. This locally calculated volatility depicts short-term fluctuation intensity, and enables a context-dependent assessment of current market behavior.

To classify this local volatility, a comparison is made with the median of all available volatilities in the training dataset which was used for fitting the volatility classification indicator. If the current volatility is greater than the median, the market is classified as highly volatile. Otherwise, the market is classified as low volatility.

This threshold definition is deliberately based on a dynamic, data-dependent approach rather than using a fixed absolute threshold. This automatically adapts the volatility classification to each specific market, and can therefore theoretically be applied to other financial markets.

Figure 2 shows the volatility classification on the same base data used in Figure 1.




Figure 2: Volatility Classification

3.3.3 Combining Trend and Volatility Classification

After the two separate classifications in Subsection 3.3.1 and Subsection 3.3.2 the results can be combined. The results of the combination can be seen in Figure 3.




Figure 3: Regime Classification

While classification of the market regimes some discrepancies may occur in individual cases between the algorithmically determined category and the visually perceived category. Such misclassifications are particularly possible during transition phases or in the case of short-term outliers. For example, the downtrend, highly volatile market regime after 10:00 PM, the algorithm detected a downtrend, but visually, an uptrend is perceived. At this time, the market is in a transition, which causes this misclassification.

Crucially, the classification distinguishes consistently and meaningfully, which is the case when viewed as a whole. The robustness and usefulness of the approach does not arise from absolute freedom from errors, but rather from the systematic reduction of uncertainty compared to a purely visual or subjective assessment.

3.4 Dividing Market Regimes in Durations

Based on the combination of trend direction and volatility level, the duration of each market regime is also taken into account to enable a more differentiated classification. The regime duration is divided into three quantiles (33%, 66%, 100%) based on its distribution, allowing for classification into short-term, medium-term, and long-term phases. The notation in the current work for the respective regime durations is $Q_{0.33}$, $Q_{0.66}$ and Q_1 .

It should be noted that the duration used to distinguish between $Q_{0.33}$, $Q_{0.66}$ and Q_1 is calculated per regime across all durations, not within a single occurrence of a regime. For example, if a low-volatility uptrend emerges, it may only be in $Q_{0.33}$ and then transition into a high-volatility uptrend that initially occurs in $Q_{0.33}$ and then in $Q_{0.66}$.

Table 2 shows the duration in minutes that a regime must last to reach the respective quantile.

Regime		Duration in Minutes		
		$Q_{0.33}$	$Q_{0.66}$	Q_1
Down	Low	14	37	319
	High	17	49	442
Sideways	Low	3	9	74
	High	4	11	116
Up	Low	17	50	367
	High	14	37	274

Table 2: Regime Durations

The combination of these three dimensions results in a total of 18 different market regimes. This fine-grained segmentation makes it possible to more precisely capture and differentiate different market conditions. For example, a short-term, volatile uptrend from a long-term, calm downtrend.

Even though this makes the market appear more fragmented, and regimes can change more frequently, this finer subdivision is analytically useful. It allows for context-sensitive market behavior to be examined, and differences in the dynamics of effectiveness of trading strategies in specific regime types to be systematically analyzed.

Instead of smoothing out reality with overly broad categories, a more nuanced picture is deliberately drawn that takes into account both the direction, intensity, and stability of market behavior. The resulting increased complexity is not a disadvantage, but rather a prerequisite for more meaningful, context-based analyses.

4 Money- and Risk-Management

Successful trading is not only based on a good strategy, but also on a disciplined management of capital and risk. Even the best prediction is useless if losses grow uncontrolled or a majority of the capital is risked by a few wrong decisions. This is where money and risk management come into play. They define clear rules regarding how much to invest per trade, what level of risk is acceptable, and how losses can be limited. The goal is to protect capital over the long term, minimizing drawdowns, and profit from positive expected values in a controlled manner. This chapter introduces fundamental concepts, metrics, and methods helping to make rational and sustainable decisions.

4.1 Calculating the Position Size

An important part of risk management is calculating the position size. It is helpful to maximize the potential returns, as well as minimizing the financial risk. Many traders are only willing to risk 1% or 2% of the available capital per trade, to prevent a series of losing trades from decimating the available capital too much [21].

The distance between the estimated entry price and the estimated stop-loss price represents the maximum distance a market can move in an unprofitable direction before a position is automatically closed.

This allows the calculation of the position size to be carried out in three steps:

1. **Determining the risk per trade:** First, it must be determined how much of the available capital should be risked. If 1% of \$10,000 is to be risked, the maximum risk is \$100. It is important to note that the fraction of 1% does not have to be fixed. Thus, it is possible to risk more if the entry signal is very clear. If the entry signal is less clear, it can also be risked less. Only a fixed upper limit should be defined.
2. **Calculating the risk per unit:** To calculate the risk per share the absolute distance between the estimated entry price, and the estimated stop-loss price must be calculated. This represents the risk per unit.
3. **Calculate the position size:** Dividing the risked capital by the risk per unit represents the number of units to buy or sell.

In total, these three steps can be combined into one formula [22]:

$$PositionSize = \frac{AvailableBalance * RiskPerTrade}{RiskPerUnit}$$

So if the available account balance is \$10,000.00, the risk per trade is 1%, and the distance from the estimated entry price to the stop-loss is \$5. The position size is calculated by:

$$PositionSize = \frac{\$10.000 * 1\%}{\$5} = \frac{\$100}{\$5} = 20[Units]$$

It is important to note that numbers can result with many decimal places, and many brokers only allow positions with a certain number of decimal places. If this is the case, the position size must be subsequently rounded to the maximum number of decimal places supported.

4.2 Validating an Entry Signal

Not every entry signal generated by a trading strategy is necessarily profitable. To be profitable in the long term, it is important to ensure that entry signals that are too risky or unrealistic in advance are filtered out, thus preventing positions from being opened. This chapter presents some techniques that can be used to validate entry signals.

4.2.1 Risk-Reward Ratio

The risk-reward ratio (RRR) is a fundamental key figure in trading. It describes the ratio between the potential profit (reward), and the potential loss (risk) of a single trade. The RRR helps in deciding whether an entry signal is too risky or not. It ensures that not only the win rate determines the success of a strategy, but also the ratio of profit to loss in each individual trade [23].

The RRR is calculated by:

$$RRR = \frac{PossibleProfit}{PossibleLoss} = \frac{|OpenPrice - TakeProfitPrice|}{|OpenPrice - StopLossPrice|}$$

For example, if a long trade is opened at \$100, with a take-profit at \$110, and a stop-loss at \$95, the result is:

$$RRR = \frac{|\$100 - \$110|}{|\$100 - \$95|} = \frac{\$10}{\$5} = 2$$

This means that for every dollar risked, a potential profit of two dollars is targeted.

A RRR greater 1 is generally considered positive because the expected profit is higher than the potential loss. However, the RRR should not be viewed in isolation. The essential factor is the combination of RRR, and win rate:

1. **High RRR, low win rate:** e.g., $RRR = 3$ with only a 30% probability of winning \Rightarrow potentially profitable.
2. **Low RRR, high win rate:** e.g., $RRR = 0.5$ with an 80% win rate \Rightarrow also potentially profitable.

The following rule of thumb clarifies when a strategy has a positive expected value in the long term:

$$ExpectedValue = PossibleProfit * WinRatio - PossibleLoss * (1 - WinRatio)$$

Only when this expected value is above zero, a trading strategy is statistically profitable.

The RRR is not only a mathematical metric, but a central component of risk management. It helps traders systematically plan how much they are willing to lose per trade, relative to the expected profit. By consistently applying a minimum RRR (e.g., ≥ 1.5), many inefficient setups can be eliminated in advance [24].

4.2.2 Maximum Account Risk

A central aspect of risk management in algorithmic trading systems is limiting overall risk at the portfolio level. In this work, a fixed risk threshold of a maximum of 10% of the total portfolio value has been defined. This threshold serves as an overarching safety limit to prevent excessive risk accumulating through individual positions or simultaneous strategies.

In this context, overall risk refers to the maximum potential loss of all open positions, weighted by their respective stop-loss levels and position size. The 10% limit means that, aggregated across all positions, the potential loss in the event of a simultaneous stop-out may not exceed one-tenth of the current portfolio value.

The risk limitation serves as a protective mechanism against systemic error in the trading system, as even in the case of erroneous trades or sudden market turmoil, the total loss remains within a calculable maximum risk. Therefore, it increases the portfolio resilience against extreme scenarios.

4.2.3 Minimum Take-Profit

In every trading strategy, transaction fees play an important role. Especially in short-term trading, transaction fees can turn seemingly profitable trades negative if they are not adequately considered. A common mistake is setting the take-profit level too narrowly, resulting in a profit that is smaller than the costs incurred. To trade profitably and sustainably, it is therefore essential that the take-profit at least covers the fees incurred, but ideally, significantly higher.

Assuming a trader sets a take-profit target of 0.1% on a position, while round-trip transaction costs amount to 0.15%. Even with a 100% win rate, the trader would lose money.

4.3 Dealing with Trading Fees

As shown in Table 1 ByBit charges two different fees named maker-, and taker-fee. The maker fee is charged when a limit order is placed in the order book, thereby creating liquidity. In contrast, the taker fee is charged when a market order is executed. This removes liquidity from the market. It is better for a broker if a market is as liquid as possible [25]. Therefore, maker orders incur lower fees than taker orders.

Especially in higher-frequency trading, it is better to charge as few fees as possible. Therefore, it is better for a trader to execute as many limit orders as possible. Two orders are required for a complete trade: one for entry, and one for exit. But typically, three

orders are placed (entry, stop-loss, take-profit), with either the take-profit or stop-loss order being executed.

If a market order is to be executed as an entry order, it is possible to convert it into a limit order by setting the order price slightly below (for buy positions) or slightly above (for short positions) the current price. The same procedure can be followed for the take-profit order.

However, this procedure should not be used for a stop-loss order. If this is converted to a limit position, there is no longer any guarantee that the stop order will be executed at all, as there is no longer any guarantee that the price will rise above or below the order level. This means that a position may remain open significantly longer than intended. If the price then continues to move in an unprofitable direction, the worst-case scenario is that the position is automatically closed by the broker because there is no longer any capital in the account.

The same problem can theoretically occur with a take-profit order. The difference here is that the set stop market order still defines the maximum risk. Thus, while this particular trade may not be profitable, it is impossible to lose all the capital. It is similar with the opening order. If it is not executed, the entire trade is not going to be executed. This means a missed potential profit, but there is no risk of losing capital.

If all orders are executed as planned, and the price moves in a profitable direction, the conversion of the orders will result in a reduction of fees by a factor of 2.75.

5 Performance Comparison

The performance of trading strategies can be compared in many ways. This chapter provides an overview of the most important performance metrics.

5.1 Equity Curve

The equity curve is one of the simplest and fastest ways to compare trading strategies. It is a visual or graphical representation of the account equity over time. A gradually upward sloping equity curve is more preferred than a curve which is very volatile [26].

Figure 4 shows two different equity curves. Even if the more volatile one (orange) is better in the beginning, there are many resets, which causes a worse result compared to the more stable one (blue).




Figure 4: Equity Curve Comparison

5.2 Maximum Drawdown

The maximum drawdown (MDD) measures the largest percentage loss of an investment from its highest point (peak) to its lowest point (trough) within a specific period. It shows the maximum amount of a trading strategy lost before the investment recovers. The MDD is an important risk measure because it highlights the potential losses in times of crisis. A high drawdown indicates high volatility or poor risk management [27].

For each time-point in the equity curve, the drawdown (DD) at time t can be calculated by [28]:

$$DD_t = \frac{Equity_t - PeakValue_{Before\ t}}{PeakValue_{Before\ t}}$$

The maximum drawdown is then the minimum of all drawdowns.

Figure 5 shows the drawdown curves for both equity curves from Subsection 5.1. It is noticeable that the maximum drawdown for the more volatile equity curve is much bigger, which also indicates a not optimal trading strategy.




Figure 5: Drawdown Comparison

5.3 Maximum Gain and Maximum Loss

The maximum loss indicates the largest loss within the tested period. This metric is crucial for assessing the tail risk, i.e., the risk of extreme negative events [29]. A strategy with a very high maximum loss may be profitable on average, but is potentially difficult to trade and may not be sustainable in live operations. Especially in combination with the drawdown, the maximum loss provides an indication of the financial burden a strategy can entail.

The maximum gain describes the highest profit achieved with a single trade. This metric allows conclusions to be drawn about the "opportunity profile" of a strategy, meaning how large the potential profits can be under optimal conditions [30].

However, a very high maximum gain is not necessarily a positive. In many cases, a single upward outlier is the result of random market movements and can distort overall

performance. Strategies that rely heavily on such individual results tend to be more unstable and more susceptible to market changes [30].

The combination of maximum loss and maximum gain makes it possible to evaluate strategies in terms of their return-risk asymmetry. A strategy that enables high profits but also carries a high risk of loss must be evaluated differently than a strategy with a smaller range but consistent performance.

5.4 Standard Deviation of Profits

The standard deviation of profits is a statistical measure that indicates how much the individual results of a trading strategy disperse around the average profit or loss per trade. It is thus a measure of the volatility of individual results and allows conclusions to be drawn about the stability, consistency, and risk of a trading strategy [31].

Formally, the standard deviation of profits is defined as follows:

$$\sigma = \sqrt{\frac{1}{n} \sum_{i=1}^n (p_i - \bar{p})^2}$$

Where:

- p_i is the profit (or loss) of the i th trade,
- \bar{p} is the average profit across all trades,
- n is the number of trades.

The standard deviation is particularly valuable when comparing strategies with similar average profits. It helps assess the consistency of performance [31]:

- A strategy with a low standard deviation consistently produces similar results and is therefore more robust against market noise and random effects.
- A strategy with a high standard deviation exhibits strong fluctuations between winning and losing trades. It is potentially more susceptible to overfitting and more difficult to trade psychologically.

A low standard deviation means that the results of individual trades are close to the average, while a high standard deviation indicates highly fluctuating results. In particular, strategies with few but very large wins and many small losses (or vice versa) usually have a high standard deviation, which increases the risk of miscalculations.

5.5 Win Rate

Another interesting metric for trading strategy comparison is the win rate (or win ratio). It is the percentage of profitable trades in relation to the total number of trades and is calculated as follows:

$$WinRate = \frac{NumberOfWinningTrades}{TotalNumberOfTrades} * 100$$

For example, if for a total of 20 trades, 16 have been profitable, the win rate is $\frac{16}{20} * 100 = 80\%$.

The greatest limitation of the win rate is that a high win rate does not always indicate a profitable trading strategy [32]. For example, if the win rate is 80% with an average gain of 20\$ per profitable trade and an average loss of 90\$ per unprofitable trade, the weighted average is $\frac{80*20-(100-80)*90}{100} = -2\$$. Assuming a different strategy with 40% win rate, but with an average gain of 70\$ per profitable trade and an average loss of 15\$ per unprofitable trade, its weighted average is $\frac{40*70-(100-40)*15}{100} = 19\$$. The second strategy has a much lower win rate, but its long-term outcome is much higher, compared to the first one.

6 Exploratory Data Analysis

Exploratory data analysis uses statistical metrics to gain an understanding of the raw data. In this work, exploratory data analysis also includes data partitioning for training, validation, and testing machine learning models and trading strategies, eliminating data drift, adding additional features, as well as scaling and dimension reduction.

It should be noted that apart from the data splitting, which is a general step, the following steps are only relevant for the machine learning models and not for the classic trading strategies introduced later.

6.1 Dividing the Data

Unlike classic machine learning processes, where data is split in three subsets, named train-, validation-, and test-set, here the data is split in four subsets. The fourth data-set is used for backtesting the real trading strategy, and is therefore not part of the machine learning process but plays an important role in developing the final trading strategy. In contrast, the test-set is only part of the machine learning process and is not used for evaluating or backtesting trading strategies. This distinction was made so that the critical phases (testing machine learning models and testing trading strategies) cannot influence each other.

Set	From	To	No. of Datapoints	% of All Data
Complete	08/05/2022 10:00 UTC+2	06/17/2025 11:30 UTC+2	1,507,598	
Train	08/05/2022 10:00 UTC+2	04/30/2024 23:59 UTC+2	913,684	60.6%
Validation	05/01/2024 00:00 UTC+2	09/30/2024 23:59 UTC+2	220,320	14.6%
Test	10/01/2024 00:00 UTC+2	12/31/2024 23:59 UTC+2	132,482	8.8%
Backtest	01/01/2025 00:00 UTC+2	06/17/2025 11:30 UTC+2	241,111	15.9%

Table 3: Data Split

After the splitting, the four subsets have the following summaries:

	Open	High	Low	Close	Volume
	Open	High	Low	Close	Volume
count	913684.0	913684.0	913684.0	913684.0	913684.0
mean	1933.4	1933.95	1932.85	1933.4	7.58
std	619.6	619.92	619.28	619.6	105.09
min	1074.35	1077.45	1064.05	1074.35	0.0
25%	1578.44	1578.89	1577.99	1578.44	0.0
50%	1801.52	1801.81	1801.13	1801.52	0.05
75%	2083.4	2083.84	2083.05	2083.4	2.02
max	4098.5	4099.48	4096.35	4098.5	31350.65

Table 4: Train Data

	Open	High	Low	Close	Volume
count	220320.0	220320.0	220320.0	220320.0	220320.0
mean	3050.36	3051.3	3049.4	3050.36	2.89
std	473.4	473.45	473.34	473.4	21.45
min	2111.8	2160.4	2088.13	2111.8	0.0
25%	2617.31	2618.04	2616.71	2617.31	0.0
50%	3063.4	3064.31	3062.43	3063.4	0.12
75%	3462.22	3463.29	3461.21	3462.22	1.32
max	3974.68	3976.96	3969.94	3974.68	4972.18

Table 5: Validation Data

	Open	High	Low	Close	Volume
count	132482.0	132482.0	132482.0	132482.0	132482.0
mean	3093.93	3095.27	3092.58	3093.94	4.42
std	531.85	532.29	531.4	531.85	18.42
min	2309.01	2311.73	2307.73	2309.01	0.0
25%	2541.8	2542.65	2540.94	2541.8	0.08
50%	3143.69	3145.38	3142.02	3143.7	0.68
75%	3487.9	3489.46	3486.33	3487.9	2.92
max	4107.28	4112.68	4102.6	4107.28	1341.05

Table 6: Test Data

	Open	High	Low	Close	Volume
count	241111.0	241111.0	241111.0	241111.0	241111.0
mean	2432.57	2433.77	2431.35	2432.57	7.0
std	574.72	574.95	574.49	574.72	32.11
min	1386.6	1395.8	1382.99	1386.6	0.0
25%	1887.22	1888.28	1886.1	1887.22	0.18
50%	2510.29	2511.4	2509.1	2510.29	1.26
75%	2732.0	2733.21	2730.7	2732.0	4.94
max	3742.33	3745.13	3739.65	3742.33	2534.66

Table 7: Backtest Data

6.2 Using Log>Returns

In the summary statistics of the subsets ([Table 4](#), [Table 5](#), [Table 6](#), [Table 7](#)) it is noticeable that the mean values change over time. This becomes also clear when visualizing the data ([Figure 6](#)).

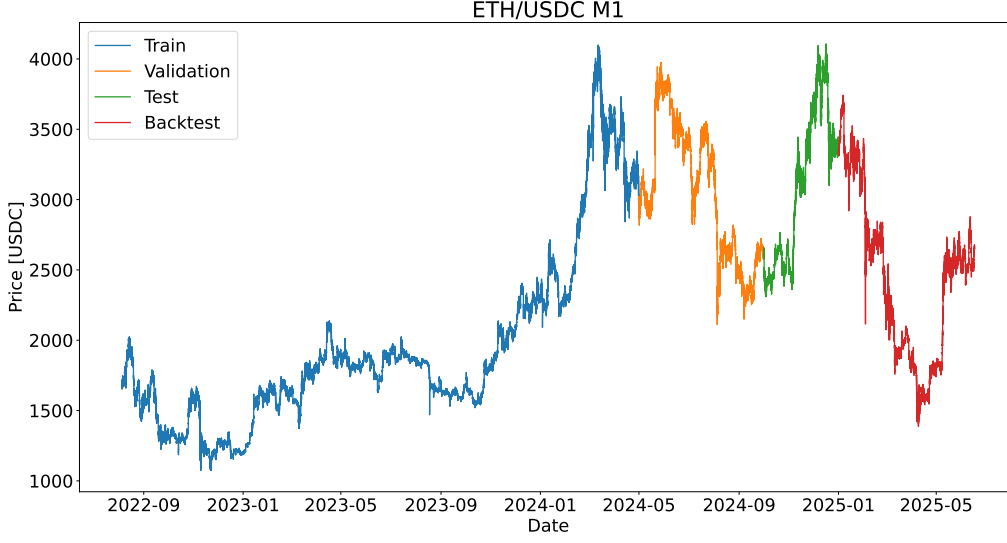


Figure 6: Price Fluctuation of ETH in USDC

To avoid this data drift, the price is transformed to its logarithmic returns (also called log-returns). These are calculated as follows [33]:

$$\text{LogReturn}_t = \ln\left(\frac{\text{Price}_t}{\text{Price}_{t-1}}\right)$$

After the transformation, the means and standard deviations in the subsets are very similar, and the data does no longer drift over time. Table 8 shows the means and standard deviations after the logarithmic return transformation. This was also confirmed by an augmented Dickey-Fuller test (ADF test). The ADF test is a common statistical procedure for investigating whether a time series contains a unit root, which would indicate non-stationarity [34]. In this context, the null hypothesis that the time series is non-stationary, i.e., exhibits a stochastic trend, was tested. The test result yielded a p-value of 0.0, which is below the 5% significance level, so the null hypothesis could be rejected. This indicates that the time series is stationary and that there is no significant data drift over the period under consideration. Additionally, Figure 7 shows visually that the data drift is eliminated.

Data Set	Mean	Standard Deviation
Train	6.5E-7	8.7E-4
Validation	-6.1E-7	9.4E-4
Test	1.8E-6	9.4E-4
Backtest	-9.5E-7	1.1E-3

Table 8: Statistics after Logarithmic Returns Transformation




Figure 7: Log Returns of ETH in USDC

If an exact price is to be restored from the logarithmic returns, this can be done with the previous price using the following formula, which can be derived by transformation:

$$Price_t = Price_{t-1} * e^{LogReturn_t}$$

If a series of logarithmic returns is present, any price in the series can also be calculated using the following formula, where $Price_0$ is the price at the start of the series and t is the distance to the start of the series:

$$Price_t = Price_0 * \prod_{i=0}^t e^{LogReturn_i} = Price_0 * e^{\sum_{i=0}^t LogReturn_i}$$

The transformation is applied to the open, high, low, and close prices, and the original prices are replaced by the logarithmic returns, so that all subsequent actions are carried out with the prices on the logarithmic returns.

6.3 Additional Features

To provide the machine learning models with more context about the price, additional features from different categories were added to the raw transformed OHLCV data.

6.3.1 Trend Following Indicators

In financial analysis trend following indicators play an essential role while modeling a predicting future price movements. This occurs because markets move in trends that are

repeatedly interrupted by outliers. This results in a zigzag movement that nevertheless moves in one direction. Trend-following indicators can be used to filter out these outliers [35].

The exponential moving average (EMA) is one type of trend following indicator. It is a variation of the classic simple moving average (SMA), placing more emphasis on newer prices. It is often used by traders in length of 10-, 50-, and 200-period.

An important aspect of distinguishing between EMA and SMA is that many traders assume that new data better reflects the current trend, while others argue that giving more weight to recent prices can lead to distortion [36].

Since the aim of this paper is short-term predictions and the EMA reacts faster to price changes than the SMA, the EMA could provide more relevant context for the model. The EMA was added in 5, 10, 20, 30, 50, and 200 period to the data.

Another trend following indicator is the moving average convergence/divergence (MACD) which does not only help to identify price trends, but also helps to measure the trend momentum. It shows the relationship between two exponential moving averages. To calculate the MACD line, an EMA(12) is subtracted from an EMA(26). Additionally, a signal line is calculated as an EMA(9) of the MACD line [37].

Although the MACD can signal possible reversals, it is also known for generating a significant number of false positives. This tends to occur particularly in sideways or ranging markets, where price movements lack strong momentum or clear directional trends. In such environments, the MACD lines frequently cross without a meaningful shift in market sentiment, leading traders to act on misleading signals. Consequently, relying solely on MACD in these conditions can result in wrong entries or exits [37]. To mitigate this, many traders combine MACD with additional indicators, such as support and resistance levels or trend confirmation tools, to improve decision-making and reduce noise.

Despite its susceptibility to false signals in sideways markets, the MACD can still be useful as a feature in a machine learning model. This is because the indicator provides valuable information about momentum and potential trend changes, especially during periods of clear price movements. In a machine learning context, the MACD is not considered in isolation but processed in combination with many other features. This allows the model to learn under which market conditions the MACD is more reliable and when it should be ignored. The nonlinear relationships and interactions with other features allow the model to identify context-dependent patterns from the MACD signals, which can be useful for predicting price movements.

6.3.2 Volatility Indicators

To measure volatility, there are other indicators and techniques to measure the volatility, in addition to those described in [Subsection 3.2](#).

One indicator is the average true range (ATR) which decomposes the entire range of an asset price for a period. It is calculated by determining the so-called true range (TR) for each candlestick - the maximum of: current high minus low; distance from the previous closing price up, and down. The ATR is then the moving average of these TR values, usually over 14 periods.

The ATR has two main limitations. The first is that an ATR value must always be set into comparison to previous ATR values, because one single value is not enough to tell if a trend is going to reverse. The second limitation is that the ATR does not tell anything about the direction of the price [38]. The ATR was added in 5, 7, 10, 14, and 18 periods to the data.

Another volatility indicator are Bollinger Bands, which consist of three lines. The middle line is a SMA of the closing prices, the lower line is calculated by subtracting a certain number of price standard deviations from the middle line, and the upper line is calculated by adding a certain number of price standard deviations to the middle line. Usually the double of the price standard deviation is added, and subtracted from the middle line [39].

The higher the volatility of the market is in the last closing prices, the wider the band gets. If the price of the market rises near the upper band, traders see the market as overbought. Similarly, if the market falls near the lower band, the market could be oversold. This allows to generate possible entry and exit signals [39]. The three Lines were added to the data with a 15, 20, and 25 period SMA with double price standard deviation.

6.3.3 Momentum Indicators

Momentum measures the strength and direction of a price movement over a certain period of time. Momentum indicators are useful because they give insights into the strength of trending prices. Therefore, they can indicate possible reversals in the trend direction [40].

A common momentum indicator is the relative strength index (RSI). It measures the speed and magnitude of an asset price by comparing the average gains and losses of the asset, and can be used to detect overbought and oversold conditions. The RSI ranges between zero and 100. Usually an RSI over 70 indicates an overbought, and an RSI below 30 indicates an oversold market. Commonly the default RSI period to compare the average gains, and losses is 14 [41]. The RSI was added in periods 7, 14, and 20 to the data.

To depict relative trend strength, a sophisticated momentum indicator was constructed that compares the log returns of two different time frames. This feature allows the model to distinguish phases of accelerating price movements from stable or declining trends. The use of logarithmic returns simultaneously achieves scale independence, and improved comparability, which is particularly advantageous for modeling financial market-related time series. This indicator was added for time frames M2, M3, M6, M9, and M12.

6.3.4 Price Transformation Indicators

Apart from the mentioned indicators shifted logarithmic returns for the last six minutes have been added to provide additional context about the last price movements in a compact form. This could help the models to recognize trend reversals, volatility changes or short-term patterns.

Lastly, the logarithmic returns of other time frames (M2, M3, M6, M9, and M12) have been added to the data providing another more stable trend context which helps to correctly classify short-term price movements. This creates a balanced feature set that takes into account both rapid reactions and long-term patterns.

6.4 Further Transformations

Due to the previous steps, the data contains 56 features (transformed OHLCV price data and additional indicators). Two further transformations are then performed on the features. This includes scaling and dimensionality reduction. It is important to note that before these steps, the data is clustered into the respective regimes using the algorithm described in [Section 3](#), and both scaling and dimensionality reduction are performed in isolation for each regime. This brings advantages in both steps.

When scaling, the value range of the scaler can be better utilized. This is evident, for example, with volatility indicators such as the ATR. This indicator generates higher values in volatile phases than in less volatile ones. If the scaler is fitted to the entire data and then a differentiation between market regimes is made, the entire value range of the scaler cannot be utilized in the respective regimes. However, once the differentiation into regimes is made, the entire value range is utilized for each regime.

Dimensionality reduction also offers an advantage. Here, different features may be more dominant than others in different regimes. This allows to examine which features are most important for each regime individually, preventing dilution.

6.4.1 Scaling the Data

Scaling data is an important step in exploratory data analysis. This is because, on the one hand, machine learning models perform better with scaled data [\[42\]](#), and, on the other hand, principal component analysis ([Subsection 6.4.2](#)) is sensitive to unscaled data [\[43\]](#).

The use of the scikit-learn `MinMaxScaler` [\[44\]](#), which scales data in a fixed range of $[0; 1]$ is useful and widely used. Neural networks, especially those with activation functions such as ReLU, which are introduced in [Subsection 7.2](#) benefits greatly from inputs with a uniform range of values. The scaling stabilizes the gradient flow, shortens training time and reduces vanishing gradients [\[45\]](#).

For these reasons, all features were scaled with the `MinMaxScaler` and are used exclusively in the scaled form in all steps related to machine learning.

6.4.2 Principal Component Analysis

The principal component analysis (PCA) is a process for dimensional reduction, by linearly transforming high dimensional datasets to a small number of uncorrelated principal components (directions of the new coordinate system). During the transformation, it can be specified how much variance in the data can be eliminated. After a transformation using PCA, it is ensured that at least the specified variance is retained [\[46\]](#).

Especially when processing numerous technical indicators or derived features in financial data, the high dimensionality can become problematic - a phenomenon known as the curse of dimensionality. This term describes the increasing challenges in modeling as the number of dimensions or features increases. Data points become increasingly sparsely disturbed, computational costs increase, and many models lose their ability to generalize. Applying PCA allows redundant or correlated information to be condensed, making the model more robust, faster, and easier to interpret. At the same time, the risk of overfitting is reduced because the model focuses on the most important structures in the dataset [47].

Figure 8 shows the cumulative explained variance for the 56 added features in Subsection 6.3 for each quantile market regime. It shows that the cumulative variance increases rapidly at the beginning. In this case, the reduction of the projection to just 3 to 6 principal components already explains at least 80% of the variance. This means that the majority of the statistically relevant structures in the dataset are retained, even though the number of features has been massively reduced. Even if 20% of the variance is lost, the benefit outweighs this: The remaining principal components capture the statistical essence of the original feature space in a significantly more compact and robust form, which is particularly well-suited for machine learning.




Figure 8: Cumulative Explained Variance

7 Deep Learning Models

This chapter presents the key components for developing and evaluating deep learning models. This includes both the selection of suitable models and the criteria for assessing their performance.

For this work, Keras version 3.10.0 was used as a high-level API for modeling and training neural networks [48]. Instead of the standard TensorFlow backend, PyTorch version 2.7.1 was used as the backend. A key reason for this decision was its better support on Windows, compared to TensorFlow [49], particularly with regard to installation and compatibility with existing CUDA drivers. By using Keras in combination with the PyTorch backend, user-friendly modeling could be combined with stable and well-supported execution on Windows systems.

The model architectures used in this work are not based on specific, citable publications, but were designed as part of an experimental and iterative development process. The goal was to construct powerful models well-suited to the given problem. Particularly with regard to processing sequential data of limited length and high variability.

Established neural building blocks such as 1D convolutional layers, GRUs, and attention mechanisms were used, the fundamentals of which are comprehensively documented in the literature. However, the specific design of the model architectures, such as the combination of multiple parallel CNN paths, the integration of GRUs after convolution steps, or the use of Global Average Pooling instead of flattening, represents a creative and pragmatic composition of its own.

Additionally, all models were automatically optimized using Optuna, so that architectural decisions were also influenced by the hyperparameter search. In many cases, initial ideas originate from non-scientific sources such as blog posts, online tutorials, or community code snippets, which are not scientifically citable.

In summary, the developed models are original variations, inspired by familiar architectural elements, but not directly adopted or reconstructed from specific publications. This allows for greater flexibility and reinforces the experimental nature of the work.

For both regression and classification models, 13 different architectures were used, which can be divided into neural networks, convolutional neural networks, and long short-term memories.

7.1 Optuna

In modern machine learning methods, the selection of hyperparameters plays a central role in model performance. Hyperparameters such as learning rate, the number of layers in a neural network, or regularization strengths directly influence the behavior and generalization of the model. The search for optimal values for these parameters, the so-called hyperparameter optimization, is often a time-consuming and computationally intensive process [50].

Optuna is a modern framework for automated hyperparameter optimization designed for efficiency, flexibility, and ease of use. It was developed to enable easy integration into existing machine learning pipelines while providing powerful, goal-oriented optimization.

The central abstraction in Optuna is the concept of a study, which defines the optimization task. It encapsulates the objective function and maintains a history of evaluated hyperparameter configurations. Within a study, multiple trials are executed. Each trial represents a single evaluation of the objective function with a specific set of hyperparameters [51].

The objective function is a user-defined Python function that receives a set of hyperparameters as input and returns a scalar value, which is typically the validation loss or accuracy, that reflects the performance of the model with those parameters. Optuna uses this value to decide how to sample the next set of hyperparameters. Thus, the objective function is the core of the optimization process, as it defines what is being optimized and how success is measured [51].

The goal of Optuna is to automatically find those hyperparameter combinations that satisfy a specific optimization criterion (e.g., minimum validation error rate or maximum accuracy). The search process should be as efficient as possible, requiring as few model training sessions as possible [52].

In this work, Optuna was used to automatically run multiple training models with different hyperparameters to find the best hyperparameters. In the relevant sections of the next chapters, the range of the hyperparameters for each model will be mentioned.

7.2 Neural Networks

Neural networks are among the central methods of machine learning and form the basis of many modern deep learning models. The simplest and most widely used form is the feedforward neural network (FNN), also known as a multilayer perceptron (MLP). Such networks are universally applicable and, with appropriate structuring, can be used for classification, regression, and pattern recognition.

A classic neural network consists of several layers of artificial neurons that forward information in a fixed order from input to output. The structure can be divided as follows [53]:

1. **Input layer:** Accepts the raw input data, e.g., a vector form of a time series or preprocessed features.
2. **Hidden layers:** One or more layers of artificial neurons, each of which calculates a weighted sum of the inputs and further processes it using an activation function. Typical activation functions are ReLU, sigmoid, or tanh.
3. **Output layer:** Provides the final result, e.g., class membership (for classification) or continuous value (for regression).

In this work, three different neural network models have been tested. The tested models are different in their complexity, depth, and methodical approach to data processing.

7.2.1 Base Model with Flatten-Architecture

The first model follows a classic feedforward approach, with one to three dense layers, each with 32 to 128 neurons and a rectified linear unit (ReLU) activation function. The ReLU activation function was used because it provides an efficient calculation and a efficient gradient propagation with fewer vanishing gradients [54]. Either before or after the dense layers, the input data are flattened [55]. This decision influences whether the model processes all timepoints as a vector early on or treats each time step separately. This architecture was used as a simple baseline model to obtain a quick training and a first baseline model.

The learning rate of the Adam optimizer ranges between 10^{-5} and 10^{-2} , while the input layer has a length of 5 to 150. All hyperparameters are managed by Optuna.

```
1 num_layers = trial.suggest_int('num_layers', 1, 3)
2 num_units = trial.suggest_int('num_units', 32, 128)
3 input_length = trial.suggest_int('input_length', 5, 150)
4 learning_rate = trial.suggest_float('learning_rate', 1e-5, 1e-2,
5                                     log=True)
6 flatten_before = trial.suggest_categorical('flatten_before',
7                                             [True, False])
8
9 model = Sequential()
10
11 model.add(InputLayer(shape=(input_length, input_dimension)))
12
13 if flatten_before:
14     model.add(Flatten())
15
16 model.add(Dense(num_units, activation='relu'))
17 for _ in range(num_layers - 1):
18     model.add(Dense(num_units, activation='relu'))
19
20 if not flatten_before:
21     model.add(Flatten())
22
23 model.add(Dense(...))
24
25 model.compile(optimizer=Adam(learning_rate=learning_rate), ...)
```

Listing A: Base FNN

7.2.2 Dropout Neural Network

In the second model, the data is first converted into a one-dimensional vector using a flatten layer. This is followed by one to three dense layers using ReLU activation. Batch normalization supports stable and rapid convergence [56], while dropout between 5% and 35% of the data is used as a regularization method to prevent overfitting [57]. All other parameters are identical to those in the simple neural network.

The parameters tuned via Optuna control the model depth, the number of nodes per layer, the training dynamics, and the number of historical time points used as input.

```
1 num_layers = trial.suggest_int('num_layers', 1, 3)
2 num_units_1 = trial.suggest_int('num_units_1', 32, 256)
3 num_units_2 = trial.suggest_int('num_units_1', 32, 256)
4 dropout = trial.suggest_float('dropout', 0.05, 0.35)
5 input_length = trial.suggest_int('input_length', 5, 150)
6 learning_rate = trial.suggest_float('learning_rate', 1e-5, 1e-2,
7                                     log=True)
8
9 model = Sequential()
10
11 model.add(InputLayer(shape=(input_length, input_dimension)))
12
13 model.add(Flatten())
14
15 model.add(Dense(num_units_1, activation='relu'))
16 model.add(BatchNormalization())
17 model.add(Dropout(dropout))
18 for _ in range(num_layers - 1):
19     model.add(Dense(num_units_2, activation='relu'))
20
21 model.add(Dropout(dropout))
22 model.add(Dense(...))
23
24 model.compile(optimizer=Adam(learning_rate=learning_rate), ...)
```

Listing B: Dropout NN

7.2.3 Residual Neural Network

The third model integrates a residual concept, originally known from image processing but also increasingly used in time series contexts. The architecture consists of a linear shortcut branch connected to the main path via an additive connection. This construction facilitates the training of deeper networks by improving gradient flow and preventing information loss through multiple layers [58].

In addition to the previous parameters, two additional variables controlled by Optuna are added: `num_units_res_prep` and `num_units_res`. These determine the complexity of the residual branch. The distinction whether flattening is done before or after the hidden layers is also included here to flexibly adapt data processing to the underlying data structure.

This architecture is particularly suitable for nonlinear and complex relationships, such as those frequently encountered in ETH price forecasting. The residual path allows the model to pass on basic information directly, while deeper layers learn abstract features.

```
1 num_layers = trial.suggest_int('num_layers', 1, 3)
2 num_units = trial.suggest_int('num_units', 32, 128)
3 num_units_res_prep = trial.suggest_int('num_units_res_prep', 32, 128)
4 num_units_res = trial.suggest_int('num_units_res', 32, 128)
5 input_length = trial.suggest_int('input_length', 5, 150)
6 learning_rate = trial.suggest_float('learning_rate', 1e-5, 1e-2,
```

```
7                                     log=True)
8 flatten_before = trial.suggest_categorical('flatten_before',
9                                           [True, False])
10
11 input_layer = Input(shape=(input_length, input_dimension))
12
13 x = Flatten()(input_layer)
14
15 res = Dense(num_units_res_prep, activation='relu')(x)
16 res = Dense(num_units_res, activation='linear')(res)
17
18 x = Dense(num_units_res, activation='relu')(x)
19 x = Add()([x, res])
20 x = Dense(num_units, activation='relu')(x)
21
22 output = Dense(...)(x)
23
24 model = Model(input_layer, output)
25
26 model.compile(optimizer=Adam(learning_rate=learning_rate), ...)
```

Listing C: Residual NN

7.3 Long Short-Term Memory Models

In addition to traditional feedforward neural networks, Recurrent Neural Networks (RNNs) based on the Long Short-Term Memory (LSTM) architecture were also used. LSTMs are specifically designed to capture temporal dependencies in sequential data and retain long-term information across multiple time steps [59]. This is particularly important for financial data such as ETH prices, as current prices are often influenced by past developments.

LSTM cells have an internal memory structure and three control mechanisms (input, forget, and output gates) that allow the network to decide which information should be stored, overwritten, or passed on [60]. This effectively mitigates the so-called vanishing gradient problem of traditional RNNs [61].

7.3.1 Base LSTM

This model consists of one or two stacked LSTM layers followed by a dense output layer. If only one LSTM layer is added, it returns the result directly. With multiple layers, the temporal sequence information is passed to the second LSTM layer.

Similar to feedforward neural networks, the number of layers, the number of neurons per layer, the learning rate, and the length of the input window used are managed via Optuna.

This model represents the basic form of a sequential predictor and is well suited for simple time series patterns.

```
1 num_layers = trial.suggest_int('num_layers', 1, 2)
```

```
2 num_units_input = trial.suggest_int('num_units_input', 32, 128)
3 num_units = trial.suggest_int('num_units', 32, 128)
4 input_length = trial.suggest_int('input_length', 30, 150)
5 learning_rate = trial.suggest_float('learning_rate', 1e-5, 1e-2,
6                                     log=True)
7
8 model = Sequential()
9 model.add(InputLayer(shape=(input_length, input_dimension)))
10 if num_layers == 1:
11     model.add(LSTM(num_units_input, return_sequences=False))
12 else:
13     model.add(LSTM(num_units_input, return_sequences=True))
14
15 for i in range(num_layers - 1):
16     if i == num_layers - 2:
17         model.add(LSTM(num_units, return_sequences=False))
18     else:
19         model.add(LSTM(num_units, return_sequences=True))
20
21 model.add(Dense(...))
22
23 model.compile(optimizer=Adam(learning_rate=learning_rate), ...)
```

Listing D: Base LSTM

7.3.2 Dropout LSTM

This variant extends the standard model with dropout layers inserted between the LSTM units, discarding between 5% and 50% of the data.

The additional parameter is also determined by Optuna and allows fine-tuned control over the regularization strength. This model is particularly useful when working with noisy or volatile price trends, such as cryptocurrencies.

```
1 num_layers = trial.suggest_int('num_layers', 1, 2)
2 num_units_input = trial.suggest_int('num_units_input', 32, 128)
3 num_units = trial.suggest_int('num_units', 32, 128)
4 input_length = trial.suggest_int('input_length', 30, 150)
5 dropout = trial.suggest_float('dropout', 0.05, 0.5)
6 learning_rate = trial.suggest_float('learning_rate', 1e-5, 1e-2,
7                                     log=True)
8
9 model = Sequential()
10 model.add(InputLayer(shape=(input_length, input_dimension)))
11 if num_layers == 1:
12     model.add(LSTM(num_units_input, return_sequences=False))
13 else:
14     model.add(LSTM(num_units_input, return_sequences=True))
15
16 model.add(Dropout(dropout))
17 for i in range(num_layers - 1):
18     if i == num_layers - 2:
```

```
19         model.add(LSTM(num_units, return_sequences=False))
20     else:
21         model.add(LSTM(num_units, return_sequences=True))
22     model.add(Dropout(dropout))
23
24 model.add(Dense(...))
25
26 model.compile(optimizer=Adam(learning_rate=learning_rate), ...)
```

Listing E: Dropout LSTM

7.3.3 Bidirectional LSTM (BiLSTM)

The third model uses bidirectional LSTM layers. These process the input sequence forward and backward simultaneously, allowing both earlier and later information to be incorporated into the internal state [62].

This architecture can be particularly valuable in the context of ETH predictions based on historical data, as it allows for better detection of symmetric patterns or turning points in price trends.

The hyperparameter structure is similar to the previous models, but the modeling is performed using double LSTM paths.

```
1 num_layers = trial.suggest_int('num_layers', 1, 2)
2 num_units_input = trial.suggest_int('num_units_input', 32, 128)
3 num_units = trial.suggest_int('num_units', 32, 128)
4 input_length = trial.suggest_int('input_length', 30, 150)
5 learning_rate = trial.suggest_float('learning_rate', 1e-5, 1e-2,
6                                     log=True)
7
8 model = Sequential()
9 model.add(InputLayer(shape=(input_length, input_dimension)))
10
11 if num_layers == 1:
12     model.add(Bidirectional(LSTM(num_units_input,
13                                 return_sequences=False)))
14 else:
15     model.add(Bidirectional(LSTM(num_units_input,
16                                 return_sequences=True)))
17
18 for i in range(num_layers - 1):
19     if i == num_layers - 2:
20         model.add(Bidirectional(LSTM(num_units,
21                                     return_sequences=False)))
22     else:
23         model.add(Bidirectional(LSTM(num_units,
24                                     return_sequences=True)))
25
26 model.add(Dense(...))
27
28 model.compile(optimizer=Adam(learning_rate=learning_rate), ...)
```

Listing F: Bidirectional LSTM

7.3.4 Encode-Decode Model with Repeat-Vector

The fourth model is based on an encoder-decoder approach, as used in machine translation. The input sequence is first converted into a fixed context vector by an LSTM layer, which is then duplicated multiple times using RepeatVector [63]. A decoder-LSTM interprets this vector sequentially to generate a temporally structured output.

This architecture is particularly suitable for multistep forecasting, as it can map not only the next value but an entire sequence of future values. The final preprocessing is performed using TimeDistributed(Dense(...)) to calculate a separate output for each time step [64].

```
1 num_units_input = trial.suggest_int('num_units_input', 32, 128)
2 num_units = trial.suggest_int('num_units', 32, 128)
3 input_length = trial.suggest_int('input_length', 30, 150)
4 learning_rate = trial.suggest_float('learning_rate', 1e-5, 1e-2,
5                                     log=True)
6
7 model = Sequential()
8 model.add(InputLayer(shape=(input_length, input_dimension)))
9
10 model.add(LSTM(num_units_input, return_sequences=False))
11
12 model.add(RepeatVector(30))
13
14 model.add(LSTM(num_units, activation='relu', return_sequences=True))
15 model.add(TimeDistributed(Dense(...)))
16
17 model.compile(optimizer=Adam(learning_rate=learning_rate), ...)
```

Listing G: Encode-Decode LSTM

7.4 Convolutional Neural Networks

Convolutional neural networks (CNNs) represent a special architecture within deep learning that was developed specifically to process data with spatial structures, such as images. Compared to traditional feedforward networks, CNNs are significantly more efficient in handling high-dimensional inputs and have established themselves as standard methods in numerous application areas such as object detection, natural language processing, and time series analysis [65].

The key difference between CNNs and classical neural networks lies in the use of so-called convolutional layers, which exploit the local structures in the input data. Instead of connecting every neuron to all input values (as with a fully connected layer), local filters (kernels) are used that respond only to a small portion of the input space. These filters are automatically learned during the training process [65].

Time series typically exhibit temporal dependencies, repetitions, seasonal patterns, or short-term fluctuations. Classic methods for processing this data, such as LSTM networks, explicitly model such dependencies using recursive states. CNNs, on the other hand, uses an alternative approach, called local convolutions, which also serve time series to detect relevant patterns or changes over short or medium periods of time [66].

The advantages of using CNNs for time series include:

1. **Parallel processing:** CNN does not require recursive loops, e.g., compared to LSTMs.
2. **Local pattern detection:** CNNs can detect local patterns, such as peaks, fluctuations, or trend changes.
3. **Computational complexity:** CNNs requires lower computational power compared to recursive architectures.

7.4.1 Base CNN

The simplest variant consists of two stacked one-dimensional convolutional layers, followed by maxpooling, and two dense layers for output preparation. The position of the flattening layer affects the processing order between the dense and the sequence structure.

The hyperparameters tuned by Optuna are the number of units in the convolutional layers, the length of the filters (kernel size), the size of the maxpooling region, the position of the flattening layer, and the number of neurons in the dense layers.

```
1 num_units_cnn = trial.suggest_int('num_units_cnn', 32, 128)
2 num_units = trial.suggest_int('num_units', 32, 128)
3 input_length = trial.suggest_int('input_length', 5, 150)
4 kernel_size = trial.suggest_int('kernel_size', 2, 5)
5 pool_size = trial.suggest_int('pool_size', 1, 5)
6 learning_rate = trial.suggest_float('learning_rate', 1e-5, 1e-2,
7                                     log=True)
8 flatten_before = trial.suggest_categorical('flatten_before',
9                                             [True, False])
10
11 model = Sequential()
12 model.add(InputLayer(shape=(input_length, input_dimension)))
13 model.add(Conv1D(num_units_cnn, kernel_size=kernel_size,
14                 activation='relu'))
15 model.add(MaxPooling1D(pool_size=pool_size))
16 if flatten_before:
17     model.add(Flatten())
18
19 model.add(Dense(num_units, activation='relu'))
20
21 if not flatten_before:
22     model.add(Flatten())
23
24 model.add(Dense(...))
25
```

```
26 model.compile(optimizer=Adam(learning_rate=learning_rate), ...)
```

Listing H: Base CNN

7.4.2 Deep CNN

The model shown here represents an extension of a simple CNN model (Base CNN). While the simple model uses only a single 1D convolutional layer followed by a pooling and dense layer, this Deep CNN model uses two consecutive Conv1D layers. This additional depth allows the network to learn more complex and hierarchical feature representations from the input data [67]. This is particularly useful for timeseries with multi-level dependencies or subtle patterns. The model is complemented by a MaxPooling layer, which reduces the temporal resolution and counteracts overfitting. The optional placement of the Flatten layer allows for the evaluation of different configurations of feature identification, which in turn increases model flexibility. Overall, this more deeply structured CNN enables more powerful modeling than the flattened variant.

```
1 num_units_cnn = trial.suggest_int('num_units_cnn', 32, 128)
2 num_units = trial.suggest_int('num_units', 32, 128)
3 input_length = trial.suggest_int('input_length', 5, 150)
4 kernel_size = trial.suggest_int('kernel_size', 2, 5)
5 pool_size = trial.suggest_int('pool_size', 1, 5)
6 learning_rate = trial.suggest_float('learning_rate', 1e-5, 1e-2,
7                                     log=True)
8 flatten_before = trial.suggest_categorical('flatten_before',
9                                             [True, False])
10
11 model = Sequential()
12 model.add(InputLayer(shape=(input_length, input_dimension)))
13 model.add(Conv1D(num_units_cnn, kernel_size=kernel_size,
14                  activation='relu', padding='same'))
15 model.add(Conv1D(num_units_cnn, kernel_size=kernel_size,
16                  activation='relu', padding='same'))
17 model.add(MaxPooling1D(pool_size=pool_size))
18 if flatten_before:
19     model.add(Flatten())
20
21 model.add(Dense(num_units, activation='relu'))
22 if not flatten_before:
23     model.add(Flatten())
24
25 model.add(Dense(...))
26
27 model.compile(optimizer=Adam(learning_rate=learning_rate), ...)
```

Listing I: Deep CNN

7.4.3 Attention CNN

An alternative architecture uses multiple parallel convolutional layers with different kernel sizes. This allows the model to detect both short-term and medium-term patterns simultaneously. The resulting feature maps are combined using an attention layer to highlight important sequence segments [68].

This is followed by global average pooling and a dense layer for output generation.

```
1 num_units_cnn = trial.suggest_int('num_units_cnn', 32, 128)
2 num_units = trial.suggest_int('num_units', 32, 128)
3 input_length = trial.suggest_int('input_length', 5, 150)
4 kernel_size_1 = trial.suggest_int('kernel_size_1', 2, 9)
5 kernel_size_2 = trial.suggest_int('kernel_size_2', 2, 9)
6 kernel_size_3 = trial.suggest_int('kernel_size_3', 2, 9)
7 learning_rate = trial.suggest_float('learning_rate', 1e-5, 1e-2,
8                                     log=True)
9
10 input_layer = Input(shape=(input_length, input_dimension))
11
12 conv_1 = Conv1D(num_units_cnn, kernel_size=kernel_size_1,
13                 padding='same', activation='relu')(input_layer)
14 conv_2 = Conv1D(num_units_cnn, kernel_size=kernel_size_2,
15                 padding='same', activation='relu')(input_layer)
16 conv_3 = Conv1D(num_units_cnn, kernel_size=kernel_size_3,
17                 padding='same', activation='relu')(input_layer)
18
19 concat = Attention()([conv_1, conv_2, conv_3])
20
21 gap = GlobalAveragePooling1D()(concat)
22 dense1 = Dense(num_units, activation='relu')(gap)
23
24 output_layer = Dense(...)(dense1)
25
26 model = Model(input_layer, output_layer)
27
28 model.compile(optimizer=Adam(learning_rate=learning_rate), ...)
```

Listing J: Attention CNN

7.4.4 Concatenation CNN

A slightly modified version replaces the attention module with a simple concatenation of the feature maps. This allows the model to pass the extracted features directly without performing weighting [69].

This variant is computationally less expensive than the attention-based one and is well-suited when all filter outputs should be treated equally.

```
1 num_units_cnn = trial.suggest_int('num_units_cnn', 32, 128)
2 num_units = trial.suggest_int('num_units', 32, 128)
3 input_length = trial.suggest_int('input_length', 5, 150)
4 kernel_size_1 = trial.suggest_int('kernel_size_1', 2, 9)
```

```
5 kernel_size_2 = trial.suggest_int('kernel_size_2', 2, 9)
6 kernel_size_3 = trial.suggest_int('kernel_size_3', 2, 9)
7 pool_size = trial.suggest_int('pool_size', 1, 5)
8 learning_rate = trial.suggest_float('learning_rate', 1e-5, 1e-2,
9                                     log=True)
10
11 input_layer = Input(shape=(input_length, input_dimension))
12 conv_1 = Conv1D(num_units_cnn, kernel_size=kernel_size_1,
13                 padding='same', activation='relu')(input_layer)
14 conv_2 = Conv1D(num_units_cnn, kernel_size=kernel_size_2,
15                 padding='same', activation='relu')(input_layer)
16 conv_3 = Conv1D(num_units_cnn, kernel_size=kernel_size_3,
17                 padding='same', activation='relu')(input_layer)
18
19 concat = Concatenate()([conv_1, conv_2, conv_3])
20
21 gap = GlobalAveragePooling1D()(concat)
22 dense1 = Dense(num_units, activation='relu')(gap)
23
24 output_layer = Dense(...)(dense1)
25
26 model = Model(input_layer, output_layer)
27
28 model.compile(optimizer=Adam(learning_rate=learning_rate), ...)
```

Listing K: Concatenation CNN

7.4.5 CNN-GRU Hybrid Model

The Gated Recurrent Unit (GRU) model is a simplified variant of the LSTM and was developed to reduce computational effort and model complexity. Compared to LSTM, GRU has only two gates, an update gate and a reset gate, instead of three (input, forget, and output gates). This makes GRU faster to train, requires less memory, and delivers comparable results in many tasks [70].

To combine the advantages of both architectures, a hybrid CNN-GRU model was also implemented. A Conv1D layer first extracts local features from the sequence, which are then sequentially processed using a GRU module to capture long-term dependencies and trends [71].

This model is particularly well-suited for linking local patterns (using CNN) with temporal dependencies (using GRU). The final dense layers perform the transformation to the target variable as usual.

```
1 num_units_cnn = trial.suggest_int('num_units_cnn', 32, 128)
2 num_units = trial.suggest_int('num_units', 32, 128)
3 num_units_gru = trial.suggest_int('num_units_gru', 32, 128)
4 input_length = trial.suggest_int('input_length', 5, 150)
5 kernel_size = trial.suggest_int('kernel_size', 2, 5)
6 pool_size = trial.suggest_int('pool_size', 1, 5)
7 learning_rate = trial.suggest_float('learning_rate', 1e-5, 1e-2,
8                                     log=True)
```

```
9
10 model = Sequential()
11 model.add(InputLayer(shape=(input_length, input_dimension)))
12 model.add(Conv1D(num_units_cnn, kernel_size=kernel_size,
13                 activation='relu'))
14 model.add(MaxPooling1D(pool_size=pool_size))
15 model.add(GRU(num_units_gru, return_sequences=False))
16 model.add(Dense(num_units, activation='relu'))
17 model.add(Dense(...))
18 model.compile(optimizer=Adam(learning_rate=learning_rate), ...)
```

Listing L: CNN + GRU

7.5 Model Training

In the following, all models mentioned in [Subsection 7.2](#), [Subsection 7.3](#), and [Subsection 7.4](#) are trained on the training data ([Table 4](#)) for 20 Optuna trials, each with 30 epochs. After each epoch, the models are evaluated on the validation data ([Table 5](#)).

As already mentioned, the model architectures are used for both regression and classification. For this purpose, a different final dense layer is used for each architecture, which is described in the respective chapters.

Similar to scaling and dimensionality reduction, the models are trained individually for each regime. This means that the training and validation data must be divided into 18 different subsets, and all model architectures are trained for 20 trials with 30 epochs using the corresponding subdata. This results in 10,800³ trained models per architecture. After each epoch, the evaluation metrics and the model are saved.

Another aspect that is determined by the regimes is how the input-output sequences can be divided. There are four possibilities here:

1. **Complete sequence within a regime:** In this case, the entire sequence, both the input and output data, would be within a single regime. This has the advantage that each model is truly assigned to only one regime and operates exclusively with data within that regime. However, this approach is impractical in live operation, as it is unknown whether the market will remain in this regime in the future in which the prediction is made. This would constitute data snooping, and the models would be useless.
2. **Only the output data within a regime:** The same problem arises here as in point 1. In live operation, it is unknown whether the market will remain in this regime in the future. Thus, data snooping occurs again, and the models become unusable.
3. **Only the input data within a regime:** It must be ensured that all input data lies within a regime. The number of regimes in the output data is irrelevant. This

³18Regimes * 20Trials * 30Epochs

ensures that the model always receives data that lies within a regime as input, but no data snooping occurs.

4. **The last N percent of the input data must be within a regime:** This has the advantage of eliminating data snooping and providing more test data compared to point 3, since more input sequences meet the requirements. However, it requires additional effort because the percentage N must be determined. This would require either a heuristic approach or a more complex process. Furthermore, compared to point 3, the model cannot guarantee that 100% of the input data is in the same regime and would therefore have to learn more patterns in the data to cope with regime changes within the input data.

For the reasons listed above, it was decided in this work to divide the data as described in point 3, since the distinction as to which model the data belongs is easy to implement, data snooping is excluded and no additional heuristics or processes have to be carried out to determine what percentage of the input data must belong to the respective regime.

7.6 Regression Models

Regression models were used to predict the next 30 minutes of logarithmic returns as a continuous sequence. To achieve a continuous sequence as output, the activation function in the last dense layer of the above-mentioned models was set to `linear` and the number of neurons to 30.

7.6.1 Loss-Function

In order to incorporate the trading context during model training and evaluation, traditional loss functions such as the root mean squared error have been avoided. Instead, a proprietary loss function was implemented that calculates realized profit and loss. The value of the loss function decreases the more profit is generated.

As mentioned above, the output of the regression models are either a sequence of the next expected logarithmic returns by minute. The first step of the loss function is to decide whether the opened position should be a long or short position. This is decided by cumulating the predictions. The absolute maximum is then determined for the cumulative values greater than and less than zero. If the absolute maximum of the values greater than zero is greater than the absolute maximum of the values less than zero, a long position is opened. Otherwise, a short position is opened. [Figure 9](#) illustrates the decision using a long position as an example.




Figure 9: Long Position Decision

After the direction decision, the stop-loss and take-profit levels must be determined. For long and short positions, there are two possibilities at which specific level the stop-loss and take-profit are set, based on the global high and low of the cumulative predicted logarithmic returns. If a long position has been previously determined and the global high occurs after the global low, the stop-loss is placed at the global low and the take-profit at the global high. However, if the global high occurs before the global low, the stop-loss is set at the lowest low before the global high. For short positions, the decision is exactly the opposite.

Figure 10 shows the stop-loss (red line) and take-profit (green line) levels for all four cases. The red areas mark the predicted unprofitable zones and the green areas mark the predicted profitable zones.




Figure 10: Long Position Decision

After the stop-loss and take-profit levels are determined, the actual logarithmic returns are used to calculate the profit or loss for each position. If neither the stop-loss nor the take-profit is reached, the profit calculation will by default return the loss when the stop-loss is reached to ensure that the model predicts the correct take-profit more often and is penalized if neither the take-profit nor the stop-loss is reached.

7.6.2 Metrics

Additionally, additional metrics in the trading context were added to the model during training. These include:

1. **Profit hit rate:** Percentage of trades that would have been profitable based on the prediction.
2. **Loss hit rate:** Percentage of trades that would have been unprofitable based on the prediction.
3. **None hit rate:** Percentage of trades that would have reached neither the stop loss nor the take profit based on the prediction.

When the three parts are added together, the result must be 1, since one of the events must occur in each prediction.

7.6.3 Model Evaluation

Despite the implementation and training of various deep learning models for accurate price prediction, combined with various technical indicators and a specially developed loss function for directly simulating trading decisions, none of the models achieved satisfactory predictive performance. The presented loss hit-rate metric was always exactly 1 for every model and in every market regime on both the validation and test data, meaning that every trade decision made based on the prediction was incorrect and closed with a loss.

The inadequate suitability of the models is also evident in an exemplary prediction. In one specific case, future prices were reconstructed based on the predicted logarithmic returns and the price known at the time. The calculation was carried out using the formula presented in [Subsection 6.2](#).

In reality, the ETH/USDC price was around USDC 1,800 at the given time. However, the model predicted such a negative return that the resulting price fell to almost USDC 0 within just 30 minutes. Such behavior is not only extremely unlikely under realistic market conditions, but economically virtually impossible, as it would be tantamount to a complete market collapse.

This incorrect forecast is not an isolated case, but rather symptomatic of the weaknesses of the regression models tested. It shows that the models are not only inaccurate but also tend to produce unrealistic extreme values, which would entail considerable risks in a trading context.

This chapter is dedicated to a critical analysis of possible causes.

A key reason for the poor model performance lies in the low predictability of short-term price movements. Minute-by-minute logarithmic returns are subject to a strong stochastic nature. They resemble a random walk and rarely exhibit robust, persistently usable patterns [72]. Thus, the models lack the structural foundation to make reliable predictions. However, many machine learning methods implicitly assume the existence of certain stationary or repeatable relationships, a circumstance that is severely limited with M1 returns.

A central methodological element of this work was the introduction of a modified loss function based on the simulation of trades. This function was explicitly designed to consider not only the statistical error but also the practical utility of a model within a trading strategy.

In practice, however, it became apparent that the model often "learned" to remain as inactive as possible to minimize losses, or developed strategies with high trading frequency in the presence of negative performance because it was unable to identify stable patterns of advantage. This behavior is due to a lack of positive feedback in the learning process: If no or only random winning patterns are present in the data, even the best loss function cannot enable meaningful optimization.

The poor model performance should not be interpreted as a failure of the method or the data selection, but rather as a relevant empirical result. It shows that applying deep learning to high-frequency cryptocurrency data does not automatically lead to profitable predictive models, even when modern model architectures, additional features, and specialized loss functions are used.

For the reasons mentioned above, the regression models will not be pursued further in the course of this work.

7.7 Classification-Models

Despite the poor results of the regression models, classification models were trained that predict an action to be performed. The action can be either buy, sell, or do nothing. For that, the model predicts a probability for each of the three classes. To achieve a direct comparison with the regression models, the same models as in [Subsection 7.6](#) were used. However, the main difference is that the final dense layer contains only three neurons, each representing an action, and uses the `softmax` activation function. The identical training and validation period was also used, and the training took place over 20 Optuna trials with 30 epochs each.

7.7.1 Loss-Function

For classification models, the categorical cross-entropy loss function was used. This choice is common and well-suited for classification problems where the model output represents a probability distribution across multiple classes. Mathematically, it is defined as follows:

$$L_{CCE} = - \sum_{i=1}^C y_i * \log(\hat{y}_i)$$

Where C denotes the number of classes, y_i the actual value (usually encoded as a one-hot vector), and \hat{y}_i the probability predicted by the model for class i . This function penalizes the model particularly severely if it assigns a high probability to an incorrect class and rewards correct predictions with high confidence.

Categorical cross-entropy measures the difference between the actual target distribution and the probability distribution predicted by the model [\[54\]](#).

In the context of trading, where the model must choose between three mutually exclusive actions, categorical cross-entropy is especially suitable. It encourages the model to make confident and accurate predictions by penalizing uncertain or incorrect outputs. This is important in trading, where executing the wrong action can lead to financial loss, and taking no action when an opportunity exists can result in missed profit. By using the loss function that directly evaluates the alignment between predicted probabilities and the true action, the model is guided toward making precise, confident decisions, which is a crucial requirement in financial markets, where clear and timely signals are essential for effective strategy execution.

7.7.2 Model Evaluation

[Table 9](#) shows the best performing classification models on the validation data, as well as their associated loss and accuracy on the validation and test data for each market regime.

Market Regime			Model	Loss		Accuracy	
				Validation	Test	Validation	Test
Down	Low	$Q_{0.33}$	Concatenation CNN	0.69	0.70	0.53	0.47
		$Q_{0.66}$	Concatenation CNN	0.70	0.70	0.55	0.52
		Q_1	Attention CNN	0.89	0.86	0.58	0.43
	High	$Q_{0.33}$	Attention CNN	0.78	0.89	0.53	0.51
		$Q_{0.66}$	Attention CNN	0.84	0.68	0.54	0.48
		Q_1	LSTM	0.81	0.79	0.55	0.49
Sideways	Low	$Q_{0.33}$	Attention CNN	0.75	0.63	0.52	0.49
		$Q_{0.66}$	Concatenation CNN	0.72	0.67	0.52	0.51
		Q_1	Attention CNN	0.95	0.66	0.55	0.50
	High	$Q_{0.33}$	Attention CNN	0.73	0.69	0.54	0.49
		$Q_{0.66}$	Attention CNN	0.99	0.76	0.58	0.51
		Q_1	Concatenation CNN	0.66	0.71	0.54	0.52
Up	Low	$Q_{0.33}$	Concatenation CNN	0.69	0.69	0.54	0.54
		$Q_{0.66}$	Concatenation CNN	0.69	0.69	0.54	0.51
		Q_1	Base CNN	0.75	0.65	0.53	0.49
	High	$Q_{0.33}$	Concatenation CNN	0.70	0.70	0.52	0.48
		$Q_{0.66}$	Attention CNN	0.81	0.95	0.53	0.48
		Q_1	BiLSTM	0.69	0.69	0.56	0.51

Table 9: Validation and Test Metrics for Classification Models

It is notable that some models perform very well on the validation data. On the test data, the model performance of all models is worse than on the validation data, but the metrics still show that a profitable trading strategy is possible.

As explained in [Subsection 5.5](#), a high win rate is not the only criterion that determines whether a trading strategy is successful.

It is also striking that identical model architectures perform significantly better in classification than in regression. This once again highlights that not only the model architecture, but also other factors, such as the loss function or the task, determine whether a model can learn from the training data.

8 Trading Strategies

In the context of automated trading, systematic trading strategies play a central role. They enable decisions to be made based on clearly defined rules or mathematical models, rather than on subjective assessments or human intuition.

A trading strategy defines a methodical approach by which financial instruments are bought or sold to achieve a specific goal. Typically, this is to maximize profit while limiting risk. Such strategies are often based on indicators from technical analysis or fundamental market information [73].

Essential components of a strategy are signal generation (entry/exit criteria), risk management (e.g., stop-loss, position sizes), and performance evaluation based on metrics such as cumulative profit, or maximum drawdown. The development process of a successful strategy typically comprises several phases, ranging from idea generation and backtesting to optimization and validation on previously unseen market data [74].

Similar to deep learning hyperparameters, trading strategies also involve adjustable parameters that must be defined before execution. In this work, these parameters are defined in a range, consisting of a minimal value, a maximum value, and a step size. This setup allows for a structured and repeatable parameter search, enabling the systematic exploration of different strategy configurations.

This chapter provides an overview of the trading strategies compared in this paper. Both the models trained in [Section 7](#) and classic trading strategies are presented. An overview of the metrics used to compare the trading strategies is also provided. The presented strategies also have their parameters defined, which will be tested later. It should be noted that each strategy also contains a parameter that defines whether positions are opened with a fixed stop level or a trailing stop. Since this parameter is included in all strategies, it is not listed every time.

8.1 AI Trading Strategies

AI trading strategies treat the model outputs differently depending on whether the underlying model is based on regression or classification. This variation arises from the fundamental nature of the models themselves and influences how their predictions are interpreted.

8.1.1 Regression AI Strategy

As described in [Subsection 7.6.3](#), the trained regression models are not suitable for developing a trading strategy. Nevertheless, this chapter describes a trading strategy that could be implemented using a predicted series of logarithmic returns. However, it is not evaluated or tested.

From the predicted series of logarithmic returns, the expected price in t minutes can be calculated by the formula mentioned in [Subsection 6.2](#).

Based on the predicted price, the stop-loss and take-profit can be determined identically to the loss-function of the regression models ([Subsection 7.6.1](#)). This creates an

entry signal with fixed stop-loss and take-profit level.

However, since it is possible that the price could exceed the stop level or the take-profit level might not be reached, additional parameters are introduced into the strategy that shift the two levels by a specific number of points. A third parameter is also introduced, which specifies a fixed distance to the stop-loss if no stop-loss has been predicted. This can happen, for example, if the initial prediction is positive and the price in the prediction does not fall below the current price.

Parameter Name	Min Value	Max Value	Step Size
Take-Profit Delta	-20.0	20.0	2.0
Stop-Loss Delta	-20.0	20.0	2.0
Stop-Loss not Predicted Delta	1.0	20.0	20.0

Table 10: AI Regression Model Strategy Parameters

8.1.2 Classification AI Strategy

As described in [Subsection 7.7](#) the classification models predict a probability for an action, which can be either buy, sell or do nothing. The strategy evaluates these predicted probabilities and takes action only when certain conditions are met. Specifically, a long or short entry signal is generated only if the corresponding predicted probability exceeds a predefined minimum threshold. This ensures that trades are only placed when the model is sufficiently confident in its prediction.

If neither the buy nor the sell probability reached the minimum threshold, the strategy chooses to do nothing or stay out of the market. This mechanism helps to filter out uncertain or low-confidence predictions, reducing the likelihood of unnecessary or potential unprofitable trades.

In addition to the action selection logic, the strategy also relies on predefined stop-loss and take-profit levels to manage risk and lock profits.

Parameter Name	Min Value	Max Value	Step Size
Take-Profit Distance	5.0	100.0	5.0
Stop-Loss Distance	5.0	100.0	5.0
Min. Probability for Entry	0.3	0.9	0.1

Table 11: AI Classification Model Strategy Parameters

8.2 Classic Trading Strategies

Technical analysis (TA) is based on the assumption that market movements do not develop randomly, but follow certain patterns that have repeated themselves in the past. In contrast to fundamental analysis, which deals with the intrinsic value of a financial instrument, technical analysis focuses exclusively on past price and volume movements to draw conclusions about future price developments.

Technical analysis focuses on visual and computational methods for identifying trend, support and resistance levels, and reversal points. These analysis are used to develop specific trading strategies that specifically respond to specific market behaviours. These strategies are often rule-based and can be implemented both manually and algorithmically [75]. Therefore, technical analysis strategies are more likely suitable for algorithmic trading.

This chapter describes three widely used technical analysis trading strategies. The strategies introduced are not the only classic strategies tested. They are the three that performed best in the tests. The other tested strategies are also briefly listed in [Subsection 16.4](#). However, some of them have been slightly modified in this work compared to the most widely used ones.

8.2.1 Dual Simple Moving Average Strategy

A common trading strategy is the simple moving average crossover strategy. It consists of two SMA with different periods. If the short-term SMA crosses the long-term SMA above, a long position is opened. Otherwise, if the short-term SMA crosses the long-term SMA below, a short position is opened [76]. [Figure 11](#) shows exemplary two moving averages for a synthetic price, with blue arrows marking buy entry signals and red arrows marking sell entry signals.




Figure 11: Dual Simple Moving Average Strategy Example

This strategy is straightforward to understand and implement and can help to identify possible trend changes which generate entry and exit signals. One of the greatest disadvantages is that the SMA is a lagging indicator and in quickly changing market conditions, the signals can be delayed. Especially in sideways market regimes, the SMAs can cross often, which can lead to false signals [77]. Here, the market regime recognition plays an important role. Through this, the strategy can be disabled in sideways regimes and enabled if the market has a clear trend.

Often the criteria for stop-loss and take-profit levels are percentage and risk-reward ratio based, which means that the stop-loss has a fixed distance, e.g., 2% and the take-profit is e.g. double the distance of the stop-loss distance (relative to the current price) [78]. To achieve a more adaptive stop-loss and take-profit level determination, a self-developed mechanism called swing detection was added.

This detection method identifies swings in price movements and differentiates between swing high and swing low points. A swing high point at time t is a point where the closing price is greater than all closing prices in the interval $[t - N; t + N]$. Similarly, a swing low point is a point where the closing price is lower than all closing prices in the same interval. This allows detecting simple support and resistance zones, on which the stop-loss can be set. Therefore, the strategy adapts to the current market behaviour.

The strategy has six parameters, consisting of the long-term and short-term SMA period, the order N of the swing detection, which determines how significant the swing has to be, the maximum age of the last swing point, as well as a delta for the stop-loss and take-profit levels (identically to those in Subsection 8.1.1). All parameter combinations where the short-term SMA period is greater than the long-term SMA period, have been filtered out and are not tested.

Parameter Name	Min Value	Max Value	Step Size
Short-Term SMA Period	1	10	1
Long-Term SMA Period	2	20	1
Swing Order N	1	10	1
Take-Profit Delta	0.0	200.0	10.0
Stop-Loss Delta	0.0	100.0	10.0

Table 12: Dual Simple Moving Average Strategy Parameters

8.2.2 Triple Exponential Moving Average Strategy

This strategy consists of three exponential moving averages (EMA) with different periods. The shortest EMA identifies short-term trends, while the medium EMA identifies medium-term trends, and the long EMA identifies long-term trend.

The common application of the triple exponential moving average strategy generates a buy signal if the short-term EMA crosses above the medium-term EMA. Additionally, the medium-term EMA must already be above the long-term EMA to validate the buy signal. For sell signals, the short-term EMA must cross below the medium-term EMA and the medium-term EMA has to be below the long-term EMA [79].

Therefore, the long-term EMA acts as a trend filter to reduce false positives mentioned in Subsection 8.2.1. In this work, the long-term EMA is also used as a trend filter but not by its position relative to two other EMA's. During the backtesting of the classic variant, it was noticed that many false entry signals were still generated, especially in sideways markets. Therefore, the trend was filtered using a minimal slope of the long-term EMA. The slope of the EMA is calculated by the slope of the regression line which is defined by the current EMA value (at time t) and the EMA value i minutes in the past.

$$Slope = \frac{EMA_t - EMA_{t-i}}{i}$$

If the slope is greater than a predefined threshold, only long positions can be opened. On the other hand, if the slope is less than the negated threshold, only short positions can be opened. Otherwise, the strategy cannot open any new positions.

Figure 12 shows the same synthetic price as in Figure 11. Additionally, a third EMA with period 20 is added, which is used to calculate the slope. The red-filled areas indicate zones where no positions can be opened. So the entry signals in these areas are ignored. In the green-filled areas, the entry signals are executed.




Figure 12: Triple Exponential Moving Average Strategy Example

In this strategy, the stop-loss and take-profit distance have been added as fixed parameters, which are also permuted. All parameter combinations, where the short-term EMA period is greater than the medium-term EMA period, or the medium-term EMA period is greater than the long-term EMA period, are filtered out.

Parameter Name	Min Value	Max Value	Step Size
Short-Term EMA Period	3	10	2
Medium-Term EMA Period	5	30	3
Long-Term EMA Period	10	50	5
Minimum EMA Slope	0.2	1.2	0.2
EMA Slope	10	40	10
Window Length i			
Stop-Loss Distance	10	100	10
Take-Profit Distance	10	150	10

Table 13: Triple Exponential Moving Average Strategy Parameters

8.2.3 Bollinger Bands Strategy

The Bollinger Bands strategy is a mean reversion strategy, designed to exploit short-term price fluctuations. This strategy assumes that after approaching or breaking through the outer bands, the price reverts to the central line (the moving average). A buy entry signal is generated if a candle opens below the lower Bollinger Band and closes above the lower Bollinger Band. This indicates a price reversal from an oversold condition. On the other hand, a sell entry signal is generated if the candle opens above the upper Bollinger Band and closes below the upper Bollinger Band [80].

Figure 13 shows the synthetic price, identically to Figure 11 and Figure 12. The blue-filled area shows the Bollinger Band, with the central line.




Figure 13: Bollinger Bands Strategy Example

The stop-loss level is set in a predefined distance relative to the current price. The take-profit level is set at the level of the central line, with a certain delta being added (for buy signals) or subtracted (for sell signals) to the value of the middle line.

Parameter Name	Min Value	Max Value	Step Size
Bollinger Band Period	10	25	1
No. of Standard Deviations	1.5	3.0	0.5
Stop-Loss Distance	10.0	100.0	20.0
Take-Profit Delta	-5.0	20.0	2.0

Table 14: Bollinger Band Strategy Parameters

9 Trading-Engine

For backtesting and live execution of trading strategies, a Java framework was built. This chapter describes the use cases as well as the most important components of the trading engine.

9.1 Use-Cases of the Trading-Engine

The developed trading engine represents a flexible and modular software platform that covers various use cases in the field of algorithmic trading. The focus is on combining usability, adaptability, and performance. The following describes the most important use cases and features:

1. **Backtesting trading strategies:** The trading engine enables the simulated execution of trading strategies on historical market data. This allows strategies to be tested under realistic conditions, and evaluate their performance, robustness, and risk characteristics. By replicating historical market conditions, incorrect decisions can be identified early, and the trading strategy can be adjusted without real capital.
2. **Live execution in real-time operation:** In addition to backtesting, the engine also supports real-time execution of strategies in the market. By an interface to brokers, the engine can receive market data in real time, make trading decisions, and execute orders directly. This enables the engine to be used as the basis for automated trading in production environments.
3. **Using money, and risk management strategies:** An important use case is the implementation of different money and risk management approaches. Users can integrate various strategies for position sizing, stop-loss setting, or profit-taking, and examine their impact on strategies performance in backtests or in live operation. This supports the development of more stable and profitable trading approaches.
4. **Connecting to any broker:** The trading engine is designed to connect to various brokers via interchangeable interfaces. This allows the engine to be combined with different trading systems and platforms. This flexibility enables its use in a wide variety of markets and infrastructures.
5. **Development, and integration of custom trading strategies:** A key feature of the engine is the ability for users to implement their own trading strategies. A clear separation between core functionality and strategy logic allows for flexible integration and testing of individual algorithms.

Most components of the trading engine are interchangeable, and can be implemented by the user. Only the core of the framework, including the fundamental process flow and control logic, remains fixed. This allows users to tailor the engine to their individual needs, develop their own modules for data feeds, order management, strategy, or risk control, and thus ensure a high degree of flexibility in the use, and further development of the platform.

9.2 Architecture

The trading engine consists of many modules which can be clustered in three main categories:

1. Core Modules (Trading Engine)
2. Applications
3. Adapters

Figure 14 shows the components of the trading engine.

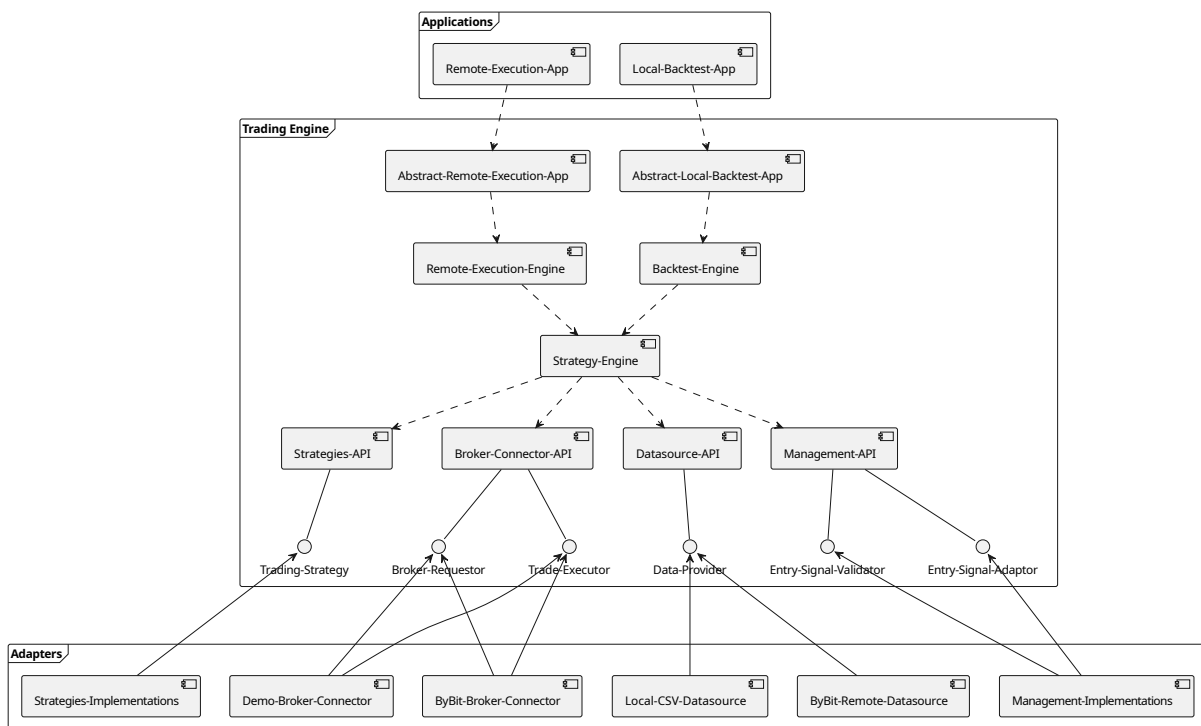


Figure 14: Trading-Engine Components

9.2.1 Plugin Architecture

To achieve the flexibility mentioned in [Subsection 9.1](#), a plugin architecture was used. Within the core, interfaces have been defined that define which external functionalities are required (Service provider interface, or SPI). To enable easy interchangeability by users, the implementations of the interfaces in the external modules (Adapters) are loaded by the Java `java.util.ServiceLoader`.

The `ServiceLoader` loads classes at runtime that implement a specific interface or abstract class. The classes to be loaded are specified via a configuration file in the `resources/META-INF/services` directory, which the `ServiceLoader` searches for in every

JAR file in the classpath. If a corresponding configuration file is found, the `ServiceLoader` can create an instance of the respective class using the default constructor [81].

As an example, a service provider interface for read access can be defined in a module `spi`.

```
1 // Contained in module spi
2 package com.example.spi.read;
3
4 public interface ReadRepository{
5     public DomainObject findAll();
6 }
```

Listing M: SPI Definition

This interface can be implemented in multiple modules. One implementation can be used to read data from files:

```
1 // Contained in module file-adapter
2 package com.example.adapter.file.read;
3
4 public class FileReadRepository implements ReadRepository {
5
6     @Override
7     public DomainObject findAll(){
8         // Logic to read a file
9     }
10 }
```

Listing N: File-Repository Implementation

Here, a configuration file named `com.example.spi.read.ReadRepository` is located in the directory `resources/META-INF/services`. This file contains the fully qualified class name of the implementation class (`com.example.adapter.file.read.FileReadRepository`).

Another implementation can be used to read data from databases:

```
1 // Contained in module db-adapter
2 package com.example.adapter.database.read;
3
4 public class DatabaseReadRepository implements ReadRepository {
5
6     @Override
7     public DomainObject findAll(){
8         // Logic to read a database
9     }
10 }
```

Listing O: Database-Repository Implementation

Similar to the `FileReadRepository` file named `com.example.spi.read.ReadRepository` is located in the directory `resources/META-INF/services`. However, this file contains `com.example.adapter.database.read.DatabaseReadRepository`.

The `ServiceLoader` searched each JAR in the classpath for files under `resources/META-INF/services` with the name `com.example.spi.read.ReadRepository` that contains the implementation of this interface, and creates them.

Since an external application has the ability to define the classpath, it can add or remove specific JARs. This allows it to determine which implementations to use. For example, it can decide to add only the `file-adapter` module or, if necessary, the `db-adapter` module if loading from a file, and a database simultaneously.

9.2.2 Core Modules

The core of the framework consists of several loosely coupled components that communicate with each other through well-defined APIs. The most central unit is the `Strategy-Engine` which contains the main logic for execution of trading strategies, and controls the cooperation of the other subsystems.

The following components form the core:

1. **Data-Provider:** Obtains market data by a configurable `Datasource-API`. The specific data source (local or remote) is integrated via adapters.
2. **Trading-Strategy:** The actual trading logic is integrated via the `Strategies-API`. This API is separate from the framework, and can be extended as required.
3. **Broker-Requestor & Trade-Executor:** Both communicate with brokers via a generic `Broker-Connector-API`, and enable both order placement, and broker status queries, such as the current account balance or all open positions.
4. **Entry-Signal-Adaptor & Entry-Signal-Validator:** These two components process, and validate generated entry signals. They are linked to the `Management-API` and enable additional risk checks, such as capital requirements or position limits. Currently the implemented management includes all techniques introduced in [Section 4](#) and is executed every time a strategy creates an entry signal.

Before the `Strategy-Engine` there are two other engines, named `Remote-Backtest-Engine`, and `Backtest-Engine`, which configure the `Strategy-Engine`. The difference between the engines is that the `Remote-Backtest-Engine` defines the `Strategy-Engine` asynchronously. This means that a remote data source can stream market data via an API. When new data arrives, the `Strategy-Engine` is notified by the data source, and the `Strategy-Engine` starts executing the trading strategy. For a local backtest, the `Strategy-Engine` is configured so that data synchronization is handled by the backtest engine, not by the data source.

Abstract apps are modules that accept and parse user configurations. For example, the task of the `Abstract-Local-Backtest-App` is to ask the user via the console which strategy should be used for the backtest. With the `Abstract-Remote-Execution-App`, the configuration is not done via the console, but primarily via environment variables. The parsed configurations are then passed to the respective engines.

9.2.3 Applications

The application layer defines specific execution environments by defining the classpath with adapters.

1. **Local-Backtest-App:** Executes trading strategies offline by using the `Local-CSV-Datasource` and the `Demo-Broker-Connector`.
2. **Remote-Execution-App:** Executes strategies in real time, and communicates with live broker interfaces.

Each application contains its own `main` method which constructs and starts the respective implementation.

9.3 Demo Broker

For local backtests, the connection to a real broker must be mocked. Therefore, the `Demo-Broker-Connector` is used to replace a real broker. He takes over the main tasks including:

1. **Order execution and position management:** Executing market, limit, take-profit, and stop-loss orders, including execution fee calculation.
2. **Adapt trailing stop positions:** Monitor trailing stop positions, and adapt the stop-level according to the most recent price.
3. **Account balance management:** Manage the available account balance as well as the margin balance.
4. **Store executed trades:** Storing the executed trades is essential for risk management and later analysis.

9.3.1 Order Execution and Position Management

In the context of trading order execution and position management are closely connected, because positions can be translated in three orders where only two of them are actually executed. As described in [Subsection 4.3](#) it is necessary that the trading engine supports market and limit orders. In real life environments, most orders are not executed at the expected price, due to slippage [82]. To simulate slippage, the trading engine can either use historical data of the underlying cryptocurrency, which is available in the lowest available timeframe (usually seconds or tick data), or simulate slippage randomly. Two problems exist with simulations using historical data. The first is that data in the seconds or tick range over a long period of time is difficult for private individuals to obtain. Furthermore, a type of latency must be simulated, which simulates the processing time and network traffic in live operation. With random simulations, the problem arises that the backtest results are no longer deterministic, thus a subsequent comparison of different strategies is

subject to error. For these reasons, the slippage factor is not taken into account during order execution. This means that market orders are always executed at the most current available price and limit orders are always executed at the specified order level.

Figure 15 shows the process of opening a single position. The margin in euro is calculated using the following formula [83].

$$Margin = \frac{PositionQuantity}{Leverage * EuroConversion} = \frac{PositionSize * OpenPrice}{Leverage * EuroConversion}$$

EuroConversion is the current price for converting the current counter currency in euro. For example, for ETH/USDC the *EuroConversion* is the current price of EUR/USDC: The conversion is done in euros because for different currency pairs with different counter currencies, the margin would be calculated in the respective counter currency. For ETH/USDC, the margin is calculated in USDC according to the formula. However, when trading ETH/BTC, the margin is calculated in BTC. To obtain a result that is always in the same unit, or in this case, the same currency, an additional conversion is implemented. It is important to note that it is also possible to implement conversion into any currency. This way, the margin could also be calculated in BTC, for example. Since USDC/EUR rates are offered by ByBit, it was easier in this case to obtain historical data directly from ByBit and obtain the EUR/USDC rate through a conversion $EUR/USDC = \frac{1}{USDC/EUR}$.

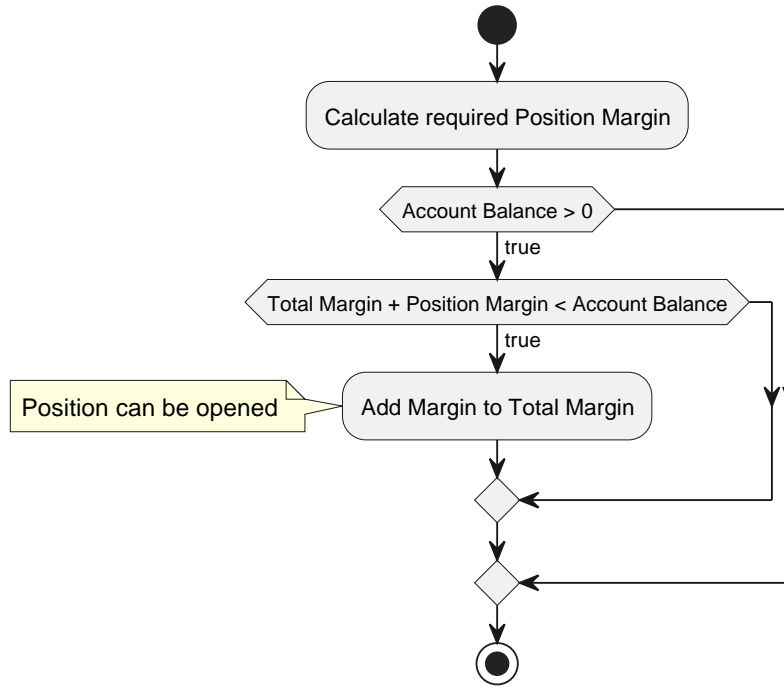


Figure 15: Opening a Single Position

Figure 16 shows the process of closing a single position. The profit in euro is calculated using the following formula [83]:

$$Profit = \frac{(ClosePrice - OpenPrice) * PositionSize}{EuroConversion} * \begin{cases} -1, & \text{if Sell-Position} \\ 1, & \text{otherwise} \end{cases}$$

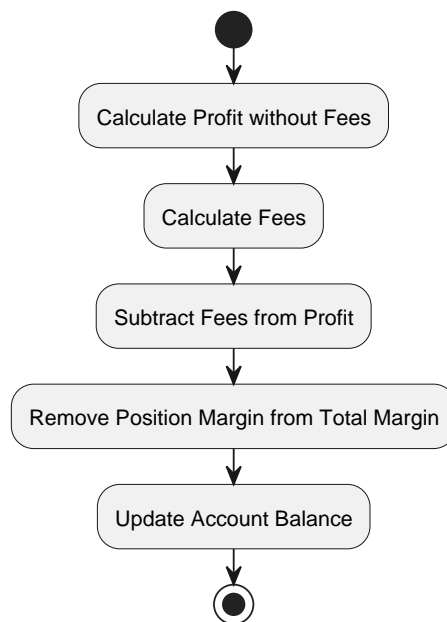


Figure 16: Closing a Position

Figure 17 shows the process of opening a position taking into account that other open positions can already exist.

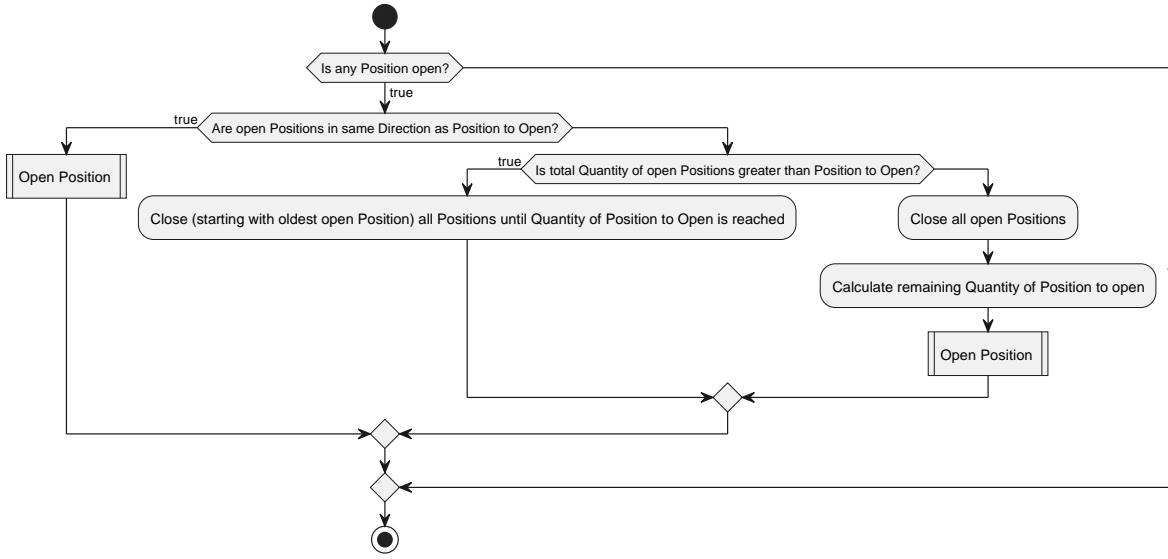


Figure 17: Opening a Position

9.3.2 Adapting Trailing Stop Positions

Trailing stop positions are positions with a dynamic stop order. The stop level moves automatically with the price as soon as it moves in the desired direction. For a buy position, the stop level moves if the price rises. If the price falls, the stop remains unchanged [84].

Figure 18 shows a synthetic closing price (black) and a trailing-stop order with 60 points distance (red) for a buy position. In the green-filled areas, the price moves in the desired direction, so the trailing-stop level follows the price. In the red-filled areas, the prices do not move in the desired direction, so the trailing-stop level does not move. An exception is the red-filled area after $t = 3$, where the price moves in the desired direction, but the trailing stop does not. This is because there is not yet a 60-point gap between the price and the stop level. Therefore, the stop only moves again once the 60-point gap is restored.




Figure 18: Trailing Stop Example

10 Backtesting Trading Strategies

The quantitative evaluation of trading strategies requires a structured and comprehensible testing environment which was introduced in [Section 9](#). This chapter describes the executed backtests, which aim to verify the performance of the develop trading strategies in [Section 8](#) under realistic conditions.

Backtesting describes a performance simulation of a trading strategy based on historical data. This involves investigating how a specific strategy would have performed in the past if it had been executed under the given market conditions. The goal is to gain initial insights into the robustness, profitability, and risk profile of the approach before it is used in live trading [\[85\]](#).

10.1 Trading Strategy Parameter Selection

In contrast to the classic machine learning process, the parameter selection for the ultimately executed strategies is performed without the training set. This also applies to strategies that use deep learning ([Subsection 8.1](#)). Unlike the training of deep learning models, which learn from the training data and adapt their internal states based on the training data, the trading strategies are initialized with fixed parameters. This corresponds to the fitting of deep learning models. However, the subsequent process is identical. The strategies are validated with their previously defined parameters on the validation set to find the best possible parameters.

10.2 Executing Backtests

Due to the large number of parameter combinations shown in [Table 15](#), backtesting all combinations is not possible due to the execution time. Therefore, only 1000 combinations per strategy were tested on the validation set. In order to cover the greatest possible variance of the parameters, the combinations were not sampled randomly from the set of all possible combinations, but the combinations were sorted in lexicographic order and sampled at constant intervals.

Strategy Name	Number of Parameters
Regression AI Strategy	8,820
Classification AI Strategy	5,600
Dual Simple Moving Average Strategy	98,000
Triple Exponential Moving Average Strategy	1,555,200
Bollinger Bands Strategy	8,320

Table 15: Number of Parameters per Strategy

10.3 Evaluating Backtests

Due to the length, the detailed results of the backtests ⁴, including the parameters ⁵ that led to the respective results, can be found in the appendix. Here, only the results of the strategies with the highest cumulative last equity combined across all market regimes are presented.

Strategy	Classification AI	Dual SMA	Triple EMA	Bollinger Bands
Number of Trades	267	3473	870	275
Maximum Gain	956	357	875	191
Maximum Loss	-245	-160	-135	-64
Profits Std.	280	37	105	41
Minimum	-2404	-23	-75	49
Equity Maximum	57908	15510	10717	2510
Last	55567	15159	9502	2510
Maximum Drawdown	4	149	126	51
Win Ratio	64	49	33	42

Table 16: Combined Strategy Results

As seen in [Table 16](#) the Classification AI strategy achieves by far the best cumulative results on the validation set, with the lowest maximum drawdown and the highest win ratio. But on the other hand, it also has the greatest loss.

To compare the out-of-sample performance ⁶, all strategies have also been tested with the best parameters on the validation set on the backtest set.

Strategy	Classification AI	Dual SMA	Triple EMA	Bollinger Bands
Number of Trades	173	5494	553	444
Maximum Gain	424	268	937	163
Maximum Loss	-144	-63	-207	-88
Profits Std.	92	9	179	32
Minimum	-2630	-4694	-948	-4378
Equity Maximum	2630	0	4048	0
Last	-2630	-4694	560	-4365
Maximum Drawdown	200	898314	131	87
Win Ratio	23	26	19	25

Table 17: Combined Out-of-Sample Strategy Results

⁴Detailed results in [Subsection 16.1](#)

⁵Detailed parameters per strategy and regime in [Subsection 16.3](#)

⁶Detailed out-of-sample results in [Subsection 16.2](#)

Table 17 shows that the strategies have a poor out-of-sample performance, resulting in either high losses or very volatile equity curves, which are not useful for real live trading. This shows that there is also a kind of overfitting in the parameter selection. While the strategies work well on the validation data, they fail in the out-of-sample tests.

10.4 Fee-Impacts

A key component of the realistic evaluation of algorithmic trading strategies is the consideration of transaction costs. These costs can have a significant impact on performance depending on trading frequency, broker model, and market conditions.

In this study, trading fees were integrated directly into the trading engine so that they are automatically considered with every order execution in backtesting. The chosen model is based on the ByBit fee structure which is based on the traded volume (0.02% for maker orders and 0.055% for taker orders). This enables realistic backtests of trading strategies. However, things like slippage were not taken into account but can be added in later improvements.

Subsection 10.4 shows the sum of all fees incurred with the best parameters on the validation data (in-sample) and on the test data (out-of-sample).

Strategy	In-Sample		Out-of-Sample
	Total Fees	Relative to PnL	Total Fees
Classification AI	3570	6.4%	1937
Dual SMA	4219	27.8%	2874
Triple EMA	3858	40.6%	4652
Bollinger Bands	1127	44.9%	1858

Table 18: Total Trading Fees

It becomes clear that a considerable amount of fees is incurred, which amounts to up to approximately 40% of the net profit (Triple EMA). However, when comparing the number of trades across strategies, an interesting insight emerges. As shown in Table 16, the Dual SMA strategy executes by far the most trades (3473 in total) on the validation set, but despite this high trading frequency, its relative fee impact is noticeable lower compared to the Triple EMA strategy, which executed only 870 trades. This indicates that not only the number of trades but also entry and exit timing, holding duration, and trade efficiency play a crucial role in determining fee sensitivity.

This highlights that fee efficiency is not strictly a function of trading frequency alone. While one might intuitively expect more trades to always lead to higher costs, the Dual SMA strategy demonstrated that a high-frequency system can still be cost-efficient if implemented with consistent execution logic and stable market conditions.

Most importantly, these findings emphasize the critical importance of incorporating fees in every phase of strategy evaluation. Conducting backtests without fees would draw an overly optimistic and misleading picture of strategy performance. For instance, a

strategy that appears profitable on a gross level may become unprofitable once realistic costs are accounted for.

For this reason, no backtests or performance evaluation in this work was conducted without accounting for trading fees. All strategies were tested under the assumption of realistic maker and taker fees, according to the selected broker ByBit. This ensures that every performance metric presented reflects realistic scenarios and therefore increases the practical relevance and robustness of the results.

11 Live-Testing

Despite the poor results from [Section 10](#), a strategy will be executed live on a demo account at ByBit for further work. The primary purpose of this live test is not to achieve profitability, but rather to validate the end-to-end functionality of the trading infrastructure under realistic market conditions.

11.1 Broker Setup

To connect to ByBit the `ByBit-Broker-Connector` and `ByBit-Remote-DataSource` components have been used.

The `ByBit-Remote-DataSource` is able to:

1. Request historical market data by the ByBit `Get Kline` API Endpoint [\[86\]](#).
2. Stream market data by the ByBit `Kline` websocket endpoint [\[87\]](#).

The `ByBit-Broker-Connector` is able to execute following actions via the ByBit API:

1. Request the current wallet balance as well as the current margin [\[88\]](#).
2. Placing Orders in the market, which includes open, take-profit, and stop-loss orders. Each Order is supported as limit and market order [\[89\]](#).
3. Request all open positions for one or multiple pairs [\[90\]](#). In the current paper this only includes positions on ETH/USDC M1.
4. Request closed trades in a specified time period [\[91\]](#).

The connection to ByBit is established by an API key and an API secret which can be created in the ByBit web interface. The distinction between a demo account and a real money account is regulated by ByBit via the base URL, which is different in each case. To use the demo depot, for REST endpoints, the base url `https://api-testnet.bybit.com` has to be used. For the real money account, the base url `https://api.bybit.com` has to be used [\[92\]](#). To stream the market data, no distinction is required because the base URL `wss://stream.bybit.com/v5/public/linear` [\[93\]](#) for streaming USDC market data is a public endpoint, which can be used without API Key. In the trading engine, the base URLs, API key, and API secret are set by environment variables to achieve an independent implementation, which is configured by the user.

11.2 Live-Test Results

Using the trading engine configuration described above, a live execution of the triple exponential moving average strategy was performed from 20th July, 2025, 05:00 UTC to 28th July, 2025, 05:00 UTC. To keep the environment as similar as possible to the local backtest environment, the balance of the demo account was set to 5000 USDC at startup.

Strategy		Triple EMA
Number of Trades		66
Maximum Gain		304
Maximum Loss		-228
Profits Std.		82
Minimum		3640
Equity	Maximum	5000
	Last	3913
Maximum Drawdown		27
Win Ratio		27

Table 19: Live-Test Statistics




Figure 19: Live-Test Results

As Table 19 and Figure 19 show, the strategy does not deliver profitable results in the live test either.

The logs ⁷ show that the connection (data streaming) was interrupted several times during execution by the broker or due to network problems. In these cases, the trading engine attempts to reconnect every ten seconds. As soon as the connection is re-established, the engine continues to run normally. The only limitation is that the internal state of the currently executing strategies is not reset, as this function does not yet exist. This is particularly relevant if a connection interruption causes several candles to be skipped, making it appear to the strategies as if there was a jump in the data. A mechanism must

⁷The logs are not published in this work due to their length.

therefore be built in that detects gaps in the data and resets the strategies. In the test carried out, however, this did not pose a problem, as the connection interruptions never lasted longer than a few seconds and no candles have been missed.

12 Dealing with Strategy-Overfitting

As has become clear in previous chapters, overfitting is one of the biggest problems with the strategies presented. This chapter therefore deals with a self-developed idea of how strategies could be adapted to the data with less risk of overfitting. It is important to note, that this is merely an idea that has not yet been tested. This would have to be addressed in future work.

It is striking that the strategies performed very well in the in-sample tests. However, they completely failed in the out-of-sample tests. This observation leads to the assumption that the parameters found are too specific to generalize well. Therefore, the conclusion was reached that the search for stable parameter spaces is more important than the search for individual parameter combinations, which at first glance, produce better results, but are not stable. Based on this insight, this chapter describes a concrete workflow for how a search for parameter spaces could work that simultaneously delivers stable and profitable results.

First, as many parameter combinations as possible should be tested through backtests. Metrics needed for the next step should be recorded for each of these. These metrics can be taken from [Section 5](#), for example, or expanded as desired. It should be noted that the larger the number of parameters, the more backtests should be performed to cover the entire parameter space as well as possible.

In the next step, optionally, a combined metric f can be created from the recorded metrics, which includes multiple aspects. This could, for example, look like this:

$$\begin{aligned} f = & + 0.25 \cdot \text{Rank}(\text{LastEquity}) \\ & + 0.15 \cdot \text{Rank}(\text{MaxEquity}) \\ & + 0.10 \cdot \text{Rank}(\text{WinRatio}) \\ & + 0.10 \cdot \text{Rank}(\text{NumTrades}) \\ & - 0.15 \cdot \text{Rank}(\text{MaxDrawdown}) \\ & - 0.10 \cdot \text{Rank}(\text{MaxLoss}) \\ & - 0.10 \cdot \text{Rank}(\text{ProfitsStd}) \\ & - 0.05 \cdot \text{Rank}(\text{MinEquity}) \end{aligned}$$

Here $\text{Rank}(X)$ is the ranking of X across all tested combinations, normalized to $[0; 1]$. This metric incorporates all metrics that were also compared in the previous strategies, eliminates scale differences, and considers whether a metric should be maximized or minimized. This results in one value for each parameter combination and backtest, which can be used to search for plateaus, simplifying the subsequent steps without losing any information.

The best 10-20% of the backtest results are then filtered using the combination metric. This has the advantage of allowing unwanted plateaus to be filtered out in advance. These unwanted plateaus can arise when various parameter combinations produce poor results, resulting in plateaus in undesirable ranges. Thus, filtering can reduce the risk of finding false plateaus. The best 10-20% range was selected because a smaller range may simply

contain many peaks in the data. A larger range can dilute the structures because average results are also included. Thus, the 10-20% range offers a good trade-off between average results and results affected by overfitting.

The next and most important step is the actual plateau search. Several methods could be used for this. The first approach involves dimensional reduction of the N -dimensional space⁸ followed by clustering. The second approach uses mathematical neighborhood analysis using gradients and variances to find smooth regions. Both approaches have advantages and disadvantages. Clustering is usually more efficient than neighborhood analysis for large N [94] and, if the correct dimensional reduction and clustering methods are chosen, it does not lose any information. However, this is also a disadvantage. If the wrong dimensional reduction method is chosen, the neighborhoods of the points may not be preserved [94], which can lead to misinterpretations of the results. At the same time, the clusters found must be transformed back into the higher-dimensional space to recover the strategy parameters, which can lead to errors. The second approach, using mathematical neighborhood analysis, is more computationally intensive, but it eliminates the need for back transformation and prevents any loss of information, as the calculations are performed directly in the original space. For this reason, it was decided to explain the second way in more detail.

The basic idea is to first find plateau candidates and then cluster them. A plateau candidate is a point that exhibits a small gradient and low variance with respect to its k nearest neighbors. The following algorithm describes how, for each tested parameter combination $x_i \in \mathbb{R}^N$ with the associated performance metric $f(x_i)$, a point is determined to be a plateau candidate.

1. **Knn-Neighborhood:** Find the k nearest points $\{x_j\}$ to x_i .
2. **Local gradient:** For x_i , calculate the local gradient g_i with

$$g_i = \frac{1}{k} * \sum_{j=i}^k \frac{|f(x_i) - f(x_j)|}{\|x_i - x_j\| + \epsilon}$$

3. **Local variance:** For x_i , calculate the local variance v_i with

$$v_i = \text{Var}(f(x_j) | x_j \in \text{KNN}(x_i))$$

4. **Plateau candidates:** Select only those x_i for which

$$f(x_i) \geq Q_{1-\alpha}, \quad g_i \leq \gamma, \quad v_i \leq \delta$$

where $Q_{1-\alpha}$ is the $(1 - \alpha)$ -quantile (e.g., $\alpha = 0.2$), and γ, δ are small limits (e.g., median of g_i, v_i)

5. **Clustering:** Connect all x_i from step 4 that are located in each other's KNNs (Connected Components or DBSCAN). Each cluster is a plateau.

⁸ N corresponds to the number of strategy parameters

Subsequently, the mean \bar{f} and the stability score $S = \frac{\bar{f}}{std(f)}$ must be calculated for each plateau. S can be interpreted that the highest possible combined metric (\bar{f}) and the lowest possible standard deviation within the cluster lead to a larger S . This means that a good strategy result, in the sense of the combined metric and high stability have been achieved simultaneously. Thus, the cluster with the highest S can be selected as the best cluster C^* .

As a final step, the strategy's final parameters must be determined from the cluster. For this purpose, for example, the best parameter vector $\vec{p}_{i^*} \in C^*$ with the highest individual f_i can be selected, which leads to the highest performance but potentially a lower tolerance. Alternatively, the mean vector \vec{p}_{center} of the cluster can be selected.

$$\vec{p}_{\text{center}} = \frac{1}{|C^*|} * \sum_{i \in C^*} \vec{p}_i$$

This vector has the maximum distance to the edges of the plateau and thus the highest tolerance to parameter fluctuations.

Thus, this workflow yields a final parameter vector that is contained in a more stable parameter space, which probably yields worse in-sample results, but generalizes significantly better than the best parameter combination on the in-sample data.

13 Conclusion

This paper investigates the performance of AI-based and classic trading strategies for the ETH/USDC market on a minute-by-minute basis. The aim was to compare both neural networks (NN, LSTM, CNN) and rule-based strategies (moving averages, Bollinger Bands) with regard to their suitability for high-frequency crypto trading.

Based on historical market data obtained from the broker ByBit from 5th August, 2022, 10:00 UTC to 17th June, 2025, 11:30 UTC, an exploratory data analysis was initially conducted. The raw prices were transformed into logarithmic returns to eliminate data drift and supplemented with technical indicators such as trend, volatility, and momentum indicators.

Subsequently, scaling using MinMaxScaler and a PCA reduction to the principal components, which explain 80% of the variance, were performed for each market regime. This resulted in a manageable feature set for the subsequent models.

To dynamically adapt to different market regimes, a regime detection algorithm was developed that uses two moving averages (SMAs) and the rolling standard deviation to classify the market environment into trend and volatility phases.

Additionally, various money and risk management techniques were evaluated. These include:

1. Position sizing based on the stop-loss distance.
2. Minimum risk-reward ratio.
3. Maximum allowable account risk.
4. Minimum take-profit to cover trading fees.

To further minimize costs, market orders were automatically converted into limit orders.

Hyperparameter optimization with Optuna played an important role in model training. However, regression models designed to predict future logarithmic returns over the next 30 minutes exhibited severe overfitting and produced economically implausible results. Classification models designed to predict buy, sell, or no action achieved on the test-set win rates minimally better than a random model and resulted in losses in simulations. In parallel, classic strategies such as dual SMA, triple EMA, and Bollinger Bands were implemented and their performance compared.

The developed trading strategies were implemented in a modular Java-based trading engine that supports backtests and live trading. The modular structure makes it possible to implement specific trading strategies, adapters to different brokers, and money and risk management mechanisms without the need to adapt the core logic.

While all trading strategy approaches showed excellent results in the in-sample back-test, their performance collapsed significantly in the out-of-sample test, leading to losses or highly volatile equity curves.

A subsequent live test primarily served to verify the technical stability of the trading engine. The engine ran without interruption, even in the event of connection interruption

by the broker. The executed triple EMA strategy did not generate any significantly profitable signals.

Overall, the work illustrates that high-frequency ETH/USDC data contains hardly any predictable patterns for the model architectures examined. Both AI models and classical approaches suffer from overfitting, which requires robust optimization methods and careful validation. While the PCA transformation facilitates model creation, it can dilute important indicator signals. Despite the lack of profitability, the work provides valuable insights into data preprocessing, regime detection, risk management, and the development of a stable trading infrastructure.

Overall, the work demonstrates that high-frequency ETH/USDC data contains hardly any predictable pattern for the model architectures examined. Both AI models and classical approaches suffer from overfitting, which requires robust optimization methods and careful validation. While PCA transformation facilitates model building, it can misrepresent important indicator signals. Despite its lack of profitability, the work provides valuable insights into data preprocessing, regime detection, risk management, and the development of a stable trading infrastructure.

14 Aim of further Works

The challenges and limitations identified in this work open up many concrete and promising starting points for further research and development projects. Particularly with regard to alternative data sources, model architectures and improved validation strategies, opportunities arise to specifically address existing weaknesses while simultaneously exploiting the full potential of data-driven trading strategies in the cryptocurrency market. The following concrete approaches can form the basis for future work:

1. **Expanding the feature set:** The integration of order book data (e.g., depth snapshots or order flow imbalance), sentiment indices from social media and news, and macroeconomic indicators can provide uncorrelated information and improve forecasting capabilities.
2. **Alternative model architectures:** Reinforcement learning agents (e.g., DQN, PPO) could learn to dynamically allocate between long, short, and cash positions. Transformer models with attention mechanisms offer the possibility of specifically weighting relevant market phases. Ensemble methods from multiple model types promise to reduce overfitting.
3. **Meta-strategies and target size adjustment:** Instead of exact return predictions, models could only forecast the trading direction or volatility classes to increase stability. A meta-strategy that switches daily between different sub-strategies based on their current performance expectations can create additional robustness.
4. **Robust optimization and validation:** Monte Carlo cross-validation with random parameter starts and scenario variation, combined with restrictive regularization (ElasticNet, DropConnect) and data augmentation, can further mitigate overfitting.

5. **Diversification of markets and time frames:** Checking trading signals on different timeframes (5-minute, 15-minute, hourly) and in other currency pairs (e.g., BTC/USDC, altcoins) enables the analysis of correlation structures and low-risk diversification.
6. **Composite strategies:** Entry confirmations should only occur when multiple algorithms issue a signal simultaneously. Portfolio-level management that weights strategies according to risk budget and other risk metrics can stabilize overall returns.
7. **Stress tests and scenario analyses:** Simulating extreme market events such as flash crashes or liquidity shortages, including slippage and order execution models, provides important insights into strategy robustness.

With these extensions on the stable technical infrastructure developed in this work, it could be possible to design a significantly more robust and better generalizing trading approach. The insights gained in modeling, regime detection, risk management, and strategy integration compared with the provided suggestions for further researches form a solid foundation for the development of practice-relevant algorithms that can deliver consistent results not only in backtests but also in live operation.

15 References

- [1] TotalBitcoin.org - Ethereum: The Birth and Evolution of the Smart Contract Revolution. 2024. URL: <https://totalbitcoin.org/ethereum-evolution/> (visited on 06/18/2025).
- [2] theknowledgeacademy.com - Advantages and Disadvantages of Ethereum. 2024. URL: <https://www.theknowledgeacademy.com/blog/advantages-and-disadvantages-of-ethereum/> (visited on 06/18/2025).
- [3] Investopia.com - USD Coin (USDC): Definition, How It Works in Currency, and Value. 2024. URL: <https://www.investopedia.com/usd-coin-5210435> (visited on 06/18/2025).
- [4] Financial Times - Stablecoins are bank deposits. 2025. URL: <https://www.ft.com/content/625f9f0e-0144-4bc4-99e2-571ab64ccbe9> (visited on 06/18/2025).
- [5] circle.com - USDC. 2025. URL: <https://www.circle.com/usdc> (visited on 06/18/2025).
- [6] binance.com - What is the Outlook for the Crypto Brokerage Industry? 2020. URL: <https://www.binance.com/en/research/analysis/what-is-the-outlook-for-the-crypto-brokerage-industry/> (visited on 06/18/2025).
- [7] Investing.com - Die besten Krypto Broker – Juni 2025. 2025. URL: <https://de.investing.com/brokers/cryptocurrency-brokers/#:~:text=Zu%20den%20wichtigsten%20Aufgaben%20eines,an%20Kryptow%C3%A4hrungen%20als%20einzelne%20B%C3%B6rsen.> (visited on 06/18/2025).
- [8] ByBit.com. 2025. URL: <https://www.bybit.com/en/> (visited on 06/18/2025).
- [9] ByBit API Doc. 2025. URL: <https://www.bybit.com/en/%20https://bybit-exchange.github.io/docs/v5/intro> (visited on 06/18/2025).
- [10] IG.com. 2025. URL: https://www.ig.com/de/kryptowaehrungshandel?_gl=1*15gtia3*_up*MQ..*_gs*MQ..&gclid=CjwKCAjwpMTCBhA-EiWA_-MsmST2-xxZ7Kxhnw1PDrrdvL7kdRSPmkn0M0R2x2ShxY2-5F5G5_mIFBoCUikQAvD_BwE&gclsrc=aw.ds&gbraid=0AAAAADzDEsR-vDC7urfY10DmU7i3x7bQU (visited on 06/18/2025).
- [11] IG API Doc. 2025. URL: <https://labs.ig.com/getting-started.html> (visited on 06/18/2025).
- [12] Capital.com. 2025. URL: <https://capital.com/de-de> (visited on 06/18/2025).
- [13] Capital API Doc. 2025. URL: <https://open-api.capital.com/#section/General-information> (visited on 06/18/2025).
- [14] ByBit API Doc - Get Kline. 2025. URL: <https://bybit-exchange.github.io/docs/v5/market/kline> (visited on 06/18/2025).

- [15] Macrosynergy - Classifying market regimes. 2025. URL: <https://macrosynergy.com/research/classifying-market-regimes/> (visited on 06/24/2025).
- [16] Felix Haase and Matthias Neuenkirch. "Predictability of bull and bear markets: A new look at forecasting stock market regimes (and returns) in the US". In: *International Journal of Forecasting* 39.2 (2023), pp. 587–605. DOI: <https://doi.org/10.1016/j.ijforecast.2022.01.004>.
- [17] Arjun Prakash et al. "Structural Clustering of Volatility Regimes for Dynamic Trading Strategies". In: *Applied Mathematical Finance* 28.3 (May 2021), pp. 236–274. DOI: [10.1080/1350486x.2021.2007146](https://doi.org/10.1080/1350486x.2021.2007146).
- [18] Szabolcs Mike and J. Doynne Farmer. An empirical behavioral model of liquidity and volatility. 2007. arXiv: [0709.0159](https://arxiv.org/abs/0709.0159) [q-fin.ST].
- [19] IG.com - Moving averages: a guide to trend trading. 2025. URL: <https://www.ig.com/en/trading-strategies/moving-averages--a-guide-to-trend-trading-241031> (visited on 06/25/2025).
- [20] wikipedia.org - Volatility (finance). 2025. URL: [https://en.wikipedia.org/w/index.php?title=Volatility_\(finance\)&oldid=1291878143](https://en.wikipedia.org/w/index.php?title=Volatility_(finance)&oldid=1291878143) (visited on 06/25/2025).
- [21] brokercheck.eu - Risk Management for Traders: An Essential Guide. 2025. URL: <https://brokercheck.eu/risk-management-traders> (visited on 06/23/2025).
- [22] BritannicaMoney - Calculating position size in trading: The key to risk management. 2024. URL: <https://www.britannica.com/money/calculating-position-size> (visited on 06/23/2025).
- [23] investopia.com - Risk/Reward Ratio: What It Is, How Stock Investors Use It. 2024. URL: <https://www.investopedia.com/terms/r/riskrewardratio.asp> (visited on 06/23/2025).
- [24] bitpanda - Risk reward ratio: what is it & why is it important? 2024. URL: <https://www.bitpanda.com/academy/en/lessons/risk-reward-ratio-what-is-it-and-why-is-it-important/> (visited on 06/23/2025).
- [25] financemagnates.com - Understanding Liquidity and Market Liquidity. 2022. URL: <https://www.financemagnates.com/forex/education-centre/understanding-liquidity-and-market-liquidity/> (visited on 06/23/2025).
- [26] quantifiedstrategies.com - Trading Performance: Strategy Metrics, Risk-Adjusted Metrics, And Backtest. 2025. URL: <https://www.quantifiedstrategies.com/trading-performance/> (visited on 07/15/2025).
- [27] investopia.com - Maximum Drawdown (MDD): Definition and Formula. 2025. URL: <https://www.investopedia.com/terms/m/maximum-drawdown-mdd.asp> (visited on 07/15/2025).
- [28] investopia.com - Drawdown: What It Is, Risks, and Examples. 2025. URL: <https://www.investopedia.com/terms/d/drawdown.asp> (visited on 07/15/2025).

- [29] Hao Jiang and Bryan Kelly. “Tail Risk and Hedge Fund Returns”. In: *SSRN Electronic Journal* (Mar. 2012). DOI: [10.2139/ssrn.2021942](https://doi.org/10.2139/ssrn.2021942).
- [30] trustnet.com - What is maximum gain? 2024. URL: <https://www.trustnet.com/investing/13425786/what-is-maximum-gain> (visited on 07/16/2025).
- [31] investopia.com - Standard Deviation vs. Variance: What’s the Difference? 2025. URL: <https://www.investopedia.com/ask/answers/021215/what-difference-between-standard-deviation-and-variance.asp> (visited on 07/16/2025).
- [32] binance.com - Win Rate. 2025. URL: <https://academy.binance.com/en/glossary/win-rate> (visited on 07/16/2025).
- [33] gregorygundersen.com - Returns and Log Returns. 2022. URL: <https://gregorygundersen.com/blog/2022/02/06/log-returns/> (visited on 06/20/2025).
- [34] Peter C. B. Phillips. “Unit Roots”. In: *Macroeconometrics and Time Series Analysis*. Ed. by Steven N. Durlauf and Lawrence E. Blume. London: Palgrave Macmillan UK, 2010, pp. 347–368. DOI: [10.1057/9780230280830_37](https://doi.org/10.1057/9780230280830_37).
- [35] Investiopia.com - Trend Trading: The 4 Most Common Indicators. 2025. URL: <https://www.investopedia.com/articles/active-trading/041814/four-most-commonlyused-indicators-trend-trading.asp> (visited on 06/19/2025).
- [36] Investiopia.com - What is EMA? How to Use Exponential Moving Average With Formula. 2024. URL: <https://www.investopedia.com/terms/e/ema.asp> (visited on 06/19/2025).
- [37] Investiopia.com - What Is MACD? 2024. URL: <https://www.investopedia.com/terms/m/macd.asp> (visited on 06/19/2025).
- [38] Investiopia.com - Average True Range (ATR) Formula, What It Means, and How to Use It. 2025. URL: <https://www.investopedia.com/terms/a/atr.asp> (visited on 06/19/2025).
- [39] Investiopia.com - Bollinger Bands: What They Are, and What They Tell Investors. 2024. URL: <https://www.investopedia.com/terms/b/bollingerbands.asp> (visited on 06/19/2025).
- [40] Investiopia.com - What Is a Momentum Indicator? Definition and Common Indicators. 2025. URL: <https://www.investopedia.com/investing/momentum-and-relative-strength-index/> (visited on 06/19/2025).
- [41] Investiopia.com - Relative Strength Index (RSI) Indicator Explained With Formula. 2024. URL: <https://www.investopedia.com/terms/r/rsi.asp> (visited on 06/19/2025).
- [42] Bharathi Deepa and K. Ramesh. “Epileptic seizure detection using deep learning through min max scaler normalization”. In: *International journal of health sciences* (2022).
- [43] scikit-learn.org - Importance of Feature Scaling. 2025. URL: https://scikit-learn.org/stable/auto_examples/preprocessing/plot_scaling_importance.html (visited on 06/20/2025).

- [44] scikit-learn.org - MinMaxScaler. 2025. URL: <https://scikit-learn.org/stable/modules/generated/sklearn.preprocessing.MinMaxScaler.html> (visited on 06/20/2025).
- [45] Yoshua Bengio. “Practical recommendations for gradient-based training of deep architectures”. In: *CoRR* abs/1206.5533 (2012). arXiv: [1206.5533](https://arxiv.org/abs/1206.5533).
- [46] Wikipedia.org - Principal component analysis. 2024. URL: https://en.wikipedia.org/w/index.php?title=Principal_component_analysis&oldid=1296008542 (visited on 06/20/2025).
- [47] Wikipedia.org - Curse of dimensionality. 2024. URL: https://en.wikipedia.org/w/index.php?title=Curse_of_dimensionality&oldid=1296421397 (visited on 06/20/2025).
- [48] keras.io. 2025. URL: <https://keras.io/> (visited on 07/06/2025).
- [49] tensorflow.org - Windows Install. 2025. URL: <https://www.tensorflow.org/install/pip#windows-native> (visited on 07/06/2025).
- [50] Jasman Pardede and Khairul Rijal. “The Effect of Hyperparameters on Faster R-CNN in Face Recognition Systems”. In: *Jurnal RESTI (Rekayasa Sistem dan Teknologi Informasi)* 9 (May 2025), pp. 436–448. DOI: [10.29207/resti.v9i3.6405](https://doi.org/10.29207/resti.v9i3.6405).
- [51] optuna.io - Basic Concepts. 2025. URL: <https://optuna.readthedocs.io/en/stable/index.html> (visited on 07/06/2025).
- [52] Yosua Harmoni, Kartika Maulida Hindrayani, and Dwi Prasetya. “Optimizing Categorical Boosting Model with Optuna for Anti-Tuberculosis Drugs Classification”. In: *Indonesian Journal of Electronics, Electromedical Engineering, and Medical Informatics* 7 (May 2025), pp. 401–414. DOI: [10.35882/ijeeemi.v7i2.92](https://doi.org/10.35882/ijeeemi.v7i2.92).
- [53] Saleem Abdullah, Ihsan Ullah, and Fazal Ghani. “Heterogeneous wireless network selection using feed forward double hierarchy linguistic neural network”. In: *Artificial Intelligence Review* 57 (July 2024). DOI: [10.1007/s10462-024-10826-y](https://doi.org/10.1007/s10462-024-10826-y).
- [54] Maria Vakalopoulou et al. “Deep Learning: Basics and Convolutional Neural Networks (CNNs)”. In: *Machine Learning for Brain Disorders*. Ed. by Olivier Colliot. New York, NY: Springer US, 2023, pp. 77–115. DOI: [10.1007/978-1-0716-3195-9_3](https://doi.org/10.1007/978-1-0716-3195-9_3).
- [55] keras.io/flatten. 2025. URL: https://keras.io/api/layers/reshaping_layers/flatten/ (visited on 07/08/2025).
- [56] Sergey Ioffe and Christian Szegedy. “Batch Normalization: Accelerating Deep Network Training by Reducing Internal Covariate Shift”. In: *CoRR* abs/1502.03167 (2015). arXiv: [1502.03167](https://arxiv.org/abs/1502.03167).
- [57] keras.io/dropout. 2025. URL: https://keras.io/api/layers/regularization_layers/dropout/ (visited on 07/08/2025).
- [58] Daoping Du et al. “Underwater image restoration based on a dual-branch super-resolution residual network”. In: *Multimedia Systems* 31 (June 2025). DOI: [10.1007/s00530-025-01861-y](https://doi.org/10.1007/s00530-025-01861-y).

- [59] Thomas Fischer and Christopher Krauss. “Deep learning with long short-term memory networks for financial market predictions”. In: *European Journal of Operational Research* 270.2 (2018), pp. 654–669. DOI: <https://doi.org/10.1016/j.ejor.2017.11.054>.
- [60] Thorir Mar Ingolfsson. 2021. URL: https://thorirmar.com/post/insight_into_lstm/ (visited on 07/08/2025).
- [61] Greg Van Houdt, Carlos Mosquera, and Gonzalo Nápoles. “A review on the long short-term memory model”. In: *Artificial Intelligence Review* 53.8 (2020), pp. 5929–5955. DOI: [10.1007/s10462-020-09838-1](https://doi.org/10.1007/s10462-020-09838-1).
- [62] Yuxin Wang. “Advanced Network Traffic Prediction Using Deep Learning Techniques: A Comparative Study of SVR, LSTM, GRU, and Bidirectional LSTM Models”. In: *ITM Web of Conferences* 70 (Jan. 2025). DOI: [10.1051/itmconf/20257003021](https://doi.org/10.1051/itmconf/20257003021).
- [63] keras.io/repeat_vector. 2025. URL: https://keras.io/api/layers/reshaping_layers/repeat_vector/ (visited on 07/08/2025).
- [64] keras.io/time_distributed. 2025. URL: https://keras.io/api/layers/recurrent_layers/time_distributed/ (visited on 07/08/2025).
- [65] Pedro Lara-Benítez, Manuel Carranza-García, and José C. Riquelme. “An Experimental Review on Deep Learning Architectures for Time Series Forecasting”. In: *CoRR* abs/2103.12057 (2021). arXiv: [2103.12057](https://arxiv.org/abs/2103.12057).
- [66] Hans Weytjens and Jochen De Weerd. “Process Outcome Prediction: CNN vs. LSTM (with Attention)”. In: *CoRR* abs/2104.06934 (2021). arXiv: [2104.06934](https://arxiv.org/abs/2104.06934).
- [67] Karen Simonyan and Andrew Zisserman. Very Deep Convolutional Networks for Large-Scale Image Recognition. 2015. arXiv: [1409.1556](https://arxiv.org/abs/1409.1556) [cs.CV].
- [68] Ashish Vaswani et al. “Attention Is All You Need”. In: *CoRR* abs/1706.03762 (2017). arXiv: [1706.03762](https://arxiv.org/abs/1706.03762).
- [69] keras.io/concatenate. 2025. URL: https://keras.io/api/layers/merging_layers/concatenate/ (visited on 07/08/2025).
- [70] Junyoung Chung et al. Empirical Evaluation of Gated Recurrent Neural Networks on Sequence Modeling. 2014. arXiv: [1412.3555](https://arxiv.org/abs/1412.3555) [cs.NE].
- [71] Yinghan Li et al. A Deep Learning Algorithm Based on CNN-LSTM Framework for Predicting Cancer Drug Sales Volume. 2025. arXiv: [2506.21927](https://arxiv.org/abs/2506.21927) [cs.CE].
- [72] R. Cont. “Empirical properties of asset returns: stylized facts and statistical issues”. In: *Quantitative Finance* 1.2 (2001), pp. 223–236.
- [73] Investopedia.com - What Is a Trading Strategy? How to Develop One. 2024. URL: <https://www.investopedia.com/terms/t/trading-strategy.asp> (visited on 07/12/2025).

- [74] Investiopia.com - 10 Steps to Building a Winning Trading Plan. 2024. URL: <https://www.investopedia.com/articles/trading/04/042104.asp> (visited on 07/12/2025).
- [75] Yunus Emre Akdoğan. “WHEN MACHINES LEARN TECHNICAL ANALYSIS: AN APPLICATION ON TECHNICAL ANALYSIS WITH MACHINE LEARNING IN BORSA ISTANBUL”. In: *Trakya Üniversitesi Sosyal Bilimler Dergisi* 27 (Mar. 2025), pp. 275–302. DOI: [10.26468/trakyasobed.1514346](https://doi.org/10.26468/trakyasobed.1514346).
- [76] Sabyasachi Mazumder, Sayan Neogy, and Sahana Das. “Simple Moving Average (SMA) Crossover Strategy with Buy Sell Indicator”. In: *Asia-Pacific Journal of Management and Technology* 03 (Apr. 2023), pp. 26–40. DOI: [10.46977/apjmt.2023.v03i04.004](https://doi.org/10.46977/apjmt.2023.v03i04.004).
- [77] medium.com - SMA Dual Moving Average Trading Strategy. 2024. URL: https://medium.com/@redsword_23261/sma-dual-moving-average-trading-strategy-c1b764fe809c (visited on 07/14/2025).
- [78] investopia.com - How to Use a Moving Average to Buy Stocks. 2025. URL: <https://www.investopedia.com/articles/active-trading/052014/how-use-moving-average-buy-stocks.asp> (visited on 07/14/2025).
- [79] howtotrade.com - The Triple Moving Average Crossover Strategy (With Examples). 2025. URL: <https://howtotrade.com/trading-strategies/triple-moving-average-crossover/> (visited on 07/15/2025).
- [80] Chun-Hao Chen et al. “An Advanced Optimization Approach for Long-Short Pairs Trading Strategy Based on Correlation Coefficients and Bollinger Bands”. In: *Applied Sciences* 12 (Jan. 2022), p. 1052. DOI: [10.3390/app12031052](https://doi.org/10.3390/app12031052).
- [81] oracle.com - ServiceLoader (Java SE 17 & JDK 17). 2025. URL: <https://docs.oracle.com/javase/8/docs/api/java/util/ServiceLoader.html> (visited on 06/26/2025).
- [82] ig.com - Slippage definition. 2025. URL: <https://www.ig.com/uk/glossary-trading-terms/slippage-definition> (visited on 06/26/2025).
- [83] lbank.com - How to calculate Profit and Loss (PnL) PnL%? 2024. URL: <https://www.lbank.com/support/articles/37004050587673> (visited on 06/26/2025).
- [84] ig.com - Trailing stop orders definition. 2025. URL: <https://www.ig.com/uk/glossary-trading-terms/trailing-stops-definition> (visited on 06/26/2025).
- [85] Jiarui Ni and Chengqi Zhang. “An Efficient Implementation of the Backtesting of Trading Strategies”. In: vol. 3758. Nov. 2005, pp. 126–131. DOI: [10.1007/11576235_17](https://doi.org/10.1007/11576235_17).
- [86] bybit-exchange.github.io - Get Kline. 2025. URL: <https://bybit-exchange.github.io/docs/v5/market/kline> (visited on 07/19/2025).
- [87] bybit-exchange.github.io - Kline. 2025. URL: <https://bybit-exchange.github.io/docs/v5/websocket/public/kline> (visited on 07/19/2025).

- [88] bybit-exchange.github.io - Get Wallet Balance. 2025. URL: <https://bybit-exchange.github.io/docs/v5/account/wallet-balance> (visited on 07/19/2025).
- [89] bybit-exchange.github.io - Place Order. 2025. URL: <https://bybit-exchange.github.io/docs/v5/order/create-order> (visited on 07/19/2025).
- [90] bybit-exchange.github.io - Get Position Info. 2025. URL: <https://bybit-exchange.github.io/docs/v5/position> (visited on 07/19/2025).
- [91] bybit-exchange.github.io - Get Closed PnL. 2025. URL: <https://bybit-exchange.github.io/docs/v5/position/close-pnl> (visited on 07/19/2025).
- [92] bybit-exchange.github.io - Integration Guidance. 2025. URL: <https://bybit-exchange.github.io/docs/v5/guide> (visited on 07/19/2025).
- [93] bybit-exchange.github.io - Connect. 2025. URL: <https://bybit-exchange.github.io/docs/v5/ws/connect> (visited on 07/19/2025).
- [94] Jiuzhou Fu and Dongfang Zhao. MPAD: A New Dimension-Reduction Method for Preserving Nearest Neighbors in High-Dimensional Vector Search. 2025. arXiv: [2504.16335](https://arxiv.org/abs/2504.16335) [cs.IR].

16 Appendix

16.1 Strategy Results on Validation Data




Figure 20: Classification AI Strategy Results




Figure 21: Dual Simple Moving Average Strategy Results




Figure 22: Triple Exponential Moving Average Strategy Results




Figure 23: Bollinger Bands Strategy Results

16.2 Out-of-Sample Strategy Results




Figure 24: Out-of-Sample Classification AI Strategy Results




Figure 25: Dual Simple Out-of-Sample Moving Average Strategy Results




Figure 26: Triple Exponential Out-of-Sample Moving Average Strategy Results

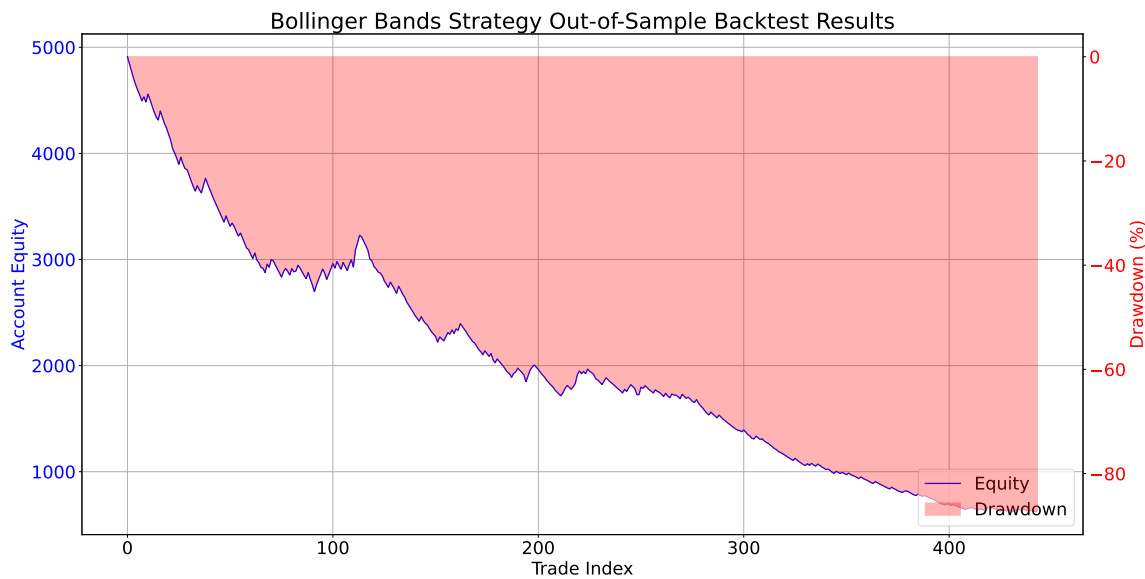


Figure 27: Out-of-Sample Bollinger Bands Strategy Results

16.3 Best Strategy Parameters

16.3.1 Classification AI Strategy Backtest Results

Regime	Downtrend + Low Volatility		Downtrend + High Volatility	
	$Q_{0.33}$	$Q_{0.66}$	$Q_{0.33}$	$Q_{0.66}$
Min. Probability for Entry	0.3	0.6	-	-
Trailing Stop Orders	false	false	-	-
Stop-Loss Delta	11	16	-	-
Take-Profit Delta	96	61	-	-
Number of Trades	29	30	-	-
Maximum Gain	518	371	-	-
Maximum Loss	-111	-95	-	-
Profits Std. Minimum	212	108	-	-
Equity Maximum	-47	0	-	-
Equity Last	3833	7978	-	-
Maximum Drawdown	3532	7978	-	-
Win Ratio	23	1	-	-
	48	96	-	-

Table 20: Classification AI Strategy Results I

Regime	Sideways Trend + Low Volatility		Sideways Trend + High Volatility	
	$Q_{0.33}$	$Q_{0.66}$	$Q_{0.33}$	$Q_{0.66}$
Min. Probability for Entry	0.3	0.3	-	-
Trailing Stop Orders	false	false	-	-
Stop-Loss Delta	11	11	-	-
Take-Profit Delta	81	81	-	-
Number of Trades	4	13	-	-
Maximum Gain	530	530	-	-
Maximum Loss	-109	-152	-	-
Profits Std. Minimum	257	255	-	-
Equity Maximum Last	-65	-301	-	-
Maximum Drawdown	701	1999	-	-
Win Ratio	592	1705	-	-
	15	14	-	-
	50	46	-	-

Table 21: Classification AI Strategy Results II

Regime	Uptrend + Low Volatility		Uptrend + High Volatility	
	$Q_{0.33}$	$Q_{0.66}$	$Q_{0.33}$	$Q_{0.66}$
Min. Probability for Entry	0.4	0.3	0.8	-
Trailing Stop Orders	false	true	false	-
Stop-Profit Delta	11	26	11	-
Take-Profit Delta	86	76	86	-
Number of Trades	23	58	56	-
Maximum Gain	568	250	469	-
Maximum Loss	-66	-53	-245	-
Profits Std.	279	125	268	-
Minimum	-193	-819	-619	-
Maximum	6275	6348	11467	-
Equity Last	6275	6348	9427	-
Maximum Drawdown	36	0	100	-
Win Ratio	60	68	55	-

Table 22: Classification AI Strategy Results III

16.3.2 Dual Simple Moving Average Strategy Backtest Results

Regime	Downtrend + Low Volatility		Downtrend + High Volatility	
	$Q_{0.33}$	$Q_{0.66}$	$Q_{0.33}$	$Q_{0.66}$
Short-Term Long-Term Swing Order N Swing Max Age Trailing Stop Order Take-Profit Delta Stop-Loss Delta	11	3	3	-
	13	19	5	-
	5	7	7	-
	10	20	40	-
	false	false	true	-
	140	155	95	-
Number of Trades Maximum Gain Maximum Loss Profits Std. Minimum Maximum Last Maximum Drawdown Win Ratio	65	50	65	-
	267	501	2	-
	135	211	218	-
	-81	-160	-41	-
	37	34	130	-
	-11	-122	-41	-
	1146	2250	177	-
	1098	2044	177	-
	108	259	0	-
	52	49	50	-

Table 23: Dual Simple Moving Average Strategy Results I

Regime	Sideways Trend + Low Volatility		Sideways Trend + High Volatility	
	$Q_{0.33}$	$Q_{0.66}$	$Q_{0.33}$	$Q_{0.66}$
Short-Term	3	3	3	5
Long-Term	19	19	5	11
Swing Order N	1	3	9	3
Swing Max Age	30	10	30	10
Trailing Stop Order	false	false	false	false
Take-Profit Delta	155	125	65	155
Stop-Loss Delta	65	5	20	35
Number of Trades	94	111	1	1
Maximum Gain	124	357	3	4
Maximum Loss	-40	-67	3	4
Profits Std.	23	37	0	0
Minimum	-11	-8	0	0
Maximum	585	418	3	4
Last	573	385	3	4
Maximum Drawdown	125	116	0	0
Win Ratio	45	33	100	100

Table 24: Dual Simple Moving Average Strategy Results II

Regime	Uptrend + Low Volatility		Uptrend + High Volatility	
	$Q_{0.33}$	$Q_{0.66}$	$Q_{0.33}$	$Q_{0.66}$
Short-Term	3	5	5	11
Long-Term	19	11	11	13
Swing Order N	5	7	7	3
Swing Max Age	30	10	10	20
Trailing Stop Order	false	false	false	true
Take-Profit Delta	200	185	155	140
Stop-Loss Delta	80	50	50	35
Number of Trades	1468	526	1	1
Maximum Gain	189	234	160	83
Maximum Loss	-121	-114	160	83
Profits Std.	27	52	0	0
Equity	-149	-218	0	0
Minimum	2860	5120	160	83
Maximum	2490	5025	160	83
Last	427	408	0	0
Maximum Drawdown	48	53	100	83
Win Ratio				

Table 25: Dual Simple Moving Average Strategy Results III

16.3.3 Triple Exponential Moving Average Strategy Backtest Results

	Regime	Downtrend + Low Volatility		Downtrend + High Volatility	
		$Q_{0.33}$	$Q_{0.66}$	$Q_{0.33}$	$Q_{0.66}$
EMA Period	Short-Term	3	9	7	5
	Medium-Term	5	14	14	23
	Long-Term	10	25	40	45
	Minimum EMA Slope	0.2	0.4	0.8	1.0
EMA Slope Window Length i	30	20	20	40	10
	Trailing Stop Order	false	false	false	false
Stop-Loss Distance		10	40	10	50
Take-Profit Distance		140	130	100	150
Number of Trades		531	159	6	3
Maximum Gain		475	174	205	111
Maximum Loss		-90	-78	-60	111
Profits Std.		70	75	123	0
	Minimum	-568	0	-54	0
Equity	Maximum	2323	2471	357	335
	Last	800	2088	191	335
Maximum Drawdown		151	97	46	0
Win Ratio		28	44	33	100

Table 26: Triple Exponential Moving Average Strategy Results I

Regime	Sideways Trend + Low Volatility		Sideways Trend + High Volatility	
	$Q_{0.33}$	$Q_{0.66}$	$Q_{0.33}$	$Q_{0.66}$
Short-Term	9	5	9	3
EMA Period	14	23	14	5
Medium-Term	25	45	25	10
Long-Term	0.4	1.0	0.4	0.2
Minimum EMA Slope	20	30	40	30
EMA Slope Window Length i	false	false	true	true
Trailing Stop Orders	20	10	10	30
Stop-Loss Distance	150	150	20	150
Take-Profit Distance				
Number of Trades	114	2	1	1
Maximum Gain	264	717	42	164
Maximum Loss	-75	-79	42	164
Profits Std.	79	398	0	0
Minimum	-627	0	0	0
Equity	916	717	42	164
Maximum	561	637	42	164
Last	2552	11	0	0
Maximum Drawdown	35	50	100	100
Win Ratio				

Table 27: Triple Exponential Moving Average Strategy Results II

Regime	Uptrend + Low Volatility		Uptrend + High Volatility	
	$Q_{0.33}$	$Q_{0.66}$	$Q_{0.33}$	$Q_{0.66}$
Short-Term	7	7	7	-
EMA Period	14	14	14	-
Medium-Term	40	40	40	-
Long-Term	0.8	0.8	0.8	-
Minimum EMA Slope	40	20	20	-
EMA Slope Window Length i	false	false	false	-
Trailing Stop Orders	10	10	10	-
Stop-Loss Distance	100	150	90	-
Take-Profit Distance				
Number of Trades	2	43	2	-
Maximum Gain	484	875	170	-
Maximum Loss	484	-135	170	-
Profits Std.	0	290	0	-
Minimum	0	-88	0	-
Maximum	968	2914	341	-
Equity	968	2075	341	-
Maximum Drawdown	0	58	0	-
Win Ratio	100	18	100	-

Table 28: Triple Exponential Moving Average Strategy Results III

16.3.4 Bollinger Bands Strategy Backtest Results

Regime	Downtrend + Low Volatility		Downtrend + High Volatility	
	$Q_{0.33}$	$Q_{0.66}$	$Q_{0.33}$	$Q_{0.66}$
Bollinger Band Period	14	18	20	15
No. of Standard Deviations	2.5	3.0	2	1.5
Trailing Stop Orders	true	false	false	false
Stop-Loss Distance	10	10	10	10
Take-Profit Delta	11	9	-1	1
Number of Trades	4	15	2	1
Maximum Gain	84	109	191	111
Maximum Loss	-27	-63	127	111
Profits Std. Minimum	44	50	32	0
Profits Std. Maximum	0	-9	0	0
Equity Last	174	250	319	111
Maximum Drawdown	174	250	319	111
Win Ratio	33	98	0	0
	75	53	100	100

Table 29: Bollinger Bands Strategy Results I

Regime	Sideways Trend + Low Volatility		Sideways Trend + High Volatility	
	$Q_{0.33}$	$Q_{0.66}$	$Q_{0.33}$	$Q_{0.66}$
Bollinger Band Period	15	23	-	23
No. of Standard Deviations	2.0	2.5	-	1.5
Trailing Stop Orders	false	false	-	false
Stop-Loss Distance	10	10	-	10
Take-Profit Delta	-3	15	-	1
Number of Trades	2	25	7	1
Maximum Gain	73	142	140	75
Maximum Loss	57	-64	-51	75
Profits Std. Minimum	7	52	67	0
Equity Maximum Last	0	-64	0	-31
Maximum Drawdown	130	398	340	75
Win Ratio	130	269	340	75
	0	52	51	0
	100	40	71	100

Table 30: Bollinger Bands Strategy Results II

Regime	Uptrend + Low Volatility		Uptrend + High Volatility	
	$Q_{0.33}$	$Q_{0.66}$	$Q_{0.33}$	$Q_{0.66}$
Bollinger Band Period	19	15	13	16
No. of Standard Deviations	1.5	1.5	1.5	1.5
Trailing Stop Orders	false	false	false	false
Stop-Loss Distance	30	10	10	10
Take-Profit Delta	9	1	15	9
Number of Trades	1	8	196	1
Maximum Gain	121	114	96	107
Maximum Loss	121	-57	-62	107
Profits Std. Minimum	0	66	25	0
Equity Maximum Last	0	-57	0	0
Maximum Drawdown	121	253	521	107
Win Ratio	121	253	128	107
	0	59	90	0
	100	62	36	100

Table 31: Bollinger Bands Strategy Results III

16.4 Other Tested Strategies

16.4.1 Other Tested Strategies Descriptions

Strategy Name	Short Description
Bollinger Bands with minimum slope strategy	Like the Bollinger Bands strategy, but with an additional SMA, which must show a minimum slope in order to open a position at all.
2 EMA with ATR strategy	Like the Dual SMA strategy but with EMAs. Additionally, an ATR is added which must show a minimum value in order to open a position at all.
2 EMA strategy	Like the Dual SMA strategy but with EMAs
2 EMA with RSI strategy	Like the Dual SMA strategy but with EMAs, Additionally, an RSI is added which must be greater than a threshold to open a long position and lower than a threshold to open a short position.
MACD strategy	If the MACD line crosses below the MACD signal line a short position is opened. Otherwise, if the MACD line crosses above the MACD signal line a long position is opened. The stop and limit distance are predefined but variable.
Parabolic SAR strategy	If the SAR value crosses below the current price a long position is opened. Otherwise, if the SAR value crosses above the current price a short position is opened. The stop loss is set at the SAR value at position opening and the limit level is set at double the distance from the current price to the current SAR value. Before a new position is opened by the described algorithm, all positions in the other direction are closed.
RSI strategy	If the RSI crosses above 70 a short position is opened. Otherwise, if the RSI drops below 30 a long position is opened. The stop and limit distances are predefined but variable.

Table 32: Description of Other Tested Strategies

16.4.2 Other Tested Strategies Results

Strategy Name		Bollinger Bands with minimum slope	2 EMA with ATR	2 EMA
Results for best strategy parameters	Number of Trades	20	901	1848
	Maximum Gain	168	149	138
	Maximum Loss	-162	-104	-55
	Profits Std.	81	34	17
	Minimum Equity	-19	-42	-32
	Maximum Equity	628	3075	1685
	Last Equity	628	2837	-31
	Maximum Drawdown	73	3	1301
	Win Ratio	65	37	30

Table 33: Results of Other Tested Strategies I

Strategy Name		2 EMA with RSI	MACD	Parabolic SAR	RSI
Results for best strategy parameters	Number of Trades	1022	1579	1082	2402
	Maximum Gain	159	397	82	187
	Maximum Loss	-82	-85	-214	-113
	Profits Std.	40	24	16	17
	Minimum Equity	138	-434	-2126	75
	Maximum Equity	5660	1723	30	2699
	Last Equity	5402	-433	-2126	2005
	Maximum Drawdown	66	125	7186	66
	Win Ratio	44	27	27	48

Table 34: Results of Other Tested Strategies II